

Hogwarts School of Witchcraft and Wizardry

Advanced Databases

Group B1

Oliwer Figura, Łukasz Machnik, Zofia Wiora

April 2025

1 Data model

1.1 Database Description

The Hogwarts School of Witchcraft and Wizardry Database System is designed to efficiently manage all aspects of the magical educational institution. This comprehensive system tracks students, teachers, academic records, house points, and extracurricular activities.

1.2 Purpose and Scope

This database serves multiple critical functions:

- Student information management including personal details, house affiliations, and dormitory assignments
- Keeping track of grades
- House points tracking for the annual House Cup competition
- Quidditch team roster management
- Teacher records and subject assignments

1.3 Database Schema

Figure 1 shows the Entity-Relationship diagram of the Hogwarts database. The schema illustrates the relationships between the nine primary tables that form the core of the system.

1.4 Key Entities and Relationships

The database consists of the following key tables:

1. **Students:** Contains personal information about each student, including their house and dormitory assignments.
2. **Houses:** The four houses of Hogwarts (Gryffindor, Hufflepuff, Ravenclaw, and Slytherin) with their attributes.
3. **Teachers:** Faculty information including employment dates and specializations.
4. **Grades:** Academic performance records with subject, teacher, student, and evaluation date.
5. **Subjects:** Course information including year level requirements and assigned classrooms.
6. **Students_Subjects:** Many-to-many relationship table connecting students to their enrolled subjects.
7. **Dormitories:** Housing information with gender and house affiliations.
8. **Points:** Records of house points awarded or deducted, with detailed attribution.
9. **Quidditch_Team_Members:** Manages Quidditch team rosters with positions and captain status.

2 Workload

2.1 Transactions querying the data

2.1.1 Grades analysis

This transaction takes all the students and for each of them displays the most interesting statistics, such as position in quidditch team, number of received points, and average of all grades.

Parameters: house_name, year_range_start, year_range_end, date_start, date_end, grade_range_start, grade_range_end, teacher_names

1. IF

- student's house_id value belongs to a house named house_name
- student's year value belongs to a range [year_range_start, year_range_end]
- in the time period [date_start, date_end] student obtained only grades from range [grade_range_start, grade_range_end]

2. SHOW Student's

- name
- surname
- gender
- number of received points in the specified time period
- position in quidditch team (if applicable, NULL otherwise)
- average of all grades awarded by teachers included in the list teacher_names in the specified time period
- (a) for this student for each subject calculate average grade in the specified time period
 - (b) SHOW name of subject with the highest average and this average
- (a) for this student for each subject calculate average grade in the specified time period
 - (b) SHOW name of subject with the lowest average and this average

3. ORDERED BY

- average grade
- descending

Tables used: Quidditch_Team_Members, Students, Houses, Points, Teachers, Grades, Subjects, Students_Subjects

2.1.2 House points summary

This transaction shows all houses and for each of them displays useful data connected to the House Cup competition.

Parameters: `houses_list`, `teachers_list`, `only_players`, `date_start`, `date_end`

1. IF

- house's name value belongs to a house listed in `houses_list`
- student is assigned to any quidditch team (only applicable if `only_players` is set to true)
- points were awarded in the specified time period

2. SHOW

- house name
- sum of points awarded by teachers listed in `teachers_list` collected by students assigned to this house
- a teacher that awarded the most points for this house
- a teacher that awarded the least points for this house
- a month in which the most points were awarded for this house
- a student which collected the most points for this house

3. ORDERED BY

- house name
- ascending

Tables used: `Quidditch_Team_Members`, `Students`, `Houses`, `Points`, `Teachers`

2.1.3 Best students display

This transaction is used to pick the best students divided by house, year and gender.

Parameters: `top_considered`, `date_start`, `date_end`

1. FOR EACH House

(a) FOR EACH Year

i. FOR EACH Gender

SHOW

- students ordered by points earned in specified time period
- limit number of displayed students to `top_considered`
- From this list SHOW
 - Student with the best average grade in the specified time period

Tables used: `Houses`, `Students`, `Points`

2.2 Transactions changing the data

2.2.1 Raising grades

This transaction is used to motivate students to earn points for their houses by rewarding academic improvements. Students who have received a lot of points are rewarded by raising their worst grades.

Parameters: `houses_list`, `teachers_list`, `min_points`, `raise_amount`, `number_of_affected`

For all grades received by all the students

IF

- student's house is in `houses_list`
- number of points received by student is higher than `min_points` (query)
- grade was assigned by one of the teachers from the `teachers_list`
- grade is lesser than maximum possible grade
- grade is one of the top `number_of_affected` worst grades received from the teacher from the `teachers_list` by the student (if the grades are the same, the older one is increased) (query)

UPDATE the grade by increasing its value by `raise_amount`.

Tables used: `Students`, `Houses`, `Teachers`, `Points`, `Grades`

2.2.2 Advancing students

This transaction would usually be performed at the beginning of a school year to move students to the next class, but only if they passed all the subjects. As a result of advancing to the next class, all students have to be assigned to new subjects and dormitories. Additionally, students finishing school need to be deleted, and new students need to be added.

Parameters: `new_students_list`, `minimal_average`, `max_dorm_capacity`

1. For all students

IF

- worst average of the subject of the student is higher than `minimal_average` (query based on 2.1.1)

THEN UPDATE

- increase student's year by one
- change student's dormitory value to NULL

2. For all students

IF

- student's year is higher than 7

THEN DELETE this student

3. INSERT new students from `new_students_list`
4. For all students and first dormitory fulfilling all conditions
IF

- value of student's dormitory is equal to NULL
- dormitory house is equal to student house
- dormitory year is equal to student year
- dormitory gender is equal to student gender
- list of students with dormitory value equal to dormitory id is not higher than `max_dorm_capacity` (query).

THEN UPDATE

- assign the dormitory to the student

5. DELETE all records from Points table
6. DELETE all records from Grades table
7. DELETE all records from Students_Subjects table
8. For all students and all subjects
IF

- student's year is equal to subject's year

INSERT student and subject to Students_Subjects

Tables used: Students, Dormitories, Quidditch_Team_Members, Points, Grades, Students_Subjects, Subjects

2.2.3 Points normalization

This transaction addresses potential imbalances in the house points system caused by varying teacher grading standards by applying the linear normalization of the points values.

Parameters: `min_points_value`, `max_points_value`, `start_date`, `end_date`, `teachers_list`

For all teachers

IF

- teacher is on the `teachers_list`

UPDATE value of the points by

- finding the lowest value of points awarded by the teacher between `start_date` and `end_date` (query)
- finding the highest value of points awarded by the teacher between `start_date` and `end_date` (query)
- applying linear normalization to all values of points awarded by the teacher between `start_date` and `end_date` (query) so they all fall in the range of `[min_points_value, max_points_value]`

Tables used: Points, Teachers

3 DBMS

We chose and installed the following DBMS: Oracle 21c Express Edition. We also installed Oracle SQL Developer - software that helps to manage the Oracle database and allows generating an image with all relations. At last, we created the whole database structure with SQL code visible below and generated the aforementioned image of all relations (Fig. 2)

```
CREATE TABLE Students (
    id NUMBER NOT NULL PRIMARY KEY,
    name VARCHAR2(64) NOT NULL,
    surname VARCHAR2(64) NOT NULL,
    gender CHAR NOT NULL,
    date_of_birth DATE NOT NULL,
    year NUMBER NOT NULL,
    hogsmeade_consent NUMBER(1) DEFAULT 0 NOT NULL,
    house_id NUMBER NOT NULL,
    dormitory_id NUMBER
);
```

```
CREATE TABLE Dormitories (
    id NUMBER NOT NULL PRIMARY KEY,
    gender CHAR NOT NULL,
    room_number NUMBER NOT NULL,
    house_id NUMBER NOT NULL
);
```

```
CREATE TABLE Houses (
    id NUMBER NOT NULL PRIMARY KEY,
    name VARCHAR2(16) NOT NULL,
    symbol VARCHAR2(16) NOT NULL,
    location VARCHAR2(128),
    teacher_id NUMBER
);
```

```
CREATE TABLE Quidditch_Team_Members (
    id NUMBER NOT NULL PRIMARY KEY,
    position VARCHAR2(64),
    is_captain NUMBER(1) DEFAULT 0 NOT NULL,
    student_id NUMBER NOT NULL
);
```

```
CREATE TABLE Teachers (
    id NUMBER NOT NULL PRIMARY KEY,
    name VARCHAR2(64) NOT NULL,
```

```
        surname VARCHAR2(64) NOT NULL,  
        date_of_birth DATE NOT NULL,  
        date_of_employment DATE NOT NULL  
    );  
  
CREATE TABLE Points (  
    id NUMBER NOT NULL PRIMARY KEY,  
    value NUMBER NOT NULL,  
    description VARCHAR2(1024),  
    award_date DATE NOT NULL,  
    student_id NUMBER NOT NULL,  
    teacher_id NUMBER  
);  
  
CREATE TABLE Grades (  
    id NUMBER NOT NULL PRIMARY KEY,  
    value VARCHAR2(1) NOT NULL,  
    award_date DATE NOT NULL,  
    student_id NUMBER NOT NULL,  
    subject_id NUMBER NOT NULL,  
    teacher_id NUMBER NOT NULL  
);  
  
CREATE TABLE Subjects (  
    id NUMBER NOT NULL PRIMARY KEY,  
    name VARCHAR2(64) NOT NULL,  
    classroom NUMBER NOT NULL,  
    year NUMBER NOT NULL,  
    teacher_id NUMBER NOT NULL  
);  
  
CREATE TABLE Students_Subjects (  
    student_id NUMBER NOT NULL,  
    subject_id NUMBER NOT NULL  
);  
  
ALTER TABLE Students  
    ADD FOREIGN KEY (house_id) REFERENCES Houses(id);  
  
ALTER TABLE Students  
    ADD FOREIGN KEY (dormitory_id) REFERENCES Dormitories(id);
```

```
ALTER TABLE Dormitories
    ADD FOREIGN KEY (house_id) REFERENCES Houses(id);

ALTER TABLE Houses
    ADD FOREIGN KEY (teacher_id) REFERENCES Teachers(id);

ALTER TABLE Quidditch_Team_Members
    ADD FOREIGN KEY (student_id) REFERENCES Students(id);

ALTER TABLE Points
    ADD FOREIGN KEY (student_id) REFERENCES Students(id);

ALTER TABLE Points
    ADD FOREIGN KEY (teacher_id) REFERENCES Teachers(id);

ALTER TABLE Grades
    ADD FOREIGN KEY (student_id) REFERENCES Students(id);

ALTER TABLE Grades
    ADD FOREIGN KEY (subject_id) REFERENCES Subjects(id);

ALTER TABLE Grades
    ADD FOREIGN KEY (teacher_id) REFERENCES Teachers(id);

ALTER TABLE Subjects
    ADD FOREIGN KEY (teacher_id) REFERENCES Teachers(id);

ALTER TABLE Students_Subjects
    ADD FOREIGN KEY (student_id) REFERENCES Students(id);

ALTER TABLE Students_Subjects
    ADD FOREIGN KEY (subject_id) REFERENCES Subjects(id);
```

4 Data

In order to populate the database, we've created the **Python script** generating required data.

4.1 Methods of generating the data

- **Students:** For each student, gender is randomly selected, and then a random name and surname are selected from the list of British names and surnames from Wikipedia. Date of birth is drawn from the range [2007-01-01, 2013-12-31] and the year is calculated accordingly. Students are assigned to houses randomly with a uniform distribution. Dormitories are assigned randomly in accordance with the constraints of the student's house, gender, and year, so that in one dormitory there are no more students than 6.
- **Houses:** There are four houses statically generated.
- **Teachers:** For each teacher, a name and surname are randomly selected from the list of British names and surnames from Wikipedia. Then, the date of birth is drawn from the range [1900-01-01, 2000-01-01]. The date of employment is also selected randomly, but the employee can be no younger than 24 to become a teacher.
- **Grades:** Grades are generated with a normal distribution (mean = 3, standard deviation = 1.25, min value = 0, max value = 5) and then translated to corresponding letters. Students and teachers are assigned randomly, taking into account that teacher and student must match the subject. The award date is drawn from the range [2023-09-01, 2024-06-30].
- **Subjects:** Subjects are generated from the provided list of subjects, assigned to corresponding years. Classroom and teacher are selected randomly.
- **Students_Subjects:** This table is generated based on the Students and Subjects table, so that each student is assigned to all subjects that have the same year.
- **Dormitories:** For each house, gender, and year, there are between 58 and 62 dormitories. Room number consists of year number and identification number which can occur only once in the house.
- **Points:** The value of points is randomly generated with a normal distribution (mean = 0, standard deviation = 5) and rounded to integer values. If the value is equal to 0, it is omitted. Students and teachers are assigned randomly, and the description is generated based on the name of the student and information if the number of received points is negative or positive. The content of the description is prepared using randomly selected words with various probabilities, so they feel natural to the reader. The award date is drawn from the range [2023-09-01, 2024-06-30].
- **Quidditch_Team_Members:** Every student is with probability 1% assigned to be a member of Quidditch team, then the roles are distributed with one of the members being the captain.

4.2 Number of records in each table

For each script run, the number of records generated may vary slightly. In the table, there are exact values shown for the data generated in one of the runs.

- **Students:** $10\,000 \pm 10\%$
- **Houses:** 4
- **Teachers:** $1000 \pm 10\%$
- **Grades:** from 10 to 30 grades for each subject assigned to a student

Table	Rows
Students	9 346
Houses	4
Teachers	1097
Grades	1 977 314
Subjects	76
Students_Subjects	101 583
Dormitories	3 369
Points	2 986 448
Quidditch_Team_Members	97

Table 1: Tables dimensions

- **Subjects:** 76
- **Students_Subjects:** from 10 to 12 subjects assigned to a student
- **Dormitories:** from 58 to 62 dormitories for each house, gender, and year
- **Points:** $3\ 000\ 000 \pm 10\%$
- **Quidditch_Team_Members:** about 100

4.3 Distribution of data

While generating data for **Points** and **Grades** tables, we assumed that the occurrence of each grade/point value is subject to a normal distribution.

5 Database workload

5.1 Transactions querying the data

All execution times are for default parameters.

5.1.1 Grades analysis

Query displaying students' statistics.

```

1  SELECT q.name,
2      q.surname,
3      q.gender,
4      q.total_points,
5      q.position,
6      g1.symbol AS best_grade,
```

```
7     s1.name AS best_subject,
8     g2.symbol AS worst_grade,
9     s2.name AS worst_subject
10    FROM (
11        SELECT students.name,
12              students.surname,
13              students.gender,
14              SUM(points.value) AS total_points,
15              quidditch_team_members.position,
16              MAX(averages.rounded_avg) KEEP (DENSE_RANK FIRST ORDER BY averages.accurate_avg DESC) AS max_grade,
17              MAX(averages.subject_id) KEEP (DENSE_RANK FIRST ORDER BY averages.accurate_avg DESC) AS best_subject,
18              MAX(averages.rounded_avg) KEEP (DENSE_RANK FIRST ORDER BY averages.accurate_avg ASC) AS min_grade,
19              MAX(averages.subject_id) KEEP (DENSE_RANK FIRST ORDER BY averages.accurate_avg ASC) AS worst_subject
20        FROM students
21          INNER JOIN houses
22            ON students.house_id = houses.id
23          INNER JOIN points
24            ON students.id = points.student_id
25          INNER JOIN grades
26            ON students.id = grades.student_id
27          LEFT JOIN quidditch_team_members
28            ON students.id = quidditch_team_members.student_id
29          INNER JOIN (
30              SELECT grades.student_id,
31                  grades.subject_id,
32                  AVG(grades_enum.value) AS ACCURATE_AVG,
33                  ROUND(AVG(grades_enum.value), 0) AS ROUNDED_AVG
34            FROM grades
35              INNER JOIN grades_enum
36                ON grades.value = grades_enum.symbol
37              GROUP BY grades.student_id,
38                  grades.subject_id
39          ) averages ON averages.student_id = students.id
40          WHERE houses.name = 'Gryffindor'
41          AND students.year >= 1
42          AND students.year <= 7
43          AND points.award_date >= TO_DATE('2023-09-01', 'YYYY-MM-DD')
44          AND points.award_date <= TO_DATE('2023-09-07', 'YYYY-MM-DD')
45          AND students.id NOT IN (
46              SELECT students.id
47              FROM students
48                INNER JOIN grades
```

```
49         ON students.id = grades.student_id
50     INNER JOIN grades_enum
51         ON grades.value = grades_enum.symbol
52 WHERE grades.award_date >= TO_DATE('2023-09-01', 'YYYY-MM-DD')
53     AND grades.award_date <= TO_DATE('2023-09-07', 'YYYY-MM-DD')
54     AND (
55         grades_enum.value <
56         (
57             SELECT grades_enum.value
58             FROM grades_enum
59             WHERE grades_enum.symbol = 'T'
60         )
61         OR grades_enum.value >
62         (
63             SELECT grades_enum.value
64             FROM grades_enum
65             WHERE grades_enum.symbol = 'O'
66         )
67     )
68 )
69 GROUP BY students.id,
70 students.name,
71 students.surname,
72 students.gender,
73 quidditch_team_members.position
74 ) q
75     INNER JOIN grades_enum g1
76         ON q.max_grade = g1.value
77     INNER JOIN grades_enum g2
78         ON q.min_grade = g2.value
79     INNER JOIN subjects s1
80         ON q.best_subject = s1.id
81     INNER JOIN subjects s2
82         ON q.worst_subject = s2.id
83 ORDER BY 4 DESC;
```

5.1.2 House points summary

Query displaying points summary for all houses.

```
1  SELECT houses.name,
2      SUM(points.value),
3      t1.name AS BEST_TEACHER_NAME,
4      t1.surname AS BEST_TEACHER_SURNAME,
5      t2.name AS WORST_TEACHER_NAME,
6      t2.surname AS WORST_TEACHER_SURNAME,
7      m.best_month, s.best_student_name,
8      s.best_student_surname
9  FROM houses
10 INNER JOIN students
11     ON houses.id = students.house_id
12 INNER JOIN points
13     ON students.id = points.student_id
14 INNER JOIN (
15     SELECT q.house_id,
16         MAX(q.teacher_id) KEEP (DENSE_RANK FIRST ORDER BY q.total_points DESC) AS FAVOURITE_TEACHER,
17         MAX(q.teacher_id) KEEP (DENSE_RANK FIRST ORDER BY q.total_points ASC) AS LEAST_FAVOURITE_TEACHER
18     FROM (
19         SELECT students.house_id,
20             points.teacher_id,
21             SUM(points.value) as TOTAL_POINTS
22         FROM students
23             INNER JOIN points
24                 ON students.id = points.student_id
25         WHERE points.award_date >= TO_DATE('2023-09-01', 'YYYY-MM-DD')
26             AND points.award_date <= TO_DATE('2024-06-30', 'YYYY-MM-DD')
27         GROUP BY students.house_id,
28                 points.teacher_id
29     ) q
30     GROUP BY q.house_id
31 ) t
32     ON t.house_id = houses.id
33 INNER JOIN (
34     SELECT q.house_id,
35         MAX(q.award_month) KEEP (DENSE_RANK FIRST ORDER BY q.total_points DESC) AS BEST_MONTH
36     FROM (
37         SELECT students.house_id,
38             EXTRACT(MONTH FROM points.award_date) AS AWARD_MONTH,
39             SUM(points.value) AS TOTAL_POINTS
40         FROM points
41             INNER JOIN students
42                 ON points.student_id = students.id
```

```
43     WHERE points.award_date >= TO_DATE('2023-09-01', 'YYYY-MM-DD')
44     AND points.award_date <= TO_DATE('2024-06-30', 'YYYY-MM-DD')
45     GROUP BY students.house_id,
46             EXTRACT(MONTH FROM points.award_date)
47   ) q
48   GROUP BY q.house_id
49 ) m
50   ON m.house_id = houses.id
51 INNER JOIN teachers t1
52   ON t1.id = t.favourite_teacher
53 INNER JOIN teachers t2
54   ON t2.id = t.least_favourite_teacher
55 INNER JOIN (
56   SELECT students.house_id,
57         MAX(students.name) KEEP (DENSE_RANK FIRST ORDER BY q.total_points DESC) AS BEST_STUDENT_NAME,
58         MAX(students.surname) KEEP (DENSE_RANK FIRST ORDER BY q.total_points DESC) AS BEST_STUDENT_SURNAME
59   FROM students
60   INNER JOIN (
61       SELECT points.student_id,
62             SUM(points.value) AS TOTAL_POINTS
63       FROM points
64       WHERE points.award_date >= TO_DATE('2023-09-01', 'YYYY-MM-DD')
65       AND points.award_date <= TO_DATE('2024-06-30', 'YYYY-MM-DD')
66       GROUP BY points.student_id
67   ) q
68   ON students.id = q.student_id
69   GROUP BY students.house_id
70 ) s
71   ON s.house_id = houses.id
72 WHERE points.award_date >= TO_DATE('2023-09-01', 'YYYY-MM-DD')
73   AND points.award_date <= TO_DATE('2024-06-30', 'YYYY-MM-DD')
74 GROUP BY houses.name,
75       t1.name,
76       t1.surname,
77       t2.name,
78       t2.surname,
79       m.best_month,
80       s.best_student_name,
81       s.best_student_surname;
```

5.1.3 Best students display

Query selecting best students from the database.

```
1 CREATE OR REPLACE PROCEDURE display_best_students(
2     p_top_count IN NUMBER DEFAULT 10,
3     p_date_start IN DATE DEFAULT TO_DATE('2023-01-01', 'YYYY-MM-DD'),
4     p_date_end IN DATE DEFAULT TO_DATE('2023-12-31', 'YYYY-MM-DD')
5 ) AS
6     v_total_students NUMBER;
7     v_start_time TIMESTAMP;
8     v_end_time TIMESTAMP;
9     v_execution_time NUMBER;
10    BEGIN
11        v_start_time := SYSTIMESTAMP;
12        DECLARE
13            v_rank NUMBER := 1;
14        BEGIN
15            FOR student_rec IN (
16                SELECT
17                    s.name || ' ' || s.surname AS student_name,
18                    h.name AS house_name,
19                    s.year,
20                    ROUND(AVG(ge.value), 2) AS avg_grade,
21                    NVL(SUM(p.value), 0) AS total_points
22                FROM
23                    Students s
24                    JOIN Houses h ON s.house_id = h.id
25                    LEFT JOIN Grades g ON s.id = g.student_id
26                    AND g.award_date BETWEEN p_date_start AND p_date_end
27                    LEFT JOIN Grades_Enum ge ON g.value = ge.symbol
28                    LEFT JOIN Points p ON s.id = p.student_id
29                    AND p.award_date BETWEEN p_date_start AND p_date_end
30                GROUP BY
31                    s.name, s.surname, h.name, s.year
32                ORDER BY
33                    avg_grade DESC, total_points DESC
34                FETCH FIRST p_top_count ROWS ONLY
35            ) LOOP
36                DBMS_OUTPUT.PUT_LINE(
37                    LPAD(v_rank, 4) || ' ' ||
38                    RPAD(student_rec.student_name, 26) || ' ' ||
39                    RPAD(student_rec.house_name, 10) || ' ' ||
```

```

40         LPAD(student_rec.year, 4) || ' ' || 
41         LPAD(TO_CHAR(student_rec.avg_grade, '90.99'), 9) || ' ' || 
42         LPAD(TO_CHAR(student_rec.total_points), 6)
43     );
44
45     v_rank := v_rank + 1;
46 END LOOP;
47
48
49 SELECT COUNT(*) INTO v_total_students FROM Students;
50
51 v_end_time := SYSTIMESTAMP;
52 v_execution_time := EXTRACT(SECOND FROM (v_end_time - v_start_time)) +
53                     EXTRACT(MINUTE FROM (v_end_time - v_start_time)) * 60;
54
END;

```

5.1.4 Best students by average grade

Query displaying top students by average grade.

```

1 WITH top_students AS (
2     SELECT
3         ROW_NUMBER() OVER (ORDER BY ROUND(AVG(ge.value), 2) DESC, NVL(SUM(p.value), 0) DESC) AS rank,
4         s.name || ' ' || s.surname AS student_name,
5         h.name AS house_name,
6         s.year,
7         ROUND(AVG(ge.value), 2) AS avg_grade,
8         NVL(SUM(p.value), 0) AS total_points
9     FROM
10        Students s
11        JOIN Houses h ON s.house_id = h.id
12        LEFT JOIN Grades g ON s.id = g.student_id
13            AND g.award_date BETWEEN TO_DATE('&date_start', 'YYYY-MM-DD') AND TO_DATE('&date_end', 'YYYY-MM-DD')
14        LEFT JOIN Grades_Enum ge ON g.value = ge.symbol
15        LEFT JOIN Points p ON s.id = p.student_id
16            AND p.award_date BETWEEN TO_DATE('&date_start', 'YYYY-MM-DD') AND TO_DATE('&date_end', 'YYYY-MM-DD')
17    GROUP BY
18        s.name, s.surname, h.name, s.year
19 )
20 SELECT
21     rank,
22     student_name,

```

```
23     house_name,
24     year,
25     avg_grade,
26     total_points
27 FROM
28   top_students
29 WHERE
30   rank <= &top_count
31 ORDER BY
32   rank;
```

5.2 Transactions changing the data

5.2.1 Raising grades

Transaction raising students' grades based on earned points.

```
1  CREATE OR REPLACE PROCEDURE raise_student_grades(
2    p_houses VARCHAR2,
3    p_teachers VARCHAR2,
4    p_min_points NUMBER,
5    p_raise_amount NUMBER,
6    p_num_affected NUMBER
7 ) AS
8   v_start TIMESTAMP;
9   v_end TIMESTAMP;
10
11  v_count NUMBER := 0;
12  v_students_found NUMBER := 0;
13  v_teachers_found NUMBER := 0;
14  v_display_limit CONSTANT NUMBER := 5;
15
16  v_houses VARCHAR2(1000);
17  v_teachers VARCHAR2(1000);
18
19  PROCEDURE short_delay IS
20    v_dummy NUMBER;
21  BEGIN
22    FOR i IN 1..10 LOOP
23      SELECT 1 INTO v_dummy FROM dual;
24    END LOOP;
25  END short_delay;
```

```
26
27 BEGIN
28     v_start := SYSTIMESTAMP;
29     v_houses := ',' || p_houses || ',';
30     v_teachers := ',' || p_teachers || ',';
31
32     DBMS_OUTPUT.PUT_LINE('Processing houses: ' || p_houses);
33
34     FOR r IN (
35         SELECT
36             s.id AS student_id,
37             s.name || ' ' || s.surname AS student_name,
38             h.name AS house_name,
39             NVL(SUM(p.value), 0) AS total_points
40         FROM
41             Students s
42             JOIN Houses h ON s.house_id = h.id
43             LEFT JOIN Points p ON s.id = p.student_id
44         WHERE
45             INSTR(v_houses, ',' || h.name || ',') > 0
46         GROUP BY
47             s.id, s.name, s.surname, h.name
48         HAVING
49             NVL(SUM(p.value), 0) >= p_min_points
50         ORDER BY
51             total_points DESC
52     ) LOOP
53         v_students_found := v_students_found + 1;
54
55         IF v_students_found <= v_display_limit THEN
56             DBMS_OUTPUT.PUT_LINE('Processing student #' || v_students_found || ': '
57             || r.student_name ||
58             ' (' || r.house_name || ', ' || r.total_points || ' points)');
59         ELSIF v_students_found = v_display_limit + 1 THEN
60             DBMS_OUTPUT.PUT_LINE('... Processing remaining students ...');
61         END IF;
62
63         short_delay;
64
65         DECLARE
66             v_teacher_count NUMBER := 0;
67             v_student_grades_updated NUMBER := 0;
```

```
68      BEGIN
69          FOR t IN (
70              SELECT id, name || ' ' || surname AS teacher_name
71                  FROM Teachers
72                  WHERE p_teachers = 'ALL' OR INSTR(v_teachers, ',' || id || ',',) > 0
73                  FETCH FIRST 5 ROWS ONLY
74          ) LOOP
75              v_teachers_found := v_teachers_found + 1;
76              v_teacher_count := v_teacher_count + 1;
77
78      DECLARE
79          v_grades_updated NUMBER := 0;
80      BEGIN
81          FOR g IN (
82              SELECT
83                  g.id AS grade_id,
84                  g.value AS current_grade,
85                  ge.value AS numeric_value,
86                  CASE
87                      WHEN g.subject_id IS NOT NULL THEN (
88                          SELECT name
89                              FROM Subjects
90                              WHERE id = g.subject_id)
91                      ELSE 'Unknown Subject'
92                  END AS subject_name
93          FROM
94              Grades g
95              JOIN Grades_Enum ge ON g.value = ge.symbol
96          WHERE
97              g.student_id = r.student_id
98              AND g.teacher_id = t.id
99              AND g.value != '0'
100         ORDER BY
101             ge.value ASC, g.award_date ASC
102             FETCH FIRST p_num_affected ROWS ONLY
103         ) LOOP
104             short_delay;
105
106         DECLARE
107             v_new_grade VARCHAR2(1);
108         BEGIN
109             v_new_grade :=
```

```
110          CASE g.current_grade
111              WHEN 'T' THEN 'D'
112              WHEN 'D' THEN 'P'
113              WHEN 'P' THEN 'A'
114              WHEN 'A' THEN 'E'
115              WHEN 'E' THEN 'O'
116              ELSE g.current_grade
117          END;
118
119      UPDATE Grades
120          SET value = v_new_grade
121          WHERE id = g.grade_id;
122
123      v_count := v_count + SQL%ROWCOUNT;
124      v_grades_updated := v_grades_updated + SQL%ROWCOUNT;
125      v_student_grades_updated := v_student_grades_updated
126          + SQL%ROWCOUNT;
127      END;
128  END LOOP;
129  IF v_grades_updated > 0 AND v_students_found <= v_display_limit
130  THEN
131      DBMS_OUTPUT.PUT_LINE('    - Updated ' || v_grades_updated ||
132                      ' grades for teacher: ' || t.teacher_name);
133  END IF;
134  END;
135  END LOOP;
136  IF v_teacher_count > 0 AND v_students_found <= v_display_limit THEN
137      DBMS_OUTPUT.PUT_LINE('    Processed ' || v_teacher_count
138                      || ' teachers for this student');
139  IF v_student_grades_updated > 0 THEN
140      DBMS_OUTPUT.PUT_LINE('    Total grades updated for this student: '
141                      || v_student_grades_updated);
142  END IF;
143  END IF;
144  END;
145  IF MOD(v_students_found, 10) = 0 THEN
146      DBMS_OUTPUT.PUT_LINE('Progress: ' || v_students_found
147                      || ' students processed, ' ||
148                      v_count || ' grades updated so far (' ||
149                      TO_CHAR(ROUND(EXTRACT(SECOND FROM
150                      (SYSTIMESTAMP - v_start)), 1)) ||
151                      ' seconds elapsed');
```

```

152     END IF;
153 END LOOP;

154
155 IF v_students_found = 0 THEN
156   DBMS_OUTPUT.PUT_LINE('No eligible students ');
157 END IF;

158
159 v_end := SYSTIMESTAMP;

160
161 DBMS_OUTPUT.PUT_LINE('==== EXECUTION SUMMARY ===');
162 DBMS_OUTPUT.PUT_LINE('Students found and processed: ' || v_students_found);
163 DBMS_OUTPUT.PUT_LINE('Teachers processed: ' || v_teachers_found);
164 DBMS_OUTPUT.PUT_LINE('Grades updated: ' || v_count);
165 DBMS_OUTPUT.PUT_LINE('Execution time: ' ||
166   TO_CHAR(ROUND(EXTRACT(SECOND FROM (v_end - v_start)) +
167     EXTRACT(MINUTE FROM (v_end - v_start)) * 60, 2)) ||
168   ' seconds');

169
170 COMMIT;
171 DBMS_OUTPUT.PUT_LINE('Changes committed.');

172
173 EXCEPTION
174   WHEN OTHERS THEN
175     DBMS_OUTPUT.PUT_LINE('Error: ' || SQLERRM);
176     ROLLBACK;
177     DBMS_OUTPUT.PUT_LINE('Changes rolled back.');
178 END;

```

5.2.2 Raising grades transaction

Transaction raising grades of students with high points.

```

1 MERGE INTO Grades g
2 USING (
3   WITH eligible_students AS (
4     SELECT
5       s.id AS student_id,
6       s.name || ' ' || s.surname AS student_name,
7       h.name AS house_name,
8       NVL(SUM(p.value), 0) AS total_points
9     FROM
10      Students s

```

```
11     JOIN Houses h ON s.house_id = h.id
12     LEFT JOIN Points p ON s.id = p.student_id
13   WHERE
14     INSTR('' || '&houses_list' || ',', ',' || h.name || ',') > 0
15   GROUP BY
16     s.id, s.name, s.surname, h.name
17   HAVING
18     NVL(SUM(p.value), 0) >= &min_points
19   ),
20   eligible_teachers AS (
21     SELECT id, name || ' ' || surname AS teacher_name
22     FROM Teachers
23   WHERE
24     UPPER('&teachers_list') = 'ALL' OR
25     INSTR('' || '&teachers_list' || ',', ',' || id || ',') > 0
26   ),
27   grades_to_update AS (
28     SELECT
29       g.id AS grade_id,
30       g.student_id,
31       g.teacher_id,
32       g.value AS current_grade,
33       CASE g.value
34         WHEN 'T' THEN 'D'
35         WHEN 'D' THEN 'P'
36         WHEN 'P' THEN 'A'
37         WHEN 'A' THEN 'E'
38         WHEN 'E' THEN 'O'
39         ELSE g.value
40       END AS new_grade,
41       ROW_NUMBER() OVER (
42         PARTITION BY g.student_id
43         ORDER BY ge.value ASC, DBMS_RANDOM.VALUE) AS grade_rank
44   FROM
45     Grades g
46     JOIN Grades_Enum ge ON g.value = ge.symbol
47     JOIN eligible_students es ON g.student_id = es.student_id
48     JOIN eligible_teachers et ON g.teacher_id = et.id
49   WHERE
50     g.value != 'O'
51     AND g.value != 'E'
52   )
```

```

53   SELECT
54     grade_id,
55     new_grade
56   FROM
57     grades_to_update
58   WHERE
59     grade_rank <= &number_of_affected
60   ) src
61   ON (g.id = src.grade_id)
62   WHEN MATCHED THEN
63   UPDATE SET g.value = src.new_grade;

```

5.2.3 Advancing students

Transaction moving students to the next class.

```

1 CREATE OR REPLACE PROCEDURE ADVANCING_STUDENTS
2 (
3   MIN_AVERAGE IN NUMBER DEFAULT 3.5
4 , MAX_DORM_CAPACITY IN NUMBER DEFAULT 6
5 ) AS
6   time_start NUMBER;
7   time_end NUMBER;
8   time_elapsed NUMBER;
9 BEGIN
10   time_start := DBMS_UTILITY.GET_TIME;
11   -- Advancing eligible students and removing them from dormitory
12   UPDATE students
13   SET students.year = students.year + 1, students.dormitory_id = NULL
14   WHERE students.id NOT IN (
15     SELECT DISTINCT failed.student_id
16     FROM (
17       SELECT grades.student_id, grades.subject_id, AVG(grades_enum.value)
18       FROM grades
19       INNER JOIN grades_enum ON grades.value = grades_enum.symbol
20       GROUP BY grades.student_id, grades.subject_id
21       HAVING AVG(grades_enum.value) < min_average
22     ) failed
23   );
24   -- Assigning students to a new dormitory
25   UPDATE students s
26   SET s.dormitory_id = (

```

```

27     SELECT available.id
28     FROM (
29         SELECT dormitories.id, COUNT(students.id) AS OCCUPANCY
30         FROM dormitories
31        LEFT JOIN students ON dormitories.id = students.dormitory_id
32        WHERE dormitories.year = s.year
33        AND dormitories.house_id = s.house_id
34        AND dormitories.gender = s.gender
35        HAVING COUNT(students.id) < max_dorm_capacity
36        GROUP BY dormitories.id
37        ORDER BY 2 ASC
38        FETCH FIRST ROW ONLY
39    ) available
40 )
41 WHERE s.dormitory_id IS NULL;
-- Deleting all grades, points and subject assignments
42 DELETE FROM points;
43 DELETE FROM grades;
44 DELETE FROM students_subjects;
-- Assigning all students to new subjects
45 INSERT INTO students_subjects (student_id, subject_id)
46     SELECT students.id, subjects.id
47     FROM students
48     INNER JOIN subjects ON students.year = subjects.year;
-- Calculate elapsed time
49 time_end := DBMS_UTILITY.GET_TIME;
50 time_elapsed := (time_end - time_start) / 100;
51 DBMS_OUTPUT.PUT_LINE('Execution time: ' || time_elapsed || ' seconds');
52
53 END ADVANCING_STUDENTS;

```

5.2.4 Points normalization

Transaction normalizing house points values.

```

1 UPDATE points pOuter
2     SET pOuter.value = ROUND(pOuter.value * 10 / (
3         SELECT MAX(pInner.value)
4             FROM points pInner
5             WHERE pInner.teacher_id = pouter.teacher_id
6     ))
7     WHERE pOuter.value > 0
8         AND pOuter.award_date >= TO_DATE('2023-09-01', 'YYYY-MM-DD')

```

```

9      AND pOuter.award_date <= TO_DATE('2024-06-30', 'YYYY-MM-DD')
10     AND teacher.id >= 0
11     AND teacher.id <= 100

```

5.2.5 Removing bad students' accomplishments

Transaction that removes all positive points acquired by students that didn't get average positive grades from 4 or more subjects.

```

1  DELETE FROM points
2  WHERE points.student_id IN(
3      SELECT q.student_id
4      FROM (
5          SELECT grades.student_id,
6              grades.subject_id,
7              AVG(grades_enum.value)
8          FROM grades
9              INNER JOIN grades_enum
10             ON grades.value = grades_enum.symbol
11            WHERE grades.award_date >= TO_DATE('2023-09-01', 'YYYY-MM-DD')
12            AND grades.award_date <= TO_DATE('2024-06-30', 'YYYY-MM-DD')
13            HAVING AVG(grades_enum.value) < 4
14            GROUP BY grades.student_id,
15                  grades.subject_id
16        ) q
17        HAVING COUNT(q.subject_id) >= 4
18        GROUP BY q.student_id
19    )
20    AND points.value > 0

```

6 Query plans

6.1 Transactions querying the data

6.1.1 Grades analysis

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time	
0	SELECT STATEMENT		3	882		144K (82)	00:00:06	
1	SORT ORDER BY		3	882		144K (82)	00:00:06	

2	NESTED LOOPS		3	882	144K (82)	00:00:06	
3	NESTED LOOPS		3	882	144K (82)	00:00:06	
4	NESTED LOOPS		3	741	144K (82)	00:00:06	
* 5	HASH JOIN		3	600	144K (82)	00:00:06	
* 6	HASH JOIN		3	555	144K (82)	00:00:06	
7	VIEW		3	510	144K (82)	00:00:06	
8	SORT GROUP BY		3	879	144K (82)	00:00:06	
* 9	HASH JOIN		829M	226G	70819 (63)	00:00:03	
* 10	TABLE ACCESS FULL	HOUSES	1	23	3 (0)	00:00:01	
11	MERGE JOIN ANTI		3316M	834G	48783 (47)	00:00:02	
12	MERGE JOIN		3316M	793G	46283 (49)	00:00:02	
13	MERGE JOIN		15M	3647M	12587 (3)	00:00:01	
14	MERGE JOIN OUTER		1937K	386M	2886 (6)	00:00:01	
15	MERGE JOIN		1937K	299M	2882 (6)	00:00:01	
16	SORT JOIN		1937K	96M	2632 (7)	00:00:01	
17	VIEW		1937K	96M	2632 (7)	00:00:01	
18	HASH GROUP BY		1937K	79M	2632 (7)	00:00:01	
* 19	HASH JOIN		1937K	79M	2510 (2)	00:00:01	
20	TABLE ACCESS FULL	GRADES_ENUM	6	90	3 (0)	00:00:01	
21	TABLE ACCESS FULL	GRADES	1937K	51M	2494 (2)	00:00:01	
* 22	SORT JOIN		9156	983K	2344K	250 (1)	00:00:01
* 23	TABLE ACCESS FULL	STUDENTS	9156	983K	19 (0)	00:00:01	
* 24	SORT JOIN		97	4559	4 (25)	00:00:01	
25	TABLE ACCESS FULL	QUIDDITCH_TEAM_MEMBERS	97	4559	3 (0)	00:00:01	
* 26	SORT JOIN		77469	2647K	7304K	9610 (1)	00:00:01
* 27	TABLE ACCESS FULL	POINTS	77469	2647K	8886 (1)	00:00:01	
* 28	SORT JOIN		1937K	24M	74M	11766 (2)	00:00:01
29	TABLE ACCESS FULL	GRADES	1937K	24M	2491 (2)	00:00:01	
* 30	SORT UNIQUE		4804	62452	2500 (2)	00:00:01	
31	VIEW	VW_NSO_1	4804	62452	2499 (2)	00:00:01	
* 32	HASH JOIN		4804	182K	2499 (2)	00:00:01	
* 33	TABLE ACCESS FULL	GRADES_ENUM	1	15	3 (0)	00:00:01	
* 34	TABLE ACCESS FULL	GRADES_ENUM	1	15	3 (0)	00:00:01	
* 35	TABLE ACCESS FULL	GRADES_ENUM	1	15	3 (0)	00:00:01	
* 36	TABLE ACCESS FULL	GRADES	49275	1154K	2496 (2)	00:00:01	
37	TABLE ACCESS FULL	GRADES_ENUM	6	90	3 (0)	00:00:01	
38	TABLE ACCESS FULL	GRADES_ENUM	6	90	3 (0)	00:00:01	
39	TABLE ACCESS BY INDEX ROWID	SUBJECTS	1	47	1 (0)	00:00:01	
* 40	INDEX UNIQUE SCAN	SYS_C008771	1		0 (0)	00:00:01	
* 41	INDEX UNIQUE SCAN	SYS_C008771	1		0 (0)	00:00:01	
42	TABLE ACCESS BY INDEX ROWID	SUBJECTS	1	47	1 (0)	00:00:01	

Predicate Information (identified by operation id):

```

5 - access("Q"."MIN_GRADE"="G2"."VALUE")
6 - access("Q"."MAX_GRADE"="G1"."VALUE")
9 - access("STUDENTS"."HOUSE_ID"="HOUSES"."ID")
10 - filter("HOUSES"."NAME"='Gryffindor')
19 - access("GRADES"."VALUE"="GRADES_ENUM"."SYMBOL")
22 - access("AVERAGES"."STUDENT_ID"="STUDENTS"."ID")
      filter("AVERAGES"."STUDENT_ID"="STUDENTS"."ID")

```

```

23 - filter("STUDENTS"."YEAR">>=1 AND "STUDENTS"."YEAR"><=7)
24 - access("STUDENTS"."ID"="QUIDDITCH_TEAM_MEMBERS"."STUDENT_ID"(+))
    filter("STUDENTS"."ID"="QUIDDITCH_TEAM_MEMBERS"."STUDENT_ID"(+))
26 - access("STUDENTS"."ID"="POINTS"."STUDENT_ID")
    filter("STUDENTS"."ID"="POINTS"."STUDENT_ID")
27 - filter("POINTS"."AWARD_DATE">>=TO_DATE(' 2023-09-01 00:00:00', 'yyyy-mm-dd hh24:mi:ss') AND
    "POINTS"."AWARD_DATE"><=TO_DATE(' 2023-09-07 00:00:00', 'yyyy-mm-dd hh24:mi:ss'))
28 - access("STUDENTS"."ID"="GRADES"."STUDENT_ID")
    filter("STUDENTS"."ID"="GRADES"."STUDENT_ID")
30 - access("STUDENTS"."ID"="ID")
    filter("STUDENTS"."ID"="ID")
32 - access("GRADES"."VALUE"="GRADES_ENUM"."SYMBOL")
33 - filter("GRADES_ENUM"."VALUE">< (SELECT "GRADES_ENUM"."VALUE" FROM "GRADES_ENUM" "GRADES_ENUM" WHERE
    "GRADES_ENUM"."SYMBOL"='T') OR "GRADES_ENUM"."VALUE">> (SELECT "GRADES_ENUM"."VALUE" FROM "GRADES_ENUM"
    "GRADES_ENUM" WHERE "GRADES_ENUM"."SYMBOL"='O'))
34 - filter("GRADES_ENUM"."SYMBOL"='T')
35 - filter("GRADES_ENUM"."SYMBOL"='O')
36 - filter("GRADES"."AWARD_DATE">>=TO_DATE(' 2023-09-01 00:00:00', 'yyyy-mm-dd hh24:mi:ss') AND
    "GRADES"."AWARD_DATE"><=TO_DATE(' 2023-09-07 00:00:00', 'yyyy-mm-dd hh24:mi:ss'))
40 - access("Q"."BEST_SUBJECT"="S1"."ID")
41 - access("Q"."WORST_SUBJECT"="S2"."ID")

```

Note

- dynamic statistics used: dynamic sampling (level=2)
- this is an adaptive plan

Candidates for improvements:

Operations with highest cost: 0-8 (SELECT STATEMENT, SORT ORDER BY, NESTED LOOPS, HASH JOIN, VIEW, SORT GROUP BY)

Other operations with high cost: 9, 11-13, 28 (HASH JOIN, MERGE JOIN ANTI, MERGE JOIN, SORT JOIN)

6.1.2 House points summary

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT		2986K	535M		159K (1)	00:00:07
1	HASH GROUP BY		2986K	535M	583M	159K (1)	00:00:07
* 2	HASH JOIN		2986K	535M		36394 (3)	00:00:02
3	TABLE ACCESS FULL	TEACHERS	1097	20843		5 (0)	00:00:01
* 4	HASH JOIN		2986K	481M		36369 (3)	00:00:02
5	TABLE ACCESS FULL	TEACHERS	1097	20843		5 (0)	00:00:01
* 6	HASH JOIN		2986K	427M		36344 (3)	00:00:02
7	VIEW		4	64		9127 (4)	00:00:01
8	SORT GROUP BY		4	116		9127 (4)	00:00:01
9	VIEW		858	24882		9127 (4)	00:00:01
10	HASH GROUP BY		858	19734		9127 (4)	00:00:01
* 11	HASH JOIN		2986K	65M		8934 (2)	00:00:01
12	TABLE ACCESS FULL	STUDENTS	9346	65422		19 (0)	00:00:01
* 13	TABLE ACCESS FULL	POINTS	2986K	45M		8895 (1)	00:00:01
* 14	HASH JOIN		2986K	381M		27197 (3)	00:00:02

15	VIEW		4	116	9131	(4)	00:00:01
16	SORT GROUP BY		4	116	9131	(4)	00:00:01
17	VIEW		3103	89987	9131	(4)	00:00:01
18	HASH GROUP BY		3103	83781	9131	(4)	00:00:01
* 19	HASH JOIN		2986K	76M	8938	(2)	00:00:01
20	TABLE ACCESS FULL	STUDENTS	9346	65422	19	(0)	00:00:01
* 21	TABLE ACCESS FULL	POINTS	2986K	56M	8899	(1)	00:00:01
* 22	HASH JOIN		2986K	299M	18046	(3)	00:00:01
23	VIEW	VW_GBF_15	9346	812K	9131	(4)	00:00:01
* 24	HASH JOIN		9346	930K	9131	(4)	00:00:01
25	MERGE JOIN		4	380	9112	(4)	00:00:01
26	TABLE ACCESS BY INDEX ROWID	HOUSES	4	56	2	(0)	00:00:01
27	INDEX FULL SCAN	SYS_C008743	4		1	(0)	00:00:01
* 28	SORT JOIN		4	324	9110	(4)	00:00:01
29	VIEW		4	324	9109	(4)	00:00:01
30	SORT GROUP BY		4	156	9109	(4)	00:00:01
* 31	HASH JOIN		9346	355K	9108	(4)	00:00:01
32	VIEW		9346	155K	9088	(4)	00:00:01
33	HASH GROUP BY		9346	146K	9088	(4)	00:00:01
* 34	TABLE ACCESS FULL	POINTS	2986K	45M	8895	(1)	00:00:01
35	TABLE ACCESS FULL	STUDENTS	9346	200K	19	(0)	00:00:01
36	TABLE ACCESS FULL	STUDENTS	9346	65422	19	(0)	00:00:01
* 37	TABLE ACCESS FULL	POINTS	2986K	45M	8895	(1)	00:00:01

Predicate Information (identified by operation id):

```

2 - access("T2"."ID"="T"."LEAST_FAVOURITE_TEACHER")
4 - access("T1"."ID"="T"."FAVOURITE_TEACHER")
6 - access("M"."HOUSE_ID"="ITEM_2")
11 - access("POINTS"."STUDENT_ID"="STUDENTS"."ID")
13 - filter("POINTS"."AWARD_DATE">>=TO_DATE(' 2023-09-01 00:00:00', 'yyyy-mm-dd hh24:mi:ss') AND
      "POINTS"."AWARD_DATE"=<TO_DATE(' 2024-06-30 00:00:00', 'yyyy-mm-dd hh24:mi:ss'))
14 - access("T"."HOUSE_ID"="ITEM_2")
19 - access("STUDENTS"."ID"="POINTS"."STUDENT_ID")
21 - filter("POINTS"."AWARD_DATE">>=TO_DATE(' 2023-09-01 00:00:00', 'yyyy-mm-dd hh24:mi:ss') AND
      "POINTS"."AWARD_DATE"=<TO_DATE(' 2024-06-30 00:00:00', 'yyyy-mm-dd hh24:mi:ss'))
22 - access("ITEM_1"="POINTS"."STUDENT_ID")
24 - access("HOUSES"."ID"="STUDENTS"."HOUSE_ID")
28 - access("S"."HOUSE_ID"="HOUSES"."ID")
  filter("S"."HOUSE_ID"="HOUSES"."ID")
31 - access("STUDENTS"."ID"="Q"."STUDENT_ID")
34 - filter("POINTS"."AWARD_DATE">>=TO_DATE(' 2023-09-01 00:00:00', 'yyyy-mm-dd hh24:mi:ss') AND
      "POINTS"."AWARD_DATE"=<TO_DATE(' 2024-06-30 00:00:00', 'yyyy-mm-dd hh24:mi:ss'))
37 - filter("POINTS"."AWARD_DATE">>=TO_DATE(' 2023-09-01 00:00:00', 'yyyy-mm-dd hh24:mi:ss') AND
      "POINTS"."AWARD_DATE"=<TO_DATE(' 2024-06-30 00:00:00', 'yyyy-mm-dd hh24:mi:ss'))
```

Candidates for improvements:

Operations with highest cost: 0, 1 (SELECT STATEMENT, HASH GROUP BY)

Other operations with high cost: 2, 4, 6, 14 (HASH JOIN)

6.1.3 Best Students Display

This query ranks students based on their average grades and accumulated points within a specified time period. The query involves complex joins across multiple tables (Students, Houses, Grades, Grades_Enum, Points) and includes window functions for ranking.

Id	Operation	Name	Rows	Cost	Time
0	SELECT STATEMENT		10	1254	00:00:16
1	WINDOW SORT		10	1254	00:00:16
2	VIEW		9346	1250	00:00:16
* 3	FILTER				
4	WINDOW SORT		9346	1250	00:00:16
5	HASH GROUP BY		9346	1178	00:00:15
6	HASH JOIN		9346	978	00:00:12
7	TABLE ACCESS FULL	HOUSES	4	3	00:00:01
8	HASH JOIN		9346	975	00:00:12
9	TABLE ACCESS FULL	STUDENTS	9346	87	00:00:02
10	HASH JOIN OUTER		65422	882	00:00:11
11	HASH JOIN OUTER		65422	456	00:00:06
12	HASH JOIN		65422	240	00:00:03
* 13	TABLE ACCESS FULL	GRADES	65422	129	00:00:02
14	TABLE ACCESS FULL	GRADES_ENUM	6	2	00:00:01
* 15	TABLE ACCESS FULL	POINTS	65422	129	00:00:02

Predicate Information (identified by operation id):

```
-----
3 - filter(RANK<=:TOP_COUNT)
13 - filter(("G"."AWARD_DATE">>=TO_DATE(:DATE_START) AND "G"."AWARD_DATE"<=TO_DATE(:DATE_END)))
15 - filter(("P"."AWARD_DATE">>=TO_DATE(:DATE_START) AND "P"."AWARD_DATE"<=TO_DATE(:DATE_END)))
```

Candidates for improvements

- **Window Sort Operations (1, 4):** These operations have the highest cost and consume the most execution time. The first window sort is for the ROW_NUMBER function, the second to apply the final ranking.
- **Hash Group By (5):** With a cost of 1178, this operation aggregates data for calculating average grades and total points per student.
- **Multiple Hash Joins (6, 8, 10, 11, 12):** The operations join all tables involved in the query.
- **Full Table Scans (7, 9, 13, 14, 15):** The query performs full table scans on all five tables.
- **Date Range Filtering (13, 15):** Post-scan filtering on award_date columns.

6.2 Transactions changing the data

6.2.1 Raising Grades (UPDATE)

This transaction updates student grades based on specific criteria. It identifies students from selected houses with points above a threshold, then updates their worst grades to higher values. The transaction involves complex logic with multiple joins, sorting, and data modification.

Id	Operation	Name	Rows	Cost	Time
0	MERGE STATEMENT		4099	2240	00:00:28
1	MERGE	GRADES			
2	VIEW		4099	1785	00:00:22
3	WINDOW SORT PUSHED RANK		4099	1785	00:00:22
* 4	FILTER				
* 5	HASH JOIN		4099	1743	00:00:21
* 6	HASH JOIN		4099	1625	00:00:20
7	HASH JOIN		438	242	00:00:03
8	VIEW		438	139	00:00:02
9	HASH GROUP BY		438	139	00:00:02
* 10	HASH JOIN		585	119	00:00:02
* 11	TABLE ACCESS FULL	HOUSES	1	3	00:00:01
12	HASH JOIN OUTER		2336	115	00:00:02
13	TABLE ACCESS FULL	STUDENTS	9346	87	00:00:02
* 14	TABLE ACCESS FULL	POINTS	65422	28	00:00:01
15	VIEW		25	3	00:00:01
16	TABLE ACCESS FULL	TEACHERS	25	3	00:00:01
* 17	TABLE ACCESS FULL	GRADES	65422	1365	00:00:17
18	TABLE ACCESS FULL	GRADES_ENUM	6	3	00:00:01
19	BUFFER SORT		4099		
20	TABLE ACCESS BY USER ROWID	GRADES	4099	455	00:00:06

Predicate Information (identified by operation id):

```

4 - filter("GRADE_RANK" <= TO_NUMBER(:NUMBER_OF_AFFECTED))
5 - access("G"."VALUE" = "GE"."SYMBOL")
6 - access("G"."STUDENT_ID" = "ES"."STUDENT_ID" AND "G"."TEACHER_ID" = "ET"."ID")
10 - access("S"."HOUSE_ID" = "H"."ID")
11 - filter(INSTR(''||:HOUSES_LIST||',',','||"H"."NAME'||',') > 0)
14 - filter("P"."STUDENT_ID" = "S"."ID")
17 - filter("G"."VALUE" <> 'O' AND "G"."VALUE" <> 'E')

```

Candidates for improvements

- Merge Statement (0) and Window Sort Pushed Rank (3): This operations have the highest cost.

- Multiple Hash Joins (5, 6, 7, 10, 12)**: These operations join multiple tables, with operations 5 and 6 being particularly expensive.
- Full Table Access on GRADES (17)**: With a cost of 1365, this operation scans 65,422 grade records to find grades eligible for updating.
- Buffer Sort (19) and Table Access by User ROWID (20)**: These operations are part of the final MERGE update, accessing the actual grade records to be updated.
- Hash Group By (9)**: This operation aggregates points data to identify eligible students.

6.2.2 Assign subjects (INSERT)

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	INSERT STATEMENT		1016	52868		12318 (3)	00:00:01
1	LOAD TABLE CONVENTIONAL	STUDENTS SUBJECTS	1016	52868		12318 (3)	00:00:01
2	HASH UNIQUE						
3	UNION-ALL						
* 4	HASH JOIN		1015	26390		6155 (3)	00:00:01
* 5	HASH JOIN ANTI			93	1860	6152 (3)	00:00:01
* 6	TABLE ACCESS FULL	STUDENTS	9346	65422		19 (0)	00:00:01
7	VIEW	VW_NS0_1	25113	318K		6133 (3)	00:00:01
* 8	FILTER						
9	SORT GROUP BY		25113	343K	45M	6133 (3)	00:00:01
* 10	HASH JOIN			1977K	26M	2511 (2)	00:00:01
11	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01
12	TABLE ACCESS FULL	GRADES	1977K	16M		2495 (2)	00:00:01
* 13	TABLE ACCESS FULL	SUBJECTS	76	456		3 (0)	00:00:01
* 14	HASH JOIN ANTI			1	44	6158 (3)	00:00:01
* 15	HASH JOIN		97	3007		25 (0)	00:00:01
16	MERGE JOIN CARTESIAN		97	2328		6 (0)	00:00:01
* 17	TABLE ACCESS FULL	SUBJECTS	1	20		3 (0)	00:00:01
18	BUFFER SORT		97	388		3 (0)	00:00:01
19	TABLE ACCESS FULL	QUIDDITCH_TEAM_MEMBERS	97	388		3 (0)	00:00:01
* 20	TABLE ACCESS FULL	STUDENTS	8011	56077		19 (0)	00:00:01
* 21	TABLE ACCESS FULL	SUBJECTS	1	20		3 (0)	00:00:01
22	VIEW	VW_NS0_2	25113	318K		6133 (3)	00:00:01
* 23	FILTER						
24	SORT GROUP BY		25113	343K	45M	6133 (3)	00:00:01
* 25	HASH JOIN			1977K	26M	2511 (2)	00:00:01
26	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01
27	TABLE ACCESS FULL	GRADES	1977K	16M		2495 (2)	00:00:01

Predicate Information (identified by operation id):

```

4 - access("STUDENTS"."YEAR"="SUBJECTS"."YEAR")
5 - access("STUDENTS"."ID"="STUDENT_ID")
6 - filter("STUDENTS"."YEAR">>=1 AND "STUDENTS"."YEAR"<=7)
8 - filter(SUM("GRADES_ENUM"."VALUE")/COUNT("GRADES_ENUM"."VALUE")<3.5)
10 - access("GRADES"."VALUE"="GRADES_ENUM"."SYMBOL")
13 - filter("SUBJECTS"."YEAR">>=1 AND "SUBJECTS"."YEAR"<=7)

```

```

14 - access("STUDENTS"."ID"="STUDENT_ID")
15 - access("STUDENTS"."ID"="QUIDDITCH_TEAM_MEMBERS"."STUDENT_ID")
17 - filter("SUBJECTS"."NAME"='Flying')
20 - filter("STUDENTS"."YEAR"><> (SELECT "SUBJECTS"."YEAR" FROM "SUBJECTS" "SUBJECTS" WHERE
    "SUBJECTS"."NAME"='Flying'))
21 - filter("SUBJECTS"."NAME"='Flying')
23 - filter(SUM("GRADES_ENUM"."VALUE")/COUNT("GRADES_ENUM"."VALUE")<3.5)
25 - access("GRADES"."VALUE"="GRADES_ENUM"."SYMBOL")

```

Candidates for improvements

- Hash Unique (2)**: This operation has the highest cost.
- Multiple Hash Joins (4, 5, 10, 14, 15, 25)**: These operations join multiple tables, with operations 4, 5 and 14 being particularly expensive.
- Full Table Access (12, 27)**: Both with a cost of 2495 and accessing 1 977 000 rows are very costly.
- Sort Group By (9, 24)**: These operations are also one of the more expensive

6.2.3 Removing bad students' accomplishments (DELETE)

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	DELETE STATEMENT		74812	2264K		16247 (2)	00:00:01
1	DELETE	POINTS	74812	2264K		16247 (2)	00:00:01
* 2	HASH JOIN RIGHT SEMI	VW_NSO_1	468	6084		7356 (3)	00:00:01
3	VIEW		468	7956		7356 (3)	00:00:01
* 4	SORT GROUP BY		25113	416K		7356 (3)	00:00:01
5	VIEW		25113	539K	68M	7356 (3)	00:00:01
* 6	FILTER		1977K	41M		2524 (3)	00:00:01
7	HASH GROUP BY		6	30		3 (0)	00:00:01
8	HASH JOIN		1977K	32M		2508 (2)	00:00:01
9	TABLE ACCESS FULL	GRADES_ENUM					
10	TABLE ACCESS FULL	GRADES					
* 11	TABLE ACCESS FULL	POINTS	1494K	25M		8881 (1)	00:00:01

Predicate Information (identified by operation id):

```

2 - access("POINTS"."STUDENT_ID"="STUDENT_ID")
4 - filter(COUNT("Q"."SUBJECT_ID")>=4)
6 - filter(SUM("GRADES_ENUM"."VALUE")/COUNT("GRADES_ENUM"."VALUE")<4)
8 - access("GRADES"."VALUE"="GRADES_ENUM"."SYMBOL")
10 - filter("GRADES"."AWARD_DATE">>=TO_DATE(' 2023-09-01 00:00:00', 'yyyy-mm-dd
    hh24:mi:ss') AND "GRADES"."AWARD_DATE"<=TO_DATE(' 2024-06-30 00:00:00', 'yyyy-mm-dd
    hh24:mi:ss'))
11 - filter("POINTS"."VALUE">0)

```

Candidates for improvements

- Hash Joins (2, 8)**: These operations join multiple tables, which is quite expensive

- **Hash Group By (7)**: This operation groups all students for every subject they attend to calculate average for each of them.
- **Table Access Full (11)**: This operation accesses table points which is the biggest table in our database.

7 Testing application

7.1 Tester script

The script for this task was written in SQL. It runs the whole workload (given as a parameter) multiple times and saves observed times in a `workload_metrics` table. Another procedure saves query plans in a separate `workload_plans` table.

```

1  create or replace PROCEDURE WORKLOAD_TESTER
2  (
3      P_WORKLOAD IN SYS_REFCURSOR
4  ) AS
5      V_NAME VARCHAR2(32);
6      V_SQL VARCHAR2(4000);
7      V_PLAN VARCHAR2(32767);
8      V_PLAN_ID NUMBER;
9      V_START TIMESTAMP;
10     V_END TIMESTAMP;
11     V_ELAPSED NUMBER;
12 BEGIN
13     LOOP
14         FETCH P_WORKLOAD INTO V_NAME, V_SQL;
15         EXIT WHEN P_WORKLOAD%NOTFOUND;
16
17         SELECT MAX(workload_plans.id) INTO V_PLAN_ID
18         FROM workload_plans
19             INNER JOIN workload_sql ON workload_plans.sql_id = workload_sql.id
20             WHERE workload_sql.name = V_NAME;
21         EXECUTE IMMEDIATE 'ALTER SYSTEM FLUSH BUFFER_CACHE';
22         EXECUTE IMMEDIATE 'ALTER SYSTEM FLUSH SHARED_POOL';
23         V_START := SYSTIMESTAMP;
24         EXECUTE IMMEDIATE V_SQL;
25         V_END := SYSTIMESTAMP;
26         ROLLBACK;
27         V_ELAPSED := EXTRACT(SECOND FROM (V_END - V_START)) * 1000;
28         INSERT INTO workload_metrics (plan_id, execution_time) VALUES (v_plan_id, v_elapsed);
29         COMMIT;
30     END LOOP;
31 END WORKLOAD_TESTER;

```

7.2 Running tests

the workload is stored in `workload_sql` table and passed to tester script as a cursor. The tests are being run 10 times in the following way:

```

1  DECLARE
2      workload SYS_REFCURSOR;
3
4  BEGIN
5      DELETE FROM workload_metrics;
6      DELETE FROM workload_plans;
7
8      OPEN workload FOR
9          SELECT name, code FROM WORKLOAD_SQL;
10     generate_plans(workload);
11     CLOSE workload;
12
13     FOR i IN 1..10 LOOP
14         OPEN workload FOR
15             SELECT name, code FROM WORKLOAD_SQL;
16         workload_tester(workload);
17         CLOSE workload;
18     END LOOP;
19
20 END;

```

7.3 Tests results

Transaction Name	Tests Amount	Min Time	Max Time	Avg Time	Plan cost
Grades analysis	10	703	754	724	144 K
House points summary	10	754	790	771	159 K
Best Students Display	10	271	307	288	1254
Raising Grades (UPDATE)	10	2384	2707	2500	2240
Assign subjects (INSERT)	10	3830	4123	3939	12 318
Removing points (DELETE)	10	5159	6616	5422	16 247

7.4 Conclusions

Transactions querying the data in general are much faster than the transactions changing the data. That's because the second type of transactions use selects within them and also because UPDATE, INSERT and DELETE commands take more time.

8 Indexes

8.1 Index on Points(award_date)

Purpose: Optimization of transaction House Points Summary.

Query plan points to be changed: 13, 21, 34, 37

Id	Operation	Name	Rows	Bytes	TempSpc	Cost	(%CPU)	Time
0	SELECT STATEMENT		2986K	535M		159K	(1)	00:00:07
1	HASH GROUP BY		2986K	535M	583M	159K	(1)	00:00:07
* 2	HASH JOIN		2986K	535M		36394	(3)	00:00:02
3	TABLE ACCESS FULL	TEACHERS	1097	20843		5	(0)	00:00:01
* 4	HASH JOIN		2986K	481M		36369	(3)	00:00:02
5	TABLE ACCESS FULL	TEACHERS	1097	20843		5	(0)	00:00:01
* 6	HASH JOIN		2986K	427M		36344	(3)	00:00:02
7	VIEW		4	64		9127	(4)	00:00:01
8	SORT GROUP BY		4	116		9127	(4)	00:00:01
9	VIEW		858	24882		9127	(4)	00:00:01
10	HASH GROUP BY		858	19734		9127	(4)	00:00:01
* 11	HASH JOIN		2986K	65M		8934	(2)	00:00:01
12	TABLE ACCESS FULL	STUDENTS	9346	65422		19	(0)	00:00:01
* 13	TABLE ACCESS FULL	POINTS	2986K	45M		8895	(1)	00:00:01
* 14	HASH JOIN		2986K	381M		27197	(3)	00:00:02
15	VIEW		4	116		9131	(4)	00:00:01
16	SORT GROUP BY		4	116		9131	(4)	00:00:01
17	VIEW		3103	89987		9131	(4)	00:00:01
18	HASH GROUP BY		3103	83781		9131	(4)	00:00:01
* 19	HASH JOIN		2986K	76M		8938	(2)	00:00:01
20	TABLE ACCESS FULL	STUDENTS	9346	65422		19	(0)	00:00:01
* 21	TABLE ACCESS FULL	POINTS	2986K	56M		8899	(1)	00:00:01
* 22	HASH JOIN		2986K	299M		18046	(3)	00:00:01
23	VIEW	VW_GBF_15	9346	812K		9131	(4)	00:00:01
* 24	HASH JOIN		9346	930K		9131	(4)	00:00:01
25	MERGE JOIN		4	380		9112	(4)	00:00:01
26	TABLE ACCESS BY INDEX ROWID	HOUSES	4	56		2	(0)	00:00:01
27	INDEX FULL SCAN	SYS_C008743	4			1	(0)	00:00:01
* 28	SORT JOIN		4	324		9110	(4)	00:00:01
29	VIEW		4	324		9109	(4)	00:00:01
30	SORT GROUP BY		4	156		9109	(4)	00:00:01
* 31	HASH JOIN		9346	355K		9108	(4)	00:00:01
32	VIEW		9346	155K		9088	(4)	00:00:01
33	HASH GROUP BY		9346	146K		9088	(4)	00:00:01
* 34	TABLE ACCESS FULL	POINTS	2986K	45M		8895	(1)	00:00:01
35	TABLE ACCESS FULL	STUDENTS	9346	200K		19	(0)	00:00:01
36	TABLE ACCESS FULL	STUDENTS	9346	65422		19	(0)	00:00:01
* 37	TABLE ACCESS FULL	POINTS	2986K	45M		8895	(1)	00:00:01

Predicate Information (identified by operation id):

```

2 - access("T2"."ID"="T"."LEAST_FAVOURITE_TEACHER")
4 - access("T1"."ID"="T"."FAVOURITE_TEACHER")
6 - access("M"."HOUSE_ID"="ITEM_2")
11 - access("POINTS"."STUDENT_ID"="STUDENTS"."ID")

```

```

13 - filter("POINTS"."AWARD_DATE">>=TO_DATE(' 2023-09-01 00:00:00', 'yyyy-mm-dd hh24:mi:ss') AND
      "POINTS"."AWARD_DATE"<=TO_DATE(' 2024-06-30 00:00:00', 'yyyy-mm-dd hh24:mi:ss'))
14 - access("T"."HOUSE_ID"="ITEM_2")
19 - access("STUDENTS"."ID"="POINTS"."STUDENT_ID")
21 - filter("POINTS"."AWARD_DATE">>=TO_DATE(' 2023-09-01 00:00:00', 'yyyy-mm-dd hh24:mi:ss') AND
      "POINTS"."AWARD_DATE"<=TO_DATE(' 2024-06-30 00:00:00', 'yyyy-mm-dd hh24:mi:ss'))
22 - access("ITEM_1"="POINTS"."STUDENT_ID")
24 - access("HOUSES"."ID"="STUDENTS"."HOUSE_ID")
28 - access("S"."HOUSE_ID"="HOUSES"."ID")
      filter("S"."HOUSE_ID"="HOUSES"."ID")
31 - access("STUDENTS"."ID"="Q"."STUDENT_ID")
34 - filter("POINTS"."AWARD_DATE">>=TO_DATE(' 2023-09-01 00:00:00', 'yyyy-mm-dd hh24:mi:ss') AND
      "POINTS"."AWARD_DATE"<=TO_DATE(' 2024-06-30 00:00:00', 'yyyy-mm-dd hh24:mi:ss'))
37 - filter("POINTS"."AWARD_DATE">>=TO_DATE(' 2023-09-01 00:00:00', 'yyyy-mm-dd hh24:mi:ss') AND
      "POINTS"."AWARD_DATE"<=TO_DATE(' 2024-06-30 00:00:00', 'yyyy-mm-dd hh24:mi:ss'))

```

8.1.1 Conclusion

Dates are often selected as range and can have many values, that's why it's best to use **B-tree** index type in this case.

8.2 Index on Grades(student_id, award_date)

Purpose: Optimization of transaction Grades Analysis.

Query plan points to be changed: 17, 29, 36

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT		84M	17G		6392K (1)	00:04:10
1	SORT ORDER BY		84M	17G	18G	6392K (1)	00:04:10
* 2	HASH JOIN		84M	17G		2384K (1)	00:01:34
3	TABLE ACCESS FULL	SUBJECTS	76	1520		3 (0)	00:00:01
* 4	HASH JOIN		84M	15G		2383K (1)	00:01:34
5	MERGE JOIN CARTESIAN		2736	82080		64 (0)	00:00:01
6	MERGE JOIN CARTESIAN		36	360		13 (0)	00:00:01
7	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01
8	BUFFER SORT		6	30		10 (0)	00:00:01
9	TABLE ACCESS FULL	GRADES_ENUM	6	30		2 (0)	00:00:01
10	BUFFER SORT		76	1520		62 (0)	00:00:01
11	TABLE ACCESS FULL	SUBJECTS	76	1520		1 (0)	00:00:01
12	VIEW		84M	13G		2382K (1)	00:01:34
13	SORT GROUP BY		84M	9G	10G	2382K (1)	00:01:34
* 14	HASH JOIN		84M	9G	30M	26186 (4)	00:00:02
15	TABLE ACCESS FULL	GRADES	1977K	7723K		2492 (2)	00:00:01
* 16	HASH JOIN		397K	45M		19118 (2)	00:00:01
* 17	HASH JOIN		1256	128K		10204 (3)	00:00:01
18	JOIN FILTER CREATE	:BF0000	23	1426		2548 (3)	00:00:01
* 19	HASH JOIN OUTER		23	1426		2548 (3)	00:00:01
* 20	HASH JOIN ANTI		23	1035		2545 (3)	00:00:01
* 21	HASH JOIN		2337	95817		22 (0)	00:00:01
* 22	TABLE ACCESS FULL	HOUSES	1	14		3 (0)	00:00:01
* 23	TABLE ACCESS FULL	STUDENTS	9346	246K		19 (0)	00:00:01
24	VIEW	VW_NSO_1	190K	745K		2521 (3)	00:00:01
* 25	HASH JOIN		190K	3541K		2521 (3)	00:00:01
* 26	TABLE ACCESS FULL	GRADES_ENUM	1	5		3 (0)	00:00:01
* 27	TABLE ACCESS FULL	GRADES_ENUM	1	5		3 (0)	00:00:01
* 28	TABLE ACCESS FULL	GRADES_ENUM	1	5		3 (0)	00:00:01

* 29	TABLE ACCESS FULL	GRADES	1957K 26M	2505 (2) 00:00:01
30	TABLE ACCESS FULL	QUIDDITCH_TEAM_MEMBERS	97 1649	3 (0) 00:00:01
31	VIEW		502K 20M	7653 (3) 00:00:01
32	HASH GROUP BY		502K 6866K	45M 7653 (3) 00:00:01
* 33	HASH JOIN		1977K 26M	2511 (2) 00:00:01
34	TABLE ACCESS FULL	GRADES_ENUM	6 30	3 (0) 00:00:01
35	JOIN FILTER USE	:BF0000	1977K 16M	2495 (2) 00:00:01
* 36	TABLE ACCESS FULL	GRADES	1977K 16M	2495 (2) 00:00:01
* 37	TABLE ACCESS FULL	POINTS	2956K 45M	8895 (1) 00:00:01

Predicate Information (identified by operation id):

```

2 - access("Q"."WORST SUBJECT"="S2"."ID")
4 - access("Q"."BEST SUBJECT"="S1"."ID" AND "Q"."MIN GRADE"="G2"."VALUE" AND
          "Q"."MAX GRADE"="G1"."VALUE")
14 - access("STUDENTS"."ID"="GRADES"."STUDENT_ID")
16 - access("STUDENTS"."ID"="POINTS"."STUDENT_ID")
17 - access("AVERAGES"."STUDENT_ID"="STUDENTS"."ID")
19 - access("STUDENTS"."ID"="QUIDDITCH_TEAM_MEMBERS"."STUDENT_ID"(+))
20 - access("STUDENTS"."ID"="ID")
21 - access("STUDENTS"."HOUSE_ID"="HOUSES"."ID")
22 - filter("HOUSES"."NAME"='Gryffindor')
23 - filter("STUDENTS"."YEAR">>=1 AND "STUDENTS"."YEAR"><=7)
25 - access("GRADES"."VALUE"="GRADES_ENUM"."SYMBOL")
26 - filter("GRADES_ENUM"."VALUE">< (SELECT "GRADES_ENUM"."VALUE" FROM "GRADES_ENUM" "GRADES_ENUM" WHERE
          "GRADES_ENUM"."SYMBOL"='T') OR "GRADES_ENUM"."VALUE">> (SELECT "GRADES_ENUM"."VALUE" FROM "GRADES_ENUM"
          "GRADES_ENUM" WHERE "GRADES_ENUM"."SYMBOL"='O'))
27 - filter("GRADES_ENUM"."SYMBOL"='T')
28 - filter("GRADES_ENUM"."SYMBOL"='O')
29 - filter("GRADES"."AWARD_DATE"=<TO_DATE(' 2024-06-25 00:00:00 ', 'syyyy-mm-dd hh24:mi:ss') AND
          "GRADES"."AWARD_DATE">>=TO_DATE(' 2023-09-02 00:00:00 ', 'syyyy-mm-dd hh24:mi:ss'))
33 - access("GRADES"."VALUE"="GRADES_ENUM"."SYMBOL")
36 - filter(SYS_OP_BLOOM_FILTER(:BF0000,"GRADES"."STUDENT_ID"))
37 - filter("POINTS"."AWARD_DATE"=<TO_DATE(' 2024-06-25 00:00:00 ', 'syyyy-mm-dd hh24:mi:ss') AND
          "POINTS"."AWARD_DATE">>=TO_DATE(' 2023-09-02 00:00:00 ', 'syyyy-mm-dd hh24:mi:ss'))
```

Students' IDs and award dates appear multiple times in our rows. However, there are still many of them, that's why it's not obvious if bitmap indexes would be better than b-tree indexes. For that reason, we've compared query plans for both of these types.

8.2.1 Experiment 1 - B-tree index type

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT		12M	2592M		1662K (1)	00:01:05
1	SORT ORDER BY		12M	2592M	2757M	1662K (1)	00:01:05
* 2	HASH JOIN		12M	2592M		1073K (1)	00:00:42
3	TABLE ACCESS FULL	SUBJECTS	76	1520		3 (0)	00:00:01
* 4	HASH JOIN		12M	2356M		1073K (1)	00:00:42
5	MERGE JOIN CARTESIAN		2736	82080		64 (0)	00:00:01
6	MERGE JOIN CARTESIAN		36	360		13 (0)	00:00:01
7	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01
8	BUFFER SORT		6	30		10 (0)	00:00:01
9	TABLE ACCESS FULL	GRADES_ENUM	6	30		2 (0)	00:00:01
10	BUFFER SORT		76	1520		62 (0)	00:00:01
11	TABLE ACCESS FULL	SUBJECTS	76	1520		1 (0)	00:00:01
12	VIEW		12M	2003M		1073K (1)	00:00:42
13	SORT GROUP BY		12M	1425M	10G	1073K (1)	00:00:42

* 14	HASH JOIN		84M	9698M	30M	44918	(2)	00:00:02
15	INDEX FAST FULL SCAN	TESTINDEX1	1977K	7723K	1736	(2)	00:00:01	
* 16	HASH JOIN		397K	44M	38682	(1)	00:00:02	
17	NESTED LOOPS		1256	123K	29768	(1)	00:00:02	
* 18	HASH JOIN OUTER		23	1380	28087	(1)	00:00:02	
19	NESTED LOOPS ANTI		23	989	28084	(1)	00:00:02	
* 20	HASH JOIN		2337	95817	22	(0)	00:00:01	
* 21	TABLE ACCESS FULL	HOUSES	1	14	3	(0)	00:00:01	
* 22	TABLE ACCESS FULL	STUDENTS	9346	246K	19	(0)	00:00:01	
23	VIEW PUSHED PREDICATE	VW_NSO_1	2	4	12	(0)	00:00:01	
24	NESTED LOOPS		2	38	12	(0)	00:00:01	
25	NESTED LOOPS		21	38	12	(0)	00:00:01	
* 26	TABLE ACCESS FULL	GRADES_ENUM	1	5	3	(0)	00:00:01	
* 27	TABLE ACCESS FULL	GRADES_ENUM	1	5	3	(0)	00:00:01	
* 28	TABLE ACCESS FULL	GRADES_ENUM	1	5	3	(0)	00:00:01	
* 29	INDEX RANGE SCAN	TESTINDEX1	21		2	(0)	00:00:01	
* 30	TABLE ACCESS BY INDEX ROWID	GRADES	3	42	9	(0)	00:00:01	
31	TABLE ACCESS FULL	QUIDDITCH_TEAM_MEMBERS	97	1649	3	(0)	00:00:01	
32	VIEW PUSHED PREDICATE		1	41	73	(2)	00:00:01	
33	SORT GROUP BY		51	714	73	(2)	00:00:01	
* 34	HASH JOIN		212	2968	72	(0)	00:00:01	
35	TABLE ACCESS FULL	GRADES_ENUM	6	30	3	(0)	00:00:01	
36	TABLE ACCESS BY INDEX ROWID BATCHED	GRADES	212	1908	69	(0)	00:00:01	
* 37	INDEX RANGE SCAN	TESTINDEX1	212		3	(0)	00:00:01	
* 38	TABLE ACCESS FULL	POINTS	2956K	45M	8895	(1)	00:00:01	

Predicate Information (identified by operation id):

```

2 - access("Q"."WORST SUBJECT"="S2"."ID")
4 - access("Q"."BEST SUBJECT"="S1"."ID" AND "Q"."MIN GRADE"="G2"."VALUE" AND "Q"."MAX GRADE"="G1"."VALUE")
14 - access("STUDENTS"."ID"="GRADES"."STUDENT_ID")
16 - access("STUDENTS"."ID"="POINTS"."STUDENT_ID")
18 - access("STUDENTS"."ID"="QUIDDITCH_TEAM_MEMBERS"."STUDENT_ID"(+))
20 - access("STUDENTS"."HOUSE_ID"="HOUSES"."ID")
21 - filter("HOUSES"."NAME"='Gryffindor')
22 - filter("STUDENTS"."YEAR">>=1 AND "STUDENTS"."YEAR" <=7)
26 - filter("GRADES_ENUM"."VALUE" < (SELECT "GRADES_ENUM"."VALUE" FROM "GRADES_ENUM" "GRADES_ENUM" WHERE
    "GRADES_ENUM"."SYMBOL"='T') OR "GRADES_ENUM"."VALUE" > (SELECT "GRADES_ENUM"."VALUE" FROM "GRADES_ENUM" "GRADES_ENUM"
    WHERE "GRADES_ENUM"."SYMBOL"='O'))
27 - filter("GRADES_ENUM"."SYMBOL"='T')
28 - filter("GRADES_ENUM"."SYMBOL"='O')
29 - access("GRADES"."STUDENT_ID"="STUDENTS"."ID" AND "GRADES"."AWARD_DATE">>=TO_DATE(' 2023-09-02 00:00:00',
    'syyyy-mm-dd hh24:mi:ss') AND "GRADES"."AWARD_DATE" <=TO_DATE(' 2024-06-25 00:00:00', 'syyyy-mm-dd hh24:mi:ss'))
30 - filter("GRADES"."VALUE"="GRADES_ENUM"."SYMBOL")
34 - access("GRADES"."VALUE"="GRADES_ENUM"."SYMBOL")
37 - access("GRADES"."STUDENT_ID"="STUDENTS"."ID")
38 - filter("POINTS"."AWARD_DATE" <=TO_DATE(' 2024-06-25 00:00:00', 'syyyy-mm-dd hh24:mi:ss') AND
    "POINTS"."AWARD_DATE">>=TO_DATE(' 2023-09-02 00:00:00', 'syyyy-mm-dd hh24:mi:ss'))

```

8.2.2 Experiment 2 - Bitmap index type

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT		12M	2592M		1658K	(1) 00:01:05
1	SORT ORDER BY		12M	2592M	2757M	1658K	(1) 00:01:05
* 2	HASH JOIN		12M	2592M		1069K	(1) 00:00:42
3	TABLE ACCESS FULL	SUBJECTS	76	1520		3	(0) 00:00:01
* 4	HASH JOIN		12M	2356M		1069K	(1) 00:00:42
5	MERGE JOIN CARTESIAN		2736	82080		64	(0) 00:00:01
6	MERGE JOIN CARTESIAN		36	360		13	(0) 00:00:01
7	TABLE ACCESS FULL	GRADES_ENUM	6	30		3	(0) 00:00:01

8	BUFFER SORT		6	30	10	(0)	00:00:01
9	TABLE ACCESS FULL	GRADES_ENUM	6	30	2	(0)	00:00:01
10	BUFFER SORT		76	1520	62	(0)	00:00:01
11	TABLE ACCESS FULL	SUBJECTS	76	1520	1	(0)	00:00:01
12	VIEW		12M	2003M	1069K	(1)	00:00:42
13	SORT GROUP BY		12M	1425M	10G	1069K	(1) 00:00:42
* 14	HASH JOIN		84M	9698M	30M	41009	(3) 00:00:02
15	TABLE ACCESS FULL	GRADES	1977K	7723K	2492	(2)	00:00:01
* 16	HASH JOIN		397K	44M	34016	(1)	00:00:02
17	NESTED LOOPS		1256	123K	25102	(1)	00:00:01
* 18	HASH JOIN OUTER		23	1380	23946	(1)	00:00:01
19	NESTED LOOPS ANTI		23	989	23943	(1)	00:00:01
* 20	HASH JOIN		2337	95817	22	(0)	00:00:01
* 21	TABLE ACCESS FULL	HOUSES	1	14	3	(0)	00:00:01
* 22	TABLE ACCESS FULL	STUDENTS	9346	246K	19	(0)	00:00:01
23	VIEW PUSHED PREDICATE	VW_NSO_1	2	4	10	(0)	00:00:01
* 24	HASH JOIN		2	38	10	(0)	00:00:01
* 25	TABLE ACCESS FULL	GRADES_ENUM	1	5	3	(0)	00:00:01
* 26	TABLE ACCESS FULL	GRADES_ENUM	1	5	3	(0)	00:00:01
* 27	TABLE ACCESS FULL	GRADES_ENUM	1	5	3	(0)	00:00:01
28	TABLE ACCESS BY INDEX ROWID BATCHED	GRADES	21	294	7	(0)	00:00:01
29	BITMAP CONVERSION TO ROWIDS						
* 30	BITMAP INDEX RANGE SCAN	TESTINDEX1					
31	TABLE ACCESS FULL	QUIDDITCH_TEAM_MEMBERS	97	1649	3	(0)	00:00:01
32	VIEW PUSHED PREDICATE		1	41	50	(2)	00:00:01
33	SORT GROUP BY		51	714	50	(2)	00:00:01
* 34	HASH JOIN		212	2968	49	(0)	00:00:01
35	TABLE ACCESS FULL	GRADES_ENUM	6	30	3	(0)	00:00:01
36	TABLE ACCESS BY INDEX ROWID BATCHED	GRADES	212	1908	46	(0)	00:00:01
37	BITMAP CONVERSION TO ROWIDS						
* 38	BITMAP INDEX RANGE SCAN	TESTINDEX1					
* 39	TABLE ACCESS FULL	POINTS	2956K	45M	8895	(1)	00:00:01

Predicate Information (identified by operation id):

```

2 - access("Q"."WORST SUBJECT"="S2"."ID")
4 - access("Q"."BEST SUBJECT"="S1"."ID" AND "Q"."MIN_GRADE"="G2"."VALUE" AND "Q"."MAX_GRADE"="G1"."VALUE")
14 - access("STUDENTS"."ID"="GRADES"."STUDENT_ID")
16 - access("STUDENTS"."ID"="POINTS"."STUDENT_ID")
18 - access("STUDENTS"."ID"="QUIDDITCH_TEAM_MEMBERS"."STUDENT_ID"())
20 - access("STUDENTS"."HOUSE_ID"="HOUSES"."ID")
21 - filter("HOUSES"."NAME"='Gryffindor')
22 - filter("STUDENTS"."YEAR">>=1 AND "STUDENTS"."YEAR"><=7)
24 - access("GRADES"."VALUE"="GRADES_ENUM"."SYMBOL")
25 - filter("GRADES_ENUM"."VALUE">< (SELECT "GRADES_ENUM"."VALUE" FROM "GRADES_ENUM" "GRADES_ENUM" WHERE
        "GRADES_ENUM"."SYMBOL"='T') OR "GRADES_ENUM"."VALUE">> (SELECT "GRADES_ENUM"."VALUE" FROM "GRADES_ENUM" "GRADES_ENUM"
        WHERE "GRADES_ENUM"."SYMBOL"='O'))
26 - filter("GRADES_ENUM"."SYMBOL"='T')
27 - filter("GRADES_ENUM"."SYMBOL"='O')
30 - access("GRADES"."STUDENT_ID"="STUDENTS"."ID" AND "GRADES"."AWARD_DATE">>=TO_DATE(' 2023-09-02 00:00:00',
        'syyyy-mm-dd hh24:mi:ss') AND "GRADES"."AWARD_DATE" <=TO_DATE(' 2024-06-25 00:00:00', 'syyyy-mm-dd hh24:mi:ss'))
    filter("GRADES"."STUDENT_ID"="STUDENTS"."ID" AND "GRADES"."AWARD_DATE">>=TO_DATE(' 2023-09-02 00:00:00',
        'syyyy-mm-dd hh24:mi:ss') AND "GRADES"."AWARD_DATE" <=TO_DATE(' 2024-06-25 00:00:00', 'syyyy-mm-dd hh24:mi:ss'))
34 - access("GRADES"."VALUE"="GRADES_ENUM"."SYMBOL")
38 - access("GRADES"."STUDENT_ID"="STUDENTS"."ID")

```

8.2.3 Conclusion

Both indexes greatly improve the efficiency of the sample query. In this particular case the **bitmap** index is slightly better than the b-tree.

8.3 Index on Points(EXTRACT(MONTH FROM award_date))

Purpose: Optimization of transaction House Points Summary.

Query plan points to be changed: 37

Plan hash value: 1621774207

Id	Operation	Name	Rows	Bytes	TempSpc	Cost	(%CPU)	Time
0	SELECT STATEMENT		2986K	535M		159K	(1)	00:00:07
1	HASH GROUP BY		2986K	535M	583M	159K	(1)	00:00:07
* 2	HASH JOIN		2986K	535M		36394	(3)	00:00:02
3	TABLE ACCESS FULL	TEACHERS	1097	20843		5	(0)	00:00:01
* 4	HASH JOIN		2986K	481M		36369	(3)	00:00:02
5	TABLE ACCESS FULL	TEACHERS	1097	20843		5	(0)	00:00:01
* 6	HASH JOIN		2986K	427M		36344	(3)	00:00:02
7	VIEW		4	64		9127	(4)	00:00:01
8	SORT GROUP BY		4	116		9127	(4)	00:00:01
9	VIEW		4	116		9127	(4)	00:00:01
10	HASH GROUP BY		4	144		9127	(4)	00:00:01
* 11	HASH JOIN		2986K	102M		8934	(2)	00:00:01
12	TABLE ACCESS FULL	STUDENTS	9346	65422		19	(0)	00:00:01
* 13	TABLE ACCESS FULL	POINTS	2986K	82M		8895	(1)	00:00:01
* 14	HASH JOIN		2986K	381M		27197	(3)	00:00:02
15	VIEW		4	116		9131	(4)	00:00:01
16	SORT GROUP BY		4	116		9131	(4)	00:00:01
17	VIEW		3103	89987		9131	(4)	00:00:01
18	HASH GROUP BY		3103	83781		9131	(4)	00:00:01
* 19	HASH JOIN		2986K	76M		8938	(2)	00:00:01
20	TABLE ACCESS FULL	STUDENTS	9346	65422		19	(0)	00:00:01
* 21	TABLE ACCESS FULL	POINTS	2986K	56M		8899	(1)	00:00:01
* 22	HASH JOIN		2986K	299M		18046	(3)	00:00:01
23	VIEW	VW_GBF_15	9346	812K		9131	(4)	00:00:01
* 24	HASH JOIN		9346	930K		9131	(4)	00:00:01
25	MERGE JOIN		4	380		9112	(4)	00:00:01
26	TABLE ACCESS BY INDEX ROWID	HOUSES	4	56		2	(0)	00:00:01
27	INDEX FULL SCAN	SYS_C008743	4			1	(0)	00:00:01
* 28	SORT JOIN		4	324		9110	(4)	00:00:01
29	VIEW		4	324		9109	(4)	00:00:01
30	SORT GROUP BY		4	156		9109	(4)	00:00:01
* 31	HASH JOIN		9346	355K		9108	(4)	00:00:01
32	VIEW		9346	155K		9088	(4)	00:00:01
33	HASH GROUP BY		9346	146K		9088	(4)	00:00:01
* 34	TABLE ACCESS FULL	POINTS	2986K	45M		8895	(1)	00:00:01
35	TABLE ACCESS FULL	STUDENTS	9346	200K		19	(0)	00:00:01
36	TABLE ACCESS FULL	STUDENTS	9346	65422		19	(0)	00:00:01
* 37	TABLE ACCESS FULL	POINTS	2986K	45M		8895	(1)	00:00:01

Predicate Information (identified by operation id):

```

2 - access("T2"."ID"="T"."LEAST_FAVOURITE_TEACHER")
4 - access("T1"."ID"="T"."FAVOURITE_TEACHER")
6 - access("M"."HOUSE_ID"="ITEM_2")
11 - access("POINTS"."STUDENT_ID"="STUDENTS"."ID")
13 - filter("POINTS"."AWARD_DATE">>=TO_DATE(' 2023-09-02 00:00:00', 'syyyy-mm-dd hh24:mi:ss') AND
      "POINTS"."AWARD_DATE"=<TO_DATE(' 2024-06-28 00:00:00', 'syyyy-mm-dd hh24:mi:ss'))
14 - access("T"."HOUSE_ID"="ITEM_2")
19 - access("STUDENTS"."ID"="POINTS"."STUDENT_ID")
21 - filter("POINTS"."AWARD_DATE">>=TO_DATE(' 2023-09-02 00:00:00', 'syyyy-mm-dd hh24:mi:ss') AND
      "POINTS"."AWARD_DATE"=<TO_DATE(' 2024-06-28 00:00:00', 'syyyy-mm-dd hh24:mi:ss'))
```

```

22 - access("ITEM_1"="POINTS"."STUDENT_ID")
24 - access("HOUSES"."ID"="STUDENTS"."HOUSE_ID")
28 - access("S"."HOUSE_ID"="HOUSES"."ID")
      filter("S"."HOUSE_ID"="HOUSES"."ID")
31 - access("STUDENTS"."ID"="Q"."STUDENT_ID")
34 - filter("POINTS"."AWARD_DATE">>=TO_DATE(' 2023-09-02 00:00:00', 'yyyy-mm-dd hh24:mi:ss') AND
      "POINTS"."AWARD_DATE" <=TO_DATE(' 2024-06-28 00:00:00', 'yyyy-mm-dd hh24:mi:ss'))
37 - filter("POINTS"."AWARD_DATE">>=TO_DATE(' 2023-09-02 00:00:00', 'yyyy-mm-dd hh24:mi:ss') AND
      "POINTS"."AWARD_DATE" <=TO_DATE(' 2024-06-28 00:00:00', 'yyyy-mm-dd hh24:mi:ss'))

```

8.3.1 Conclusion

One of the often used functions in filtering is month extraction. As there is small amount of months, the type of this **functional** index should be **bitmap**.

9 Indexes - implementation

9.1 Implemented indexes

Indexes that have been selected after experiments, and have been tested during cost comparison had been created as follows:

```

CREATE BITMAP INDEX index1 ON grades(student_id, award_date);
CREATE INDEX index2 ON points(student_id, award_date);
CREATE INDEX index3 ON grades(EXTRACT(MONTH FROM award_date));

```

9.2 Cost comparison

We ran the whole workload before and after creating indexes, **10 times** in each case. Creating them positively influenced 4 out of 6 query plans in our workload. What is very interesting is that for 3 of them the actual running times didn't actually improve - actually quite the contrary; now they are a lot bigger. The changes of times of procedures, of which query plans didn't change, are a result of natural differences in each running time.

Procedure name	Min. measured time				Max. measured time				Average measured time				Query plan cost			
	Default	Index	Diff.	Profit	Default	Index	Diff.	Profit	Default	Index	Diff.	Profit	Default	Index	Diff.	Profit
Grades analysis	552	1078	-526	-95%	831	1560	-729	-88%	628.4	1335.9	-707.5	-113%	6392 K	1658 K	4734 K	74%
Points summary	392	757	-365	-93%	567	1098	-531	-94%	452.6	935.4	-482.8	-107%	39 M	38 M	1 M	3%
Best students	232	244	-12	-5%	303	331	-28	-9%	262.7	278.6	-15.9	-6%	22 M	19 M	3 M	14%
Raise grades	1925	1883	42	2%	2809	2712	97	3%	2168.4	2303.2	-134.8	-6%	21196	21196	0	0%
Assign subjects	2757	2835	-78	-3%	3934	4081	-147	-4%	3061.4	3459.5	-398.1	-13%	12318	12318	0	0%
Remove points	26407	3279	23128	88%	34829	31564	3265	9%	29593.6	16934.5	12659.1	43%	11394	6771	4623	41%
Whole workload	32265	10076	22189	69%	43273	41346	1927	4%	36167.1	24247.1	11920.6	33%	67437 K	58698 K	8739 K	13%

9.3 Query plans comparison

Created indexes have affected four of all six procedures' query plans, including Grade Analysis, Points Summary, Best Students Display, and Removing Points. Their changed query plans are presented below.



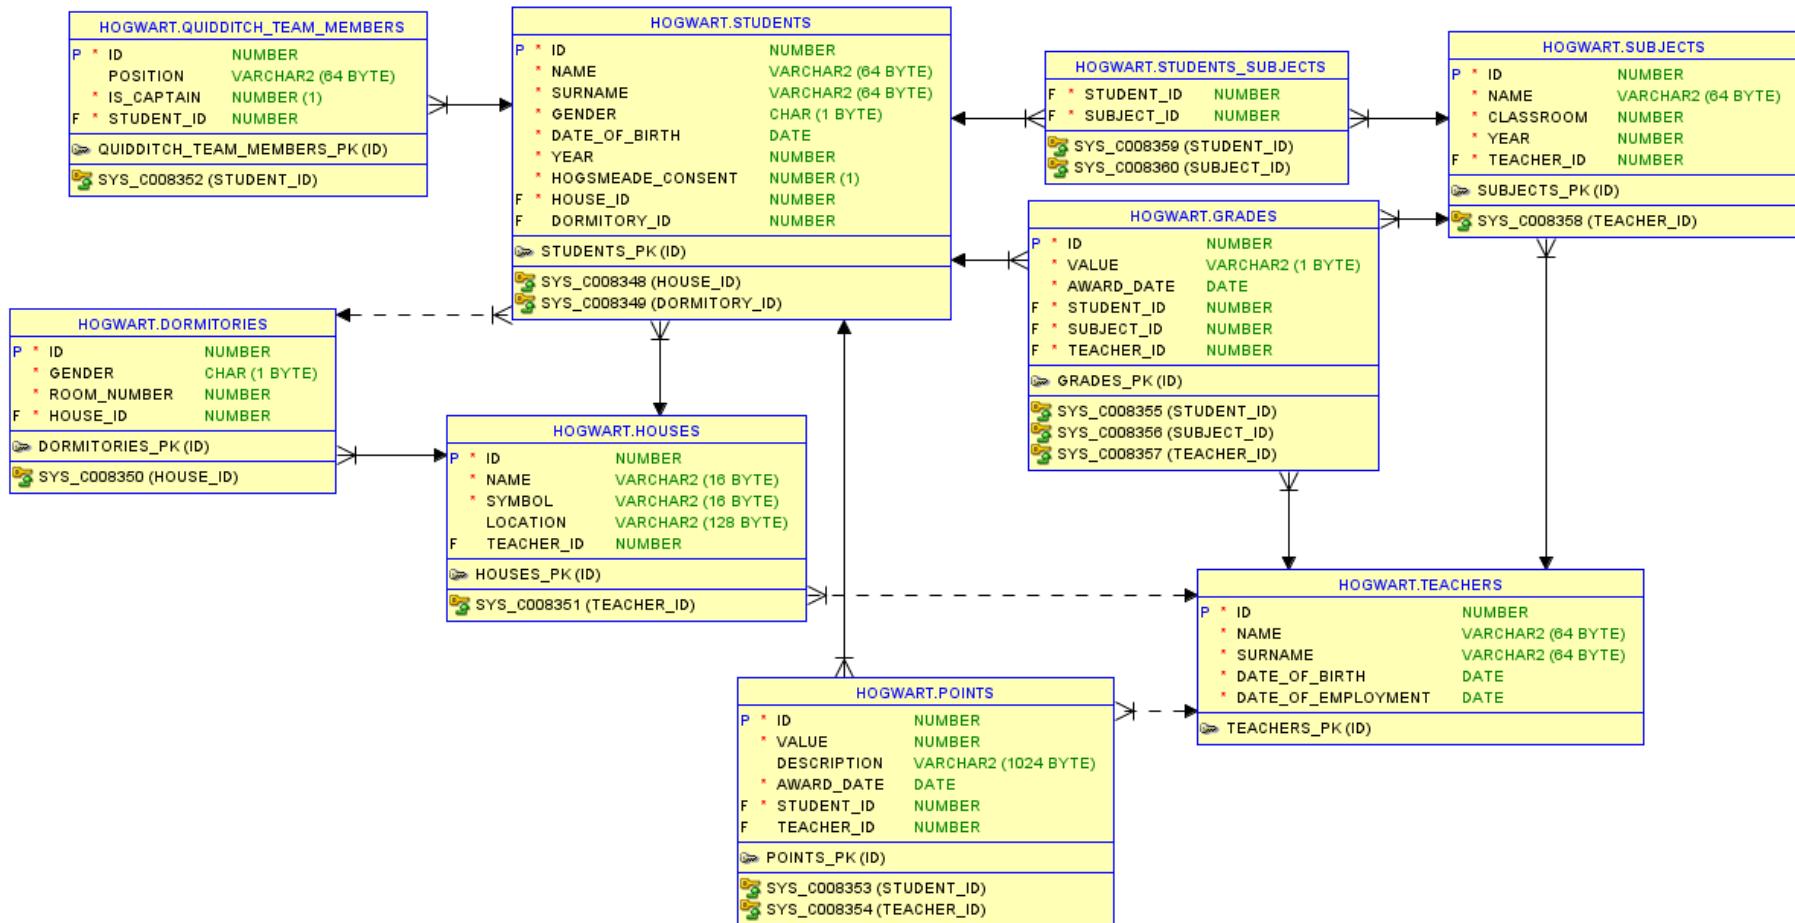


Figure 2: Database schema generated with Oracle SQL Developer Software

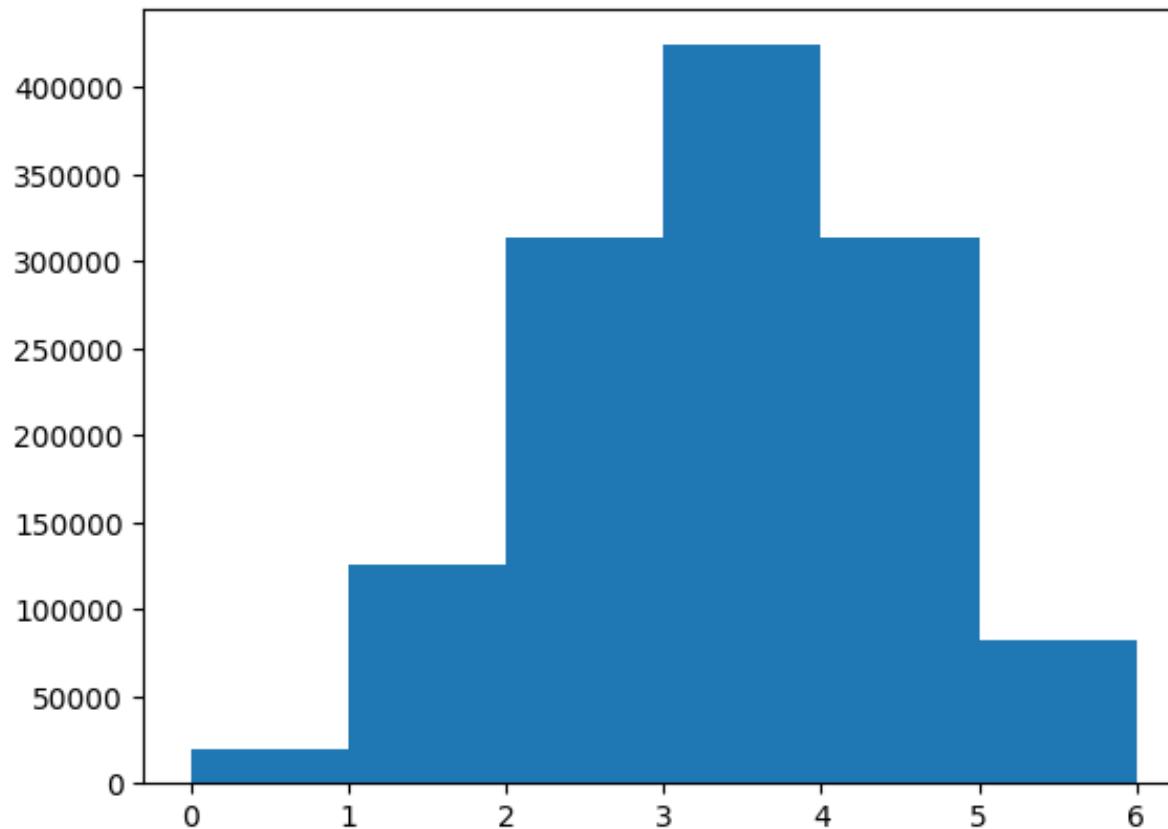


Figure 3: Grades distribution

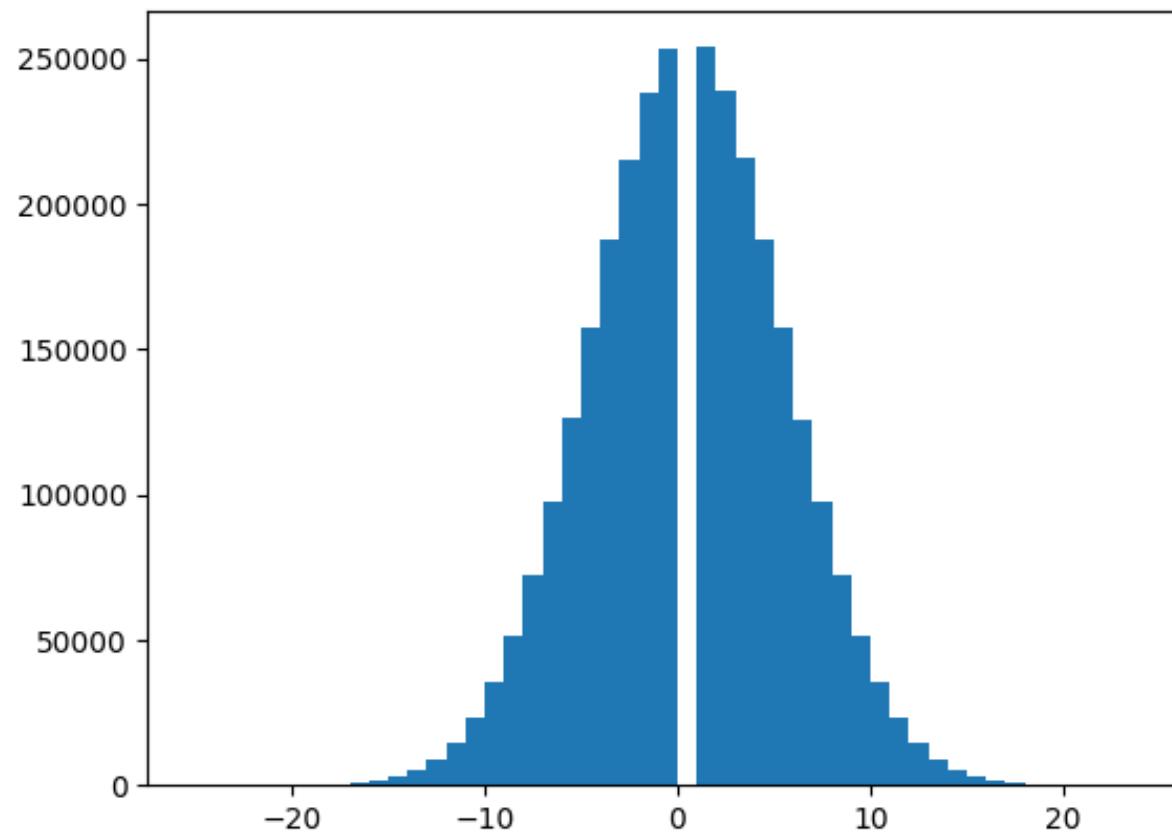


Figure 4: Points distribution

9.3.1 Grades analysis without indexes

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT		84M	17G		6392K (1)	00:04:10
1	SORT ORDER BY		84M	17G	18G	6392K (1)	00:04:10
* 2	HASH JOIN		84M	17G		2384K (1)	00:01:34
3	TABLE ACCESS FULL	SUBJECTS	76	1520		3 (0)	00:00:01
* 4	HASH JOIN		84M	15G		2383K (1)	00:01:34
5	MERGE JOIN CARTESIAN		2736	82080		64 (0)	00:00:01
6	MERGE JOIN CARTESIAN		36	360		13 (0)	00:00:01
7	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01
8	BUFFER SORT		6	30		10 (0)	00:00:01
9	TABLE ACCESS FULL	GRADES_ENUM	6	30		2 (0)	00:00:01
10	BUFFER SORT		76	1520		62 (0)	00:00:01
11	TABLE ACCESS FULL	SUBJECTS	76	1520		1 (0)	00:00:01
12	VIEW		84M	13G		2382K (1)	00:01:34
13	SORT GROUP BY		84M	9G	10G	2382K (1)	00:01:34
* 14	HASH JOIN		84M	9G	30M	26186 (4)	00:00:02
15	TABLE ACCESS FULL	GRADES	1977K	7723K		2492 (2)	00:00:01
* 16	HASH JOIN		397K	45M		19118 (2)	00:00:01
* 17	HASH JOIN		1256	128K		10204 (3)	00:00:01
18	JOIN FILTER CREATE	:BF0000	23	1426		2548 (3)	00:00:01
* 19	HASH JOIN OUTER		23	1426		2548 (3)	00:00:01
* 20	HASH JOIN ANTI		23	1035		2545 (3)	00:00:01
* 21	HASH JOIN		2337	95817		22 (0)	00:00:01
* 22	TABLE ACCESS FULL	HOUSES	1	14		3 (0)	00:00:01
* 23	TABLE ACCESS FULL	STUDENTS	9346	246K		19 (0)	00:00:01
24	VIEW	VW_NSO_1	190K	745K		2521 (3)	00:00:01
* 25	HASH JOIN		190K	3541K		2521 (3)	00:00:01
* 26	TABLE ACCESS FULL	GRADES_ENUM	1	5		3 (0)	00:00:01
* 27	TABLE ACCESS FULL	GRADES_ENUM	1	5		3 (0)	00:00:01
* 28	TABLE ACCESS FULL	GRADES_ENUM	1	5		3 (0)	00:00:01
* 29	TABLE ACCESS FULL	GRADES	1957K	26M		2505 (2)	00:00:01
30	TABLE ACCESS FULL	QUIDDITCH_TEAM_MEMBERS	97	1649		3 (0)	00:00:01
31	VIEW		502K	20M		7653 (3)	00:00:01
32	HASH GROUP BY		502K	6866K	45M	7653 (3)	00:00:01
* 33	HASH JOIN		1977K	26M		2511 (2)	00:00:01
34	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01
35	JOIN FILTER USE	:BF0000	1977K	16M		2495 (2)	00:00:01
* 36	TABLE ACCESS FULL	GRADES	1977K	16M		2495 (2)	00:00:01
* 37	TABLE ACCESS FULL	POINTS	2956K	45M		8895 (1)	00:00:01

9.3.2 Grades analysis with indexes

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT		12M	2592M		1658K (1)	00:01:05
1	SORT ORDER BY		12M	2592M		2757M (1)	00:01:05
* 2	HASH JOIN		12M	2592M		1069K (1)	00:00:42
3	TABLE ACCESS FULL	SUBJECTS	76	1520		3 (0)	00:00:01
* 4	HASH JOIN		12M	2356M		1069K (1)	00:00:42
5	MERGE JOIN CARTESIAN		2736	82080		64 (0)	00:00:01
6	MERGE JOIN CARTESIAN		36	360		13 (0)	00:00:01
7	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01
8	BUFFER SORT		6	30		10 (0)	00:00:01
9	TABLE ACCESS FULL	GRADES_ENUM	6	30		2 (0)	00:00:01
10	BUFFER SORT		76	1520		62 (0)	00:00:01
11	TABLE ACCESS FULL	SUBJECTS	76	1520		1 (0)	00:00:01
12	VIEW		12M	2003M		1069K (1)	00:00:42
13	SORT GROUP BY		12M	1425M		1069K (1)	00:00:42
* 14	HASH JOIN		84M	9698M		30M (3)	00:00:02
15	TABLE ACCESS FULL	GRADES	1977K	7723K		2492 (2)	00:00:01
* 16	HASH JOIN		397K	44M		34016 (1)	00:00:02
17	NESTED LOOPS		1256	123K		25102 (1)	00:00:01
* 18	HASH JOIN OUTER		23	1380		23946 (1)	00:00:01
19	NESTED LOOPS ANTI		23	989		23943 (1)	00:00:01
* 20	HASH JOIN		2337	95817		22 (0)	00:00:01
* 21	TABLE ACCESS FULL	HOUSES	1	14		3 (0)	00:00:01
* 22	TABLE ACCESS FULL	STUDENTS	9346	246K		19 (0)	00:00:01
23	VIEW PUSHED PREDICATE	VW_NSO_1	2	4		10 (0)	00:00:01
24	HASH JOIN		2	38		10 (0)	00:00:01
25	TABLE ACCESS FULL	GRADES_ENUM	1	5		3 (0)	00:00:01
26	TABLE ACCESS FULL	GRADES_ENUM	1	5		3 (0)	00:00:01
27	TABLE ACCESS FULL	GRADES_ENUM	1	5		3 (0)	00:00:01
28	TABLE ACCESS BY INDEX ROWID BATCHED	GRADES	21	294		7 (0)	00:00:01
29	BITMAP CONVERSION TO ROWIDS		212	2968		49 (0)	00:00:01
* 30	BITMAP INDEX RANGE SCAN	INDEX1	6	30		3 (0)	00:00:01
31	TABLE ACCESS FULL	QUIDDITCH_TEAM_MEMBERS	97	1649		3 (0)	00:00:01
32	VIEW PUSHED PREDICATE		1	41		50 (2)	00:00:01
33	SORT GROUP BY		51	714		50 (2)	00:00:01
* 34	HASH JOIN		212	2968		49 (0)	00:00:01
35	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01
36	TABLE ACCESS BY INDEX ROWID BATCHED	GRADES	212	1908		46 (0)	00:00:01
37	BITMAP CONVERSION TO ROWIDS		2956K	45M		8895 (1)	00:00:01
* 38	BITMAP INDEX RANGE SCAN	INDEX1	2956K	45M		8895 (1)	00:00:01
* 39	TABLE ACCESS FULL	POINTS	2956K	45M		8895 (1)	00:00:01

9.3.3 Points summary without indexes

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT		1105M	205G		39M (1)	00:25:29
* 1	HASH JOIN		1105M	205G		39M (1)	00:25:29
2	TABLE ACCESS FULL	TEACHERS	1097	20843		5 (0)	00:00:01
* 3	HASH JOIN		1105M	186G		39M (1)	00:25:29
4	MERGE JOIN CARTESIAN		6582	154K		26 (0)	00:00:01
5	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01
6	BUFFER SORT		1097	20843		23 (0)	00:00:01
7	TABLE ACCESS FULL	TEACHERS	1097	20843		4 (0)	00:00:01
8	VIEW		1105M	161G		39M (1)	00:25:29
9	SORT GROUP BY		1105M	118G	135G	39M (1)	00:25:29
* 10	HASH JOIN		1120M	120G	119M	65365 (13)	00:00:03
11	VIEW		2986K	85M		22617 (2)	00:00:01
12	HASH GROUP BY		2986K	34M	68M	22617 (2)	00:00:01
13	TABLE ACCESS FULL	POINTS	2986K	34M		8879 (1)	00:00:01
* 14	HASH JOIN		3506K	284M		13234 (2)	00:00:01
15	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01
* 16	HASH JOIN		3506K	267M	1400K	13207 (2)	00:00:01
* 17	HASH JOIN		16576	1197K		8928 (1)	00:00:01
* 18	HASH JOIN		73	4234		25 (0)	00:00:01
* 19	HASH JOIN		97	4268		22 (0)	00:00:01
20	TABLE ACCESS FULL	QUIDDITCH_TEAM_MEMBERS	97	1649		3 (0)	00:00:01
* 21	TABLE ACCESS FULL	STUDENTS	5504	145K		19 (0)	00:00:01
* 22	TABLE ACCESS FULL	HOUSES	3	42		3 (0)	00:00:01
* 23	TABLE ACCESS FULL	POINTS	2129K	32M		8888 (1)	00:00:01
24	TABLE ACCESS FULL	GRADES	1977K	11M		2492 (2)	00:00:01

9.3.4 Points summary with indexes

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT		1105M	205G		38M (1)	00:25:07
* 1	HASH JOIN		1105M	205G		38M (1)	00:25:07
2	TABLE ACCESS FULL	TEACHERS	1097	20843		5 (0)	00:00:01
* 3	HASH JOIN		1105M	186G		38M (1)	00:25:07
4	MERGE JOIN CARTESIAN		6582	154K		26 (0)	00:00:01
5	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01
6	BUFFER SORT		1097	20843		23 (0)	00:00:01
7	TABLE ACCESS FULL	TEACHERS	1097	20843		4 (0)	00:00:01
8	VIEW		1105M	161G		38M (1)	00:25:07
9	SORT GROUP BY		1105M	116G	131G	38M (1)	00:25:07
* 10	HASH JOIN		1120M	117G	56M	72682 (11)	00:00:03
* 11	TABLE ACCESS FULL	POINTS	2129K	32M		8888 (1)	00:00:01
* 12	HASH JOIN		4917K	454M		28129 (1)	00:00:02
13	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01
* 14	HASH JOIN		4917K	431M	2232K	28093 (1)	00:00:02
15	NESTED LOOPS		23243	1952K		23763 (1)	00:00:01
* 16	HASH JOIN		73	4234		25 (0)	00:00:01
* 17	HASH JOIN		97	4268		22 (0)	00:00:01
18	TABLE ACCESS FULL	QUIDDITCH_TEAM_MEMBERS	97	1649		3 (0)	00:00:01
* 19	TABLE ACCESS FULL	STUDENTS	5504	145K		19 (0)	00:00:01
* 20	TABLE ACCESS FULL	HOUSES	3	42		3 (0)	00:00:01
21	VIEW PUSHED PREDICATE		1	28		325 (1)	00:00:01
22	SORT GROUP BY		196	2352		325 (1)	00:00:01
23	TABLE ACCESS BY INDEX ROWID BATCHED	POINTS	320	3840		324 (0)	00:00:01
* 24	INDEX RANGE SCAN	INDEX2	320			4 (0)	00:00:01
25	TABLE ACCESS FULL	GRADES	1977K	11M		2492 (2)	00:00:01

9.3.5 Best students without indexes

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT		10	1330		22M (1)	00:14:52
1	SORT ORDER BY		10	1330		22M (1)	00:14:52
* 2	VIEW		10	1330		22M (1)	00:14:52
* 3	WINDOW SORT PUSHED RANK		631M	45G	52G	22M (1)	00:14:52
4	HASH GROUP BY		631M	45G	52G	22M (1)	00:14:52
* 5	HASH JOIN		631M	45G	79M	38180 (12)	00:00:02
* 6	TABLE ACCESS FULL	POINTS	2986K	45M		8895 (1)	00:00:01
* 7	HASH JOIN		1977K	115M		14265 (2)	00:00:01
* 8	HASH JOIN		9346	355K		22 (0)	00:00:01
9	TABLE ACCESS FULL	HOUSES	4	56		3 (0)	00:00:01
10	TABLE ACCESS FULL	STUDENTS	9346	228K		19 (0)	00:00:01
11	VIEW		1977K	41M		14229 (2)	00:00:01
12	HASH GROUP BY		1977K	35M	60M	14229 (2)	00:00:01
* 13	HASH JOIN		1977K	35M		2521 (3)	00:00:01
14	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01
* 15	TABLE ACCESS FULL	GRADES	1977K	26M		2505 (2)	00:00:01

9.3.6 Best students with indexes

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT		10	1330		19M (1)	00:12:42
1	SORT ORDER BY		10	1330		19M (1)	00:12:42
* 2	VIEW		10	1330		19M (1)	00:12:42
* 3	WINDOW SORT PUSHED RANK		48M	4012M	56G	19M (1)	00:12:42
4	HASH GROUP BY		48M	4012M	56G	19M (1)	00:12:42
* 5	HASH JOIN		631M	50G	79M	494K (3)	00:00:20
* 6	TABLE ACCESS FULL	POINTS	2986K	45M		8895 (1)	00:00:01
7	NESTED LOOPS		1977K	131M		469K (3)	00:00:19
* 8	HASH JOIN		9346	355K		22 (0)	00:00:01
9	TABLE ACCESS FULL	HOUSES	4	56		3 (0)	00:00:01
10	TABLE ACCESS FULL	STUDENTS	9346	228K		19 (0)	00:00:01
11	VIEW PUSHED PREDICATE		1	31		50 (2)	00:00:01
12	SORT GROUP BY		108	2052		50 (2)	00:00:01
* 13	HASH JOIN		212	4028		49 (0)	00:00:01
14	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01
15	TABLE ACCESS BY INDEX ROWID BATCHED	GRADES	212	2968		46 (0)	00:00:01
16	BITMAP CONVERSION TO ROWIDS						
* 17	BITMAP INDEX RANGE SCAN	INDEX1					

9.3.7 Remove points without indexes

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	DELETE STATEMENT		18383	556K	11394	(2) 00:00:01
1	DELETE	POINTS	18383	556K	11394	(2) 00:00:01
* 2	HASH JOIN RIGHT SEMI					
3	VIEW	VW_NSO_1	115	1495	2503	(2) 00:00:01
* 4	SORT GROUP BY		115	1955	2503	(2) 00:00:01
5	VIEW		2291	38947	2503	(2) 00:00:01
* 6	FILTER					
7	HASH GROUP BY		2291	50402	2503	(2) 00:00:01
* 8	HASH JOIN		45810	984K	2499	(2) 00:00:01
9	TABLE ACCESS FULL	GRADES_ENUM	6	30	3	(0) 00:00:01
* 10	TABLE ACCESS FULL	GRADES	45810	760K	2496	(2) 00:00:01
* 11	TABLE ACCESS FULL	POINTS	1494K	25M	8881	(1) 00:00:01

9.3.8 Remove points with indexes

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	DELETE STATEMENT		2078	81042	6771	(2) 00:00:01
1	DELETE	POINTS				
2	NESTED LOOPS		2078	81042	6771	(2) 00:00:01
3	NESTED LOOPS		4160	81042	6771	(2) 00:00:01
4	VIEW	VW_NSO_1	13	169	2569	(5) 00:00:01
* 5	SORT GROUP BY		13	221	2569	(5) 00:00:01
6	VIEW		248	4216	2569	(5) 00:00:01
* 7	FILTER					
8	HASH GROUP BY		248	6696	2569	(5) 00:00:01
* 9	HASH JOIN		4943	130K	2568	(5) 00:00:01
10	TABLE ACCESS FULL	GRADES_ENUM	6	30	3	(0) 00:00:01
* 11	TABLE ACCESS FULL	GRADES	4943	106K	2565	(5) 00:00:01
* 12	INDEX RANGE SCAN	INDEX2	320		3	(0) 00:00:01
* 13	TABLE ACCESS BY INDEX ROWID	POINTS	160	4160	323	(0) 00:00:01

9.4 Experiments

Two experiments were conducted, checking differences in performance after using two different types of indexes: B-tree and bitmap. Tested indexes included `index1 ON grades(student_id, award_date)` which was tested on **Grades Analysis** procedure, and `index2 ON points(student_id, award_date)` tested on **Points Summary** procedure.

Procedure name	Min. measured time				Max. measured time				Average measured time				Query plan cost			
	B-tree	Bitmap	Diff.	Profit	B-tree	Bitmap	Diff.	Profit	B-tree	Bitmap	Diff.	Profit	B-tree	Bitmap	Diff.	Profit
Grades analysis (Index 1)	607	477	130	21%	814	497	317	39%	686.2	485.5	200.7	29%	1662 K	1658 K	4 K	0.2%
Points summary (Index 2)	315	330	-15	-5%	341	374	-33	-10%	322.8	340.8	-18.0	-6%	38 M	38 M	253	0%

Query plan for `index1` showed a small advantage of using the bitmap type of index over the B-tree one. However, the measured time has proven that the difference between implementations using these two types of indexes is much higher in practice. Therefore, for `index1`, the **bitmap** type was selected.

Second experiment has shown the negligible difference between the query plan costs, however the tests shown, that the b-tree index performs better in practice. Therefore, for `index2`, the **B-tree** type was selected.

9.4.1 Grades analysis with B-tree Index

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT		12M	2592M		1662K (1)	00:01:05
1	SORT ORDER BY		12M	2592M	2757M	1662K (1)	00:01:05
* 2	HASH JOIN		12M	2592M		1073K (1)	00:00:42
3	TABLE ACCESS FULL	SUBJECTS	76	1520		3 (0)	00:00:01
* 4	HASH JOIN		12M	2356M		1073K (1)	00:00:42
5	MERGE JOIN CARTESIAN		2736	82080		64 (0)	00:00:01
6	MERGE JOIN CARTESIAN		36	360		13 (0)	00:00:01
7	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01
8	BUFFER SORT		6	30		10 (0)	00:00:01
9	TABLE ACCESS FULL	GRADES_ENUM	6	30		2 (0)	00:00:01
10	BUFFER SORT		76	1520		62 (0)	00:00:01
11	TABLE ACCESS FULL	SUBJECTS	76	1520		1 (0)	00:00:01
12	VIEW		12M	2003M		1073K (1)	00:00:42
13	SORT GROUP BY		12M	1425M	10G	1073K (1)	00:00:42
* 14	HASH JOIN		84M	9698M	30M	44918 (2)	00:00:02
15	INDEX FAST FULL SCAN	INDEX1	1977K	7723K		1736 (2)	00:00:01
* 16	HASH JOIN		397K	44M		38682 (1)	00:00:02
17	NESTED LOOPS		1256	123K		29768 (1)	00:00:02
* 18	HASH JOIN OUTER		23	1380		28087 (1)	00:00:02
19	NESTED LOOPS ANTI		23	989		28084 (1)	00:00:02
* 20	HASH JOIN		237	95817		22 (0)	00:00:01
* 21	TABLE ACCESS FULL	HOUSES	1	14		3 (0)	00:00:01
* 22	TABLE ACCESS FULL	STUDENTS	9346	246K		19 (0)	00:00:01
23	VIEW PUSHED PREDICATE	VW_NSO_1	2	2	4	12 (0)	00:00:01
24	NESTED LOOPS		2	2	38	12 (0)	00:00:01
25	NESTED LOOPS		21	38		12 (0)	00:00:01
* 26	TABLE ACCESS FULL	GRADES_ENUM	1	5		3 (0)	00:00:01
* 27	TABLE ACCESS FULL	GRADES_ENUM	1	5		3 (0)	00:00:01
* 28	TABLE ACCESS FULL	GRADES_ENUM	1	5		3 (0)	00:00:01
* 29	INDEX RANGE SCAN	INDEX1	21			2 (0)	00:00:01
* 30	TABLE ACCESS BY INDEX ROWID	GRADES	3	42		9 (0)	00:00:01
31	TABLE ACCESS FULL	QUIDDITCH_TEAM_MEMBERS	97	1649		3 (0)	00:00:01
32	VIEW PUSHED PREDICATE		1	41		73 (2)	00:00:01
33	SORT GROUP BY		51	714		73 (2)	00:00:01
* 34	HASH JOIN		212	2968		72 (0)	00:00:01
35	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01
36	TABLE ACCESS BY INDEX ROWID BATCHED	GRADES	212	1908		69 (0)	00:00:01
* 37	INDEX RANGE SCAN	INDEX1	212			3 (0)	00:00:01
* 38	TABLE ACCESS FULL	POINTS	2956K	45M		8895 (1)	00:00:01

9.4.2 Grades analysis with Bitmap Index

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT		12M	2592M		1658K (1)	00:01:05
1	SORT ORDER BY		12M	2592M	2757M	1658K (1)	00:01:05
* 2	HASH JOIN		12M	2592M		1069K (1)	00:00:42
3	TABLE ACCESS FULL	SUBJECTS	76	1520		3 (0)	00:00:01
* 4	HASH JOIN		12M	2356M		1069K (1)	00:00:42
5	MERGE JOIN CARTESIAN		2736	82080		64 (0)	00:00:01
6	MERGE JOIN CARTESIAN		36	360		13 (0)	00:00:01
7	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01
8	BUFFER SORT		6	30		10 (0)	00:00:01
9	TABLE ACCESS FULL	GRADES_ENUM	6	30		2 (0)	00:00:01
10	BUFFER SORT		76	1520		62 (0)	00:00:01
11	TABLE ACCESS FULL	SUBJECTS	76	1520		1 (0)	00:00:01
12	VIEW		12M	2003M		1069K (1)	00:00:42
13	SORT GROUP BY		12M	1425M	10G	1069K (1)	00:00:42
* 14	HASH JOIN		84M	9698M	30M	41009 (3)	00:00:02
15	TABLE ACCESS FULL	INDEX1	1977K	7723K		2492 (2)	00:00:01
* 16	HASH JOIN		397K	44M		34016 (1)	00:00:02
17	NESTED LOOPS		1256	123K		25102 (1)	00:00:01
* 18	HASH JOIN OUTER		23	1380		23946 (1)	00:00:01
19	NESTED LOOPS ANTI		23	989		23943 (1)	00:00:01
* 20	HASH JOIN		237	95817		22 (0)	00:00:01
* 21	TABLE ACCESS FULL	HOUSES	1	14		3 (0)	00:00:01
* 22	TABLE ACCESS FULL	STUDENTS	9346	246K		19 (0)	00:00:01
23	VIEW PUSHED PREDICATE	VW_NSO_1	2	2	4	10 (0)	00:00:01
24	HASH JOIN		2	38		10 (0)	00:00:01
* 25	TABLE ACCESS FULL	GRADES_ENUM	1	5		3 (0)	00:00:01
* 26	TABLE ACCESS FULL	GRADES_ENUM	1	5		3 (0)	00:00:01
* 27	TABLE ACCESS FULL	GRADES_ENUM	1	5		3 (0)	00:00:01
28	TABLE ACCESS BY INDEX ROWID BATCHED	GRADES	21	294		7 (0)	00:00:01
29	BITMAP CONVERSION TO ROWIDS		1			1	
* 30	BITMAP INDEX RANGE SCAN	INDEX1	1			1	
31	TABLE ACCESS FULL	QUIDDITCH_TEAM_MEMBERS	97	1649		3 (0)	00:00:01
32	VIEW PUSHED PREDICATE		1	41		50 (2)	00:00:01
33	SORT GROUP BY		51	714		50 (2)	00:00:01
* 34	HASH JOIN		212	2968		49 (0)	00:00:01
35	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01
36	TABLE ACCESS BY INDEX ROWID BATCHED	GRADES	212	1908		46 (0)	00:00:01
37	BITMAP CONVERSION TO ROWIDS		1			1	
* 38	BITMAP INDEX RANGE SCAN	INDEX1	1			1	
* 39	TABLE ACCESS FULL	POINTS	2956K	45M		8895 (1)	00:00:01

9.4.3 Points summary with B-tree Index

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT		1105M	205G		38M (1)	00:25:07
* 1	HASH JOIN		1105M	205G		38M (1)	00:25:07
2	TABLE ACCESS FULL	TEACHERS	1097	20843		5 (0)	00:00:01
* 3	HASH JOIN		1105M	186G		38M (1)	00:25:07
4	MERGE JOIN CARTESIAN		6582	154K		26 (0)	00:00:01
5	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01
6	BUFFER SORT		1097	20843		23 (0)	00:00:01
7	TABLE ACCESS FULL	TEACHERS	1097	20843		4 (0)	00:00:01
8	VIEW		1105M	161G		38M (1)	00:25:07
9	SORT GROUP BY		1105M	116G	131G	38M (1)	00:25:06
* 10	HASH JOIN		1120M	117G	56M	72682 (11)	00:00:03
* 11	TABLE ACCESS FULL	POINTS	2129K	32M		8888 (1)	00:00:01
* 12	HASH JOIN		4917K	454M		28129 (1)	00:00:01
13	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01
* 14	HASH JOIN		4917K	431M	2232K	28093 (1)	00:00:02
15	NESTED LOOPS		23243	1952K		23763 (1)	00:00:01
* 16	HASH JOIN		73	4234		25 (0)	00:00:01
* 17	HASH JOIN		97	4268		22 (0)	00:00:01
18	TABLE ACCESS FULL	QUIDDITCH_TEAM_MEMBERS	97	1649		3 (0)	00:00:01
* 19	TABLE ACCESS FULL	STUDENTS	5504	145K		19 (0)	00:00:01
* 20	TABLE ACCESS FULL	HOUSES	3	42		3 (0)	00:00:01
21	VIEW PUSHED PREDICATE		1	28		325 (1)	00:00:01
22	SORT GROUP BY		196	2352		325 (1)	00:00:01
23	TABLE ACCESS BY INDEX ROWID BATCHED	POINTS	320	3840		324 (0)	00:00:01
* 24	INDEX RANGE SCAN	INDEX2	320			4 (0)	00:00:01
25	TABLE ACCESS FULL	GRADES	1977K	11M		2492 (2)	00:00:01

9.4.4 Points summary with Bitmap Index

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT		1105M	205G		38M (1)	00:25:07
* 1	HASH JOIN		1105M	205G		38M (1)	00:25:07
2	TABLE ACCESS FULL	TEACHERS	1097	20843		5 (0)	00:00:01
* 3	HASH JOIN		1105M	186G		38M (1)	00:25:06
4	MERGE JOIN CARTESIAN		6582	154K		26 (0)	00:00:01
5	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01
6	BUFFER SORT		1097	20843		23 (0)	00:00:01
7	TABLE ACCESS FULL	TEACHERS	1097	20843		4 (0)	00:00:01
8	VIEW		1105M	161G		38M (1)	00:25:06
9	SORT GROUP BY		1105M	116G	131G	38M (1)	00:25:06
* 10	HASH JOIN		1120M	117G	56M	54235 (15)	00:00:03
* 11	TABLE ACCESS FULL	POINTS	2129K	32M		8888 (1)	00:00:01
* 12	HASH JOIN		4917K	454M		9682 (2)	00:00:01
13	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01
* 14	HASH JOIN		4917K	431M	2232K	9647 (2)	00:00:01
15	NESTED LOOPS		23243	1952K		5316 (2)	00:00:01
* 16	HASH JOIN		73	4234		25 (0)	00:00:01
* 17	HASH JOIN		97	4268		22 (0)	00:00:01
18	TABLE ACCESS FULL	QUIDDITCH_TEAM_MEMBERS	97	1649		3 (0)	00:00:01
* 19	TABLE ACCESS FULL	STUDENTS	5504	145K		19 (0)	00:00:01
* 20	TABLE ACCESS FULL	HOUSES	3	42		3 (0)	00:00:01
21	VIEW PUSHED PREDICATE		1	28		72 (2)	00:00:01
22	SORT GROUP BY		196	2352		72 (2)	00:00:01
23	TABLE ACCESS BY INDEX ROWID BATCHED	POINTS	320	3840		71 (0)	00:00:01
24	BITMAP CONVERSION TO ROWIDS						
* 25	BITMAP INDEX RANGE SCAN	INDEX2					
26	TABLE ACCESS FULL	GRADES	1977K	11M		2492 (2)	00:00:01

10 Partitions

10.1 Proposed partitions

10.1.1 List Partitioning on GRADES.value

Details: There are only six possible values in this column, so we can divide the data into partitions using one partition for each possible value.

Purpose: Optimization of transaction **Grades Analysis**

Possible improvements: Query Plan rows 24 and 32

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT		47M	9981M		3628K (1)	00:02:22
1	SORT ORDER BY		47M	9981M	10G	3628K (1)	00:02:22
* 2	HASH JOIN	SUBJECTS	47M	9981M		1359K (1)	00:00:54
3	TABLE ACCESS FULL		76	1520		3 (0)	00:00:01
* 4	HASH JOIN		47M	9074M		1359K (1)	00:00:54
5	MERGE JOIN CARTESIAN		2736	82080		64 (0)	00:00:01
6	MERGE JOIN CARTESIAN		36	360		13 (0)	00:00:01
7	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01
8	BUFFER SORT		6	30		10 (0)	00:00:01
9	TABLE ACCESS FULL	GRADES_ENUM	6	30		2 (0)	00:00:01
10	BUFFER SORT		76	1520		62 (0)	00:00:01
11	TABLE ACCESS FULL	SUBJECTS	76	1520		1 (0)	00:00:01

12	VIEW		47M	7712M	1358K	(1)	00:00:54	
13	SORT GROUP BY		47M	5671M	6092M	(1)	00:00:54	
* 14	HASH JOIN		47M	5671M	28M	24853	(3)	00:00:01
* 15	HASH JOIN		224K	25M		19118	(2)	00:00:01
* 16	HASH JOIN		711	74655		10204	(3)	00:00:01
17	JOIN FILTER CREATE	:BF0000	13	806		2548	(3)	00:00:01
* 18	HASH JOIN OUTER		13	806		2548	(3)	00:00:01
* 19	HASH JOIN ANTI		13	585		2545	(3)	00:00:01
* 20	HASH JOIN		1323	54243		22	(0)	00:00:01
* 21	TABLE ACCESS FULL	HOUSES	1	14		3	(0)	00:00:01
* 22	TABLE ACCESS FULL	STUDENTS	5291	139K		19	(0)	00:00:01
23	VIEW	VW_NS0_1	190K	745K		2521	(3)	00:00:01
* 24	HASH JOIN		190K	3541K		2521	(3)	00:00:01
* 25	TABLE ACCESS FULL	GRADES_ENUM	1	5		3	(0)	00:00:01
* 26	TABLE ACCESS FULL	GRADES_ENUM	1	5		3	(0)	00:00:01
* 27	TABLE ACCESS FULL	GRADES_ENUM	1	5		3	(0)	00:00:01
* 28	TABLE ACCESS FULL	GRADES	1957K	26M		2505	(2)	00:00:01
29	TABLE ACCESS FULL	QUIDDITCH_TEAM_MEMBERS	97	1649		3	(0)	00:00:01
30	VIEW		502K	20M		7653	(3)	00:00:01
31	HASH GROUP BY		502K	6866K	45M	7653	(3)	00:00:01
* 32	HASH JOIN		1977K	26M		2511	(2)	00:00:01
33	TABLE ACCESS FULL	GRADES_ENUM	6	30		3	(0)	00:00:01
34	JOIN FILTER USE	:BF0000	1977K	16M		2495	(2)	00:00:01
* 35	TABLE ACCESS FULL	GRADES	1977K	16M		2495	(2)	00:00:01
* 36	TABLE ACCESS FULL	POINTS	2956K	45M		8895	(1)	00:00:01
37	TABLE ACCESS FULL	GRADES	1977K	7723K		2492	(2)	00:00:01

Predicate Information (identified by operation id):

```

2 - access("Q"."WORST SUBJECT"="S2"."ID")
4 - access("Q"."BEST SUBJECT"="S1"."ID" AND "Q"."MIN GRADE"="G2"."VALUE" AND
           "Q"."MAX GRADE"="G1"."VALUE")
14 - access("STUDENTS"."ID"="GRADES"."STUDENT_ID")
15 - access("STUDENTS"."ID"="POINTS"."STUDENT_ID")
16 - access("AVERAGES"."STUDENT_ID"="STUDENTS"."ID")
18 - access("STUDENTS"."ID"="QUIDDITCH_TEAM_MEMBERS"."STUDENT_ID"(+))
19 - access("STUDENTS"."ID"="ID")
20 - access("STUDENTS"."HOUSE_ID"="HOUSES"."ID")
21 - filter("HOUSES"."NAME"='Gryffindor')
22 - filter("STUDENTS"."YEAR"=1 OR "STUDENTS"."YEAR"=3 OR "STUDENTS"."YEAR"=5 OR "STUDENTS"."YEAR"=7)
->24 - access("GRADES"."VALUE"="GRADES_ENUM"."SYMBOL")
25 - filter("GRADES_ENUM"."VALUE" < (SELECT "GRADES_ENUM"."VALUE" FROM "GRADES_ENUM" "GRADES_ENUM" WHERE
           "GRADES_ENUM"."SYMBOL"='T') OR "GRADES_ENUM"."VALUE" > (SELECT "GRADES_ENUM"."VALUE" FROM "GRADES_ENUM"
           "GRADES_ENUM" WHERE "GRADES_ENUM"."SYMBOL"='O'))
26 - filter("GRADES_ENUM"."SYMBOL"='T')
27 - filter("GRADES_ENUM"."SYMBOL"='O')
28 - filter("GRADES"."AWARD_DATE" <= TO_DATE(' 2024-06-25 00:00:00 ', 'syyyy-mm-dd hh24:mi:ss') AND
           "GRADES"."AWARD_DATE" >= TO_DATE(' 2023-09-02 00:00:00 ', 'syyyy-mm-dd hh24:mi:ss'))
->32 - access("GRADES"."VALUE"="GRADES_ENUM"."SYMBOL")
35 - filter(SYS_OP_BLOOM_FILTER(:BF0000,"GRADES"."STUDENT_ID"))
36 - filter("POINTS"."AWARD_DATE" <= TO_DATE(' 2024-06-25 00:00:00 ', 'syyyy-mm-dd hh24:mi:ss') AND
           "POINTS"."AWARD_DATE" >= TO_DATE(' 2023-09-02 00:00:00 ', 'syyyy-mm-dd hh24:mi:ss'))

```

10.1.2 Range Partitioning on POINTS.award_date

Details: We can divide the data into partitions such that points awarded in different months are on different partitions.

Purpose: Optimization of transaction Points Summary

Possible improvements: Query Plan row 23

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT					39M (1)	00:25:29
* 1	HASH JOIN					39M (1)	00:25:29
2	TABLE ACCESS FULL	TEACHERS	1097	20843		5 (0)	00:00:01
* 3	HASH JOIN					39M (1)	00:25:29
4	MERGE JOIN CARTESIAN					26 (0)	00:00:01
5	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01
6	BUFFER SORT		1097	20843		23 (0)	00:00:01
7	TABLE ACCESS FULL	TEACHERS	1097	20843		4 (0)	00:00:01
8	VIEW					39M (1)	00:25:29
9	SORT GROUP BY					39M (1)	00:25:29
* 10	HASH JOIN					65365 (13)	00:00:03
11	VIEW					22617 (2)	00:00:01
12	HASH GROUP BY					22617 (2)	00:00:01
13	TABLE ACCESS FULL	POINTS	2986K	34M	68M	8879 (1)	00:00:01
* 14	HASH JOIN					13234 (2)	00:00:01
15	TABLE ACCESS FULL	GRADES_ENUM	3506K	284M			
* 16	HASH JOIN		6	30		3 (0)	00:00:01
* 17	HASH JOIN		3506K	267M	1400K	13207 (2)	00:00:01
* 18	HASH JOIN		16576	1197K		8928 (1)	00:00:01
* 19	HASH JOIN		73	4234		25 (0)	00:00:01
20	TABLE ACCESS FULL	QUIDDITCH_TEAM_MEMBERS	97	4268		22 (0)	00:00:01
* 21	TABLE ACCESS FULL	STUDENTS	5504	145K		19 (0)	00:00:01
* 22	TABLE ACCESS FULL	HOUSES	3	42		3 (0)	00:00:01
* 23	TABLE ACCESS FULL	POINTS	2129K	32M		8888 (1)	00:00:01
24	TABLE ACCESS FULL	GRADES	1977K	11M		2492 (2)	00:00:01

Predicate Information (identified by operation id):

```

1 - access("Q"."WORST_TEACHER"="T2"."ID")
3 - access("Q"."BEST_TEACHER"="T1"."ID" AND "Q"."AVG_GRADE"="GRADES_ENUM"."VALUE")
10 - access("AVERAGES"."STUDENT_ID"="STUDENTS"."ID")
14 - access("GRADES"."VALUE"="GRADES_ENUM"."SYMBOL")
16 - access("STUDENTS"."ID"="GRADES"."STUDENT_ID")
17 - access("STUDENTS"."ID"="POINTS"."STUDENT_ID")
18 - access("STUDENTS"."HOUSE_ID"="HOUSES"."ID")
19 - access("STUDENTS"."ID"="QUIDDITCH_TEAM_MEMBERS"."STUDENT_ID")
21 - filter("STUDENTS"."YEAR">>4 OR "STUDENTS"."YEAR"<3)
22 - filter("HOUSES"."NAME"<>'Hufflepuff')
->23 - filter("POINTS"."AWARD_DATE">>TO_DATE(' 2024-01-06 00:00:00', 'syyyy-mm-dd hh24:mi:ss') OR
      "POINTS"."AWARD_DATE" < TO_DATE(' 2023-12-06 00:00:00', 'syyyy-mm-dd hh24:mi:ss'))
```

10.1.3 Hash Partitioning on STUDENTS.id

Details: We want to evenly distribute students of different IDs between partitions.

Purpose: Optimization of transaction Points Summary

Possible improvements: Query Plan rows 10 and 16

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time	
0	SELECT STATEMENT					39M (1)	00:25:29	
* 1	HASH JOIN					39M (1)	00:25:29	
2	TABLE ACCESS FULL	TEACHERS	1097	20843		5 (0)	00:00:01	
* 3	HASH JOIN					39M (1)	00:25:29	
4	MERGE JOIN CARTESIAN		6582	154K		26 (0)	00:00:01	
5	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01	
6	BUFFER SORT		1097	20843		23 (0)	00:00:01	
7	TABLE ACCESS FULL	TEACHERS	1097	20843		4 (0)	00:00:01	
8	VIEW					39M (1)	00:25:29	
9	SORT GROUP BY				135G	39M (1)	00:25:29	
* 10	HASH JOIN		1120M	120G	119M	65365 (13)	00:00:03	
11	VIEW			2986K	85M	22617 (2)	00:00:01	
12	HASH GROUP BY			2986K	34M	22617 (2)	00:00:01	
13	TABLE ACCESS FULL	POINTS	2986K	34M		8879 (1)	00:00:01	
* 14	HASH JOIN			3506K	284M	13234 (2)	00:00:01	
15	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01	
* 16	HASH JOIN			3506K	267M	1400K	13207 (2)	00:00:01
* 17	HASH JOIN			16576	1197K		8928 (1)	00:00:01
* 18	HASH JOIN			73	4234		25 (0)	00:00:01
* 19	HASH JOIN			97	4268		22 (0)	00:00:01
20	TABLE ACCESS FULL	QUIDDITCH_TEAM_MEMBERS	97	1649		3 (0)	00:00:01	
* 21	TABLE ACCESS FULL	STUDENTS	5504	145K		19 (0)	00:00:01	
* 22	TABLE ACCESS FULL	HOUSES	3	42		3 (0)	00:00:01	
* 23	TABLE ACCESS FULL	POINTS	2129K	32M		8888 (1)	00:00:01	
24	TABLE ACCESS FULL	GRADES	1977K	11M		2492 (2)	00:00:01	

Predicate Information (identified by operation id):

```

1 - access("Q"."WORST_TEACHER"="T2"."ID")
3 - access("Q"."BEST_TEACHER"="T1"."ID" AND "Q"."AVG_GRADE"="GRADES_ENUM"."VALUE")
->10 - access("AVERAGES"."STUDENT_ID"="STUDENTS"."ID")
14 - access("GRADES"."VALUE"="GRADES_ENUM"."SYMBOL")
->16 - access("STUDENTS"."ID"="GRADES"."STUDENT_ID")
17 - access("STUDENTS"."ID"="POINTS"."STUDENT_ID")
18 - access("STUDENTS"."HOUSE_ID"="HOUSES"."ID")
19 - access("STUDENTS"."ID"="QUIDDITCH_TEAM_MEMBERS"."STUDENT_ID")
21 - filter("STUDENTS"."YEAR">>4 OR "STUDENTS"."YEAR"><3)
22 - filter("HOUSES"."NAME"<>'Hufflepuff')
23 - filter("POINTS"."AWARD_DATE">>TO_DATE(' 2024-01-06 00:00:00', 'syyyy-mm-dd hh24:mi:ss') OR
      "POINTS"."AWARD_DATE" < TO_DATE(' 2023-12-06 00:00:00', 'syyyy-mm-dd hh24:mi:ss'))
```

10.2 Proposed experiments

10.2.1 Experiment 1: Partitioning the GRADES table – List vs. Hash

Partitioning schemes being compared:

- *Scheme A (Proposed)*: List partitioning of the GRADES table based on the `value` column (6 partitions – one for each grade).
- *Scheme B (Alternative)*: Hash partitioning of the GRADES table using the `student_id` column (e.g., into 4 or 8 partitions).

Goal/Hypothesis: The idea is to check which approach performs better. The experiment will test whether a more even data spread from hashing by `student_id` helps when running queries on GRADES, especially when grades aren't evenly distributed or when queries don't hit just one list partition. This will be compared to using list partitioning on the low-cardinality `value` column.

Tested transactions: *Grades analysis*

10.2.2 Experiment 2: Partitioning the POINTS table – Range vs. Composite (Range + Hash)

Partitioning schemes being compared:

- *Scheme A (Proposed)*: Range partitioning of the POINTS table based on `award_date` (monthly).
- *Scheme B (Alternative)*: Composite partitioning – first by range on `award_date` (monthly), and then subpartitioned by hash on `teacher_id` (e.g., 4 subpartitions per month).

Goal/Hypothesis: The goal is to see whether combining range and hash partitioning improves performance when queries filter by both date and `teacher_id`. This setup might allow for more accurate pruning and better workload distribution.

Tested transactions: *Points summary*

10.2.3 Experiment 3: Logical placement of partitions (Tablespaces) in the POINTS table

Base partitioning: Range partitioning of the POINTS table by `award_date` (monthly).

Configurations being compared:

- *Configuration A*: All partitions stored in the default single tablespace.
- *Configuration B*: Some partitions (e.g., older vs. newer months) stored in different, dedicated tablespaces.

Goal/Hypothesis: The goal is to find out whether putting partitions in different tablespaces (even if they're on the same disk) changes how the database handles queries or tracks I/O stats. This could give insight into whether using separate tablespaces helps with managing data and optimizing systems with more than one storage device.

Tested transactions: *Points Summary*

11 Partitions - implementation

11.1 Implemented partitions

As declared before, we specified three partitions: GRADES.value, POINTS.award_date, STUDENTS.id. In order to test them efficiently, we created a second database with the same data, but slightly altered creation process:

```
CREATE TABLE Grades (
    id NUMBER NOT NULL PRIMARY KEY,
    value VARCHAR2(1) NOT NULL,
    award_date DATE NOT NULL,
    student_id NUMBER NOT NULL,
    subject_id NUMBER NOT NULL,
    teacher_id NUMBER NOT NULL
)
PARTITION BY LIST(value) (
    PARTITION troll VALUES ('T'),
    PARTITION dread VALUES ('D'),
    PARTITION poor VALUES ('P'),
    PARTITION accept VALUES ('A'),
    PARTITION expec VALUES ('E'),
    PARTITION outst VALUES ('O'),
    PARTITION other VALUES (DEFAULT)
);

CREATE TABLE Points (
    id NUMBER NOT NULL PRIMARY KEY,
    value NUMBER NOT NULL,
    description VARCHAR2(1024),
    award_date DATE NOT NULL,
    student_id NUMBER NOT NULL,
    teacher_id NUMBER
)
PARTITION BY RANGE(award_date) (
    PARTITION september VALUES LESS THAN (TO_DATE('2023-10-01', 'YYYY-MM-DD')),
    PARTITION october VALUES LESS THAN (TO_DATE('2023-11-01', 'YYYY-MM-DD')),
    PARTITION november VALUES LESS THAN (TO_DATE('2023-12-01', 'YYYY-MM-DD')),
    PARTITION december VALUES LESS THAN (TO_DATE('2024-01-01', 'YYYY-MM-DD')),
    PARTITION january VALUES LESS THAN (TO_DATE('2024-02-01', 'YYYY-MM-DD')),
    PARTITION february VALUES LESS THAN (TO_DATE('2024-03-01', 'YYYY-MM-DD')),
    PARTITION march VALUES LESS THAN (TO_DATE('2024-04-01', 'YYYY-MM-DD')),
    PARTITION april VALUES LESS THAN (TO_DATE('2024-05-01', 'YYYY-MM-DD')),
    PARTITION may VALUES LESS THAN (TO_DATE('2024-06-01', 'YYYY-MM-DD')),
    PARTITION june VALUES LESS THAN (TO_DATE('2024-07-01', 'YYYY-MM-DD')),
    PARTITION future VALUES LESS THAN (MAXVALUE)
);

CREATE TABLE Students (
    id NUMBER NOT NULL PRIMARY KEY,
```

```

name VARCHAR2(64) NOT NULL,
surname VARCHAR2(64) NOT NULL,
gender CHAR NOT NULL,
date_of_birth DATE NOT NULL,
year NUMBER NOT NULL,
hogsmeade_consent NUMBER(1) DEFAULT 0 NOT NULL,
house_id NUMBER NOT NULL,
dormitory_id NUMBER
)
PARTITION BY HASH(id)
PARTITIONS 4
STORE IN (users , second_disk);

```

11.2 Cost comparison

We ran the whole workload on the database with and without partitions, **5 times** in each case. Dividing the database into partitions positively influenced most of our queries (according to query plans). Similar to what was the case with indexes - actual running times on our machine were often actually bigger with than without partitions.

Procedure name	Min. measured time				Max. measured time				Average measured time				Query plan cost			
	Default	Parti.	Diff.	Profit	Default	Parti.	Diff.	Profit	Default	Parti.	Diff.	Profit	Default	Parti.	Diff.	Profit
Grades analysis	564	498	66	11.7 %	813	537	276	33.9 %	677.0	514.4	162.6	24.0 %	3628 K	18300	3609 K	99.5 %
Points summary	377	468	-91	-24.1 %	556	540	16	2.9 %	468.6	507.8	-39.2	-8.4 %	39 M	24273	38.9 M	99.9 %
Best students	221	214	7	3.2 %	297	248	49	16.5 %	250.0	229.4	20.6	8.2 %	22 M	13368	21.9 M	99.9 %
Assign subjects	2735	6249	-3514	-128.5 %	3932	8869	-4937	-125.6 %	3403.8	7079.2	-3675.4	-108.0 %	12318	4942	7376	60 %
Remove points	26694	37578	-10884	-40.8 %	33966	40328	-6362	-18.7 %	30823.8	38952.6	-8128.8	-26.4 %	11460	13365	-1905	-16.6 %
Whole workload	30591	45007	-14416	-47.1 %	39564	50522	-10958	-27.7 %	35623.2	47283.4	-11660.2	-32.7 %	64.7 M	74248	64.6 M	99.9 %

11.3 Query plans comparison

Creating partitions changed all the query plans in our workload, and the new query plans are very different from the old ones. Below we present their comparison.

11.3.1 Grades analysis without partitions

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT		47M	9981M		3628K (1)	00:02:22
1	SORT ORDER BY		47M	9981M	10G	3628K (1)	00:02:22
* 2	HASH JOIN		47M	9981M		1359K (1)	00:00:54
3	TABLE ACCESS FULL	SUBJECTS	76	1520		3 (0)	00:00:01
* 4	HASH JOIN		47M	9074M		1359K (1)	00:00:54
5	MERGE JOIN CARTESIAN		2736	82080		64 (0)	00:00:01
6	MERGE JOIN CARTESIAN		36	360		13 (0)	00:00:01
7	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01
8	BUFFER SORT		6	30		10 (0)	00:00:01
9	TABLE ACCESS FULL	GRADES_ENUM	6	30		2 (0)	00:00:01
10	BUFFER SORT		76	1520		62 (0)	00:00:01
11	TABLE ACCESS FULL	SUBJECTS	76	1520		1 (0)	00:00:01
12	VIEW		47M	7712M		1358K (1)	00:00:54
13	SORT GROUP BY		47M	5671M	6092M	1358K (1)	00:00:54
* 14	HASH JOIN		47M	5671M	28M	24853 (3)	00:00:01
* 15	HASH JOIN		224K	25M		19118 (2)	00:00:01
* 16	HASH JOIN		711	74655		10204 (3)	00:00:01
17	JOIN FILTER CREATE	:BF0000	13	806		2548 (3)	00:00:01
* 18	HASH JOIN OUTER		13	806		2548 (3)	00:00:01
* 19	HASH JOIN ANTI		13	585		2545 (3)	00:00:01
* 20	HASH JOIN		1323	54243		22 (0)	00:00:01
* 21	TABLE ACCESS FULL	HOUSES	1	14		3 (0)	00:00:01
* 22	TABLE ACCESS FULL	STUDENTS	5291	139K		19 (0)	00:00:01
23	VIEW	VW_NSO_1	190K	745K		2521 (3)	00:00:01
* 24	HASH JOIN		190K	3541K		2521 (3)	00:00:01
* 25	TABLE ACCESS FULL	GRADES_ENUM	1	5		3 (0)	00:00:01
* 26	TABLE ACCESS FULL	GRADES_ENUM	1	5		3 (0)	00:00:01
* 27	TABLE ACCESS FULL	GRADES_ENUM	1	5		3 (0)	00:00:01
* 28	TABLE ACCESS FULL	GRADES	1957K	26M		2505 (2)	00:00:01
29	TABLE ACCESS FULL	QUIDDITCH_TEAM_MEMBERS	97	1649		3 (0)	00:00:01
30	VIEW		502K	20M		7653 (3)	00:00:01
31	HASH GROUP BY		502K	6866K	45M	7653 (3)	00:00:01
* 32	HASH JOIN		1977K	26M		2511 (2)	00:00:01
33	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01
34	JOIN FILTER USE	:BF0000	1977K	16M		2495 (2)	00:00:01
* 35	TABLE ACCESS FULL	GRADES	1977K	16M		2495 (2)	00:00:01
* 36	TABLE ACCESS FULL	POINTS	2956K	45M		8895 (1)	00:00:01
37	TABLE ACCESS FULL	GRADES	1977K	7723K		2492 (2)	00:00:01

11.3.2 Grades analysis with partitions

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	Pstart	Pstop
0	SELECT STATEMENT		1	294	18300 (1)	00:00:01		
1	SORT ORDER BY		1	294	18300 (1)	00:00:01		
2	NESTED LOOPS		1	294	18299 (1)	00:00:01		
3	NESTED LOOPS		1	294	18299 (1)	00:00:01		
4	NESTED LOOPS		1	247	18299 (1)	00:00:01		
* 5	HASH JOIN		1	200	18299 (1)	00:00:01		
6	MERGE JOIN CARTESIAN		1	30	6 (0)	00:00:01		
7	TABLE ACCESS FULL	GRADES_ENUM	1	15	3 (0)	00:00:01		
8	BUFFER SORT		1	15	3 (0)	00:00:01		
9	TABLE ACCESS FULL	GRADES_ENUM	1	15	3 (0)	00:00:01		
10	VIEW		1	170	18293 (1)	00:00:01		
11	SORT GROUP BY		1	280	18293 (1)	00:00:01		
* 12	HASH JOIN		1	280	15829 (1)	00:00:01		
* 13	HASH JOIN OUTER		1	228	13364 (1)	00:00:01		
* 14	HASH JOIN		1	181	13361 (1)	00:00:01		
15	NESTED LOOPS		1	146	2460 (1)	00:00:01		
16	NESTED LOOPS		1	146	2460 (1)	00:00:01		
17	NESTED LOOPS		1	123	2460 (1)	00:00:01		
18	PARTITION LIST ALL		1	13	2460 (1)	00:00:01	1	7
19	TABLE ACCESS FULL	GRADES	1	13	2460 (1)	00:00:01	1	7
* 20	TABLE ACCESS BY GLOBAL INDEX ROWID	STUDENTS	1	110	0 (0)	00:00:01	ROWID	ROWID
* 21	INDEX UNIQUE SCAN	SYS_CO09092	1	0	0 (0)	00:00:01		
22	NESTED LOOPS SEMI		1	39	2463 (1)	00:00:01		
* 23	TABLE ACCESS FULL	GRADES_ENUM	1	15	3 (0)	00:00:01		
* 24	TABLE ACCESS FULL	GRADES_ENUM	1	15	3 (0)	00:00:01		
* 25	TABLE ACCESS FULL	GRADES_ENUM	1	15	3 (0)	00:00:01		
26	PARTITION LIST ITERATOR		1	24	2460 (1)	00:00:01	KEY	KEY
* 27	TABLE ACCESS FULL	GRADES	1	24	2460 (1)	00:00:01	KEY	KEY
* 28	INDEX UNIQUE SCAN	SYS_CO09102	1	0	0 (0)	00:00:01		
* 29	TABLE ACCESS BY INDEX ROWID	HOUSES	1	23	0 (0)	00:00:01		
30	PARTITION RANGE ITERATOR		1	35	10901 (1)	00:00:01	1	10
* 31	TABLE ACCESS FULL	POINTS	1	35	10901 (1)	00:00:01	1	10
32	TABLE ACCESS FULL	QUIDDITCH_TEAM_MEMBERS	1	47	3 (0)	00:00:01		
33	VIEW		1	52	2464 (1)	00:00:01		
34	HASH GROUP BY		1	43	2464 (1)	00:00:01		
35	NESTED LOOPS		1	43	2463 (1)	00:00:01		
36	TABLE ACCESS FULL	GRADES_ENUM	1	15	3 (0)	00:00:01		
37	PARTITION LIST ITERATOR		1	28	2460 (1)	00:00:01	KEY	KEY
* 38	TABLE ACCESS FULL	GRADES	1	28	2460 (1)	00:00:01	KEY	KEY
39	TABLE ACCESS BY INDEX ROWID	SUBJECTS	1	47	0 (0)	00:00:01		
* 40	INDEX UNIQUE SCAN	SYS_CO09130	1	0	0 (0)	00:00:01		
* 41	INDEX UNIQUE SCAN	SYS_CO09130	1	47	0 (0)	00:00:01		
42	TABLE ACCESS BY INDEX ROWID	SUBJECTS	1	47	0 (0)	00:00:01		

11.3.3 Points summary without partitions

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT		1105M	205G		39M (1)	00:25:29
* 1	HASH JOIN		1105M	205G		39M (1)	00:25:29
2	TABLE ACCESS FULL	TEACHERS	1097	20843		5 (0)	00:00:01
* 3	HASH JOIN		1105M	186G		39M (1)	00:25:29
4	MERGE JOIN CARTESIAN		6582	154K		26 (0)	00:00:01
5	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01
6	BUFFER SORT		1097	20843		23 (0)	00:00:01
7	TABLE ACCESS FULL	TEACHERS	1097	20843		4 (0)	00:00:01
8	VIEW		1105M	161G		39M (1)	00:25:29
9	SORT GROUP BY		1105M	118G	135G	39M (1)	00:25:29
* 10	HASH JOIN		1120M	120G	119M	65365 (13)	00:00:03
11	VIEW		2986K	85M		22617 (2)	00:00:01
12	HASH GROUP BY		2986K	34M	68M	22617 (2)	00:00:01
13	TABLE ACCESS FULL	POINTS	2986K	34M		8879 (1)	00:00:01
* 14	HASH JOIN		3506K	284M		13234 (2)	00:00:01
15	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01
* 16	HASH JOIN		3506K	267M	1400K	13207 (2)	00:00:01
* 17	HASH JOIN		16576	1197K		8928 (1)	00:00:01
* 18	HASH JOIN		73	4234		25 (0)	00:00:01
* 19	HASH JOIN		97	4268		22 (0)	00:00:01
20	TABLE ACCESS FULL	QUIDDITCH_TEAM_MEMBERS	97	1649		3 (0)	00:00:01
* 21	TABLE ACCESS FULL	STUDENTS	5504	145K		19 (0)	00:00:01
* 22	TABLE ACCESS FULL	HOUSES	3	42		3 (0)	00:00:01
* 23	TABLE ACCESS FULL	POINTS	2129K	32M		8888 (1)	00:00:01
24	TABLE ACCESS FULL	GRADES	1977K	11M		2492 (2)	00:00:01

11.3.4 Points summary with partitions

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	Pstart	Pstop
0	SELECT STATEMENT		1	334	24273 (1)	00:00:01		
1	NESTED LOOPS		1	334	24273 (1)	00:00:01		
2	NESTED LOOPS		1	334	24273 (1)	00:00:01		
3	NESTED LOOPS		1	253	24273 (1)	00:00:01		
* 4	HASH JOIN		1	172	24273 (1)	00:00:01		
5	TABLE ACCESS FULL	GRADES_ENUM	1	15	3 (0)	00:00:01		
6	VIEW		1	157	24270 (1)	00:00:01		
7	SORT GROUP BY		1	284	24270 (1)	00:00:01		
* 8	HASH JOIN		1	284	24269 (1)	00:00:01		
* 9	HASH JOIN		1	245	13367 (1)	00:00:01		
10	NESTED LOOPS		1	210	2466 (1)	00:00:01		
11	NESTED LOOPS		1	210	2466 (1)	00:00:01		
* 12	HASH JOIN		1	187	2466 (1)	00:00:01		
13	NESTED LOOPS		1	140	2463 (1)	00:00:01		
14	NESTED LOOPS		1	140	2463 (1)	00:00:01		
15	NESTED LOOPS		1	30	2463 (1)	00:00:01		
16	TABLE ACCESS FULL	GRADES_ENUM	1	15	3 (0)	00:00:01		
17	PARTITION LIST ITERATOR		1	15	2460 (1)	00:00:01	KEY	KEY
* 18	TABLE ACCESS FULL	GRADES	1	15	2460 (1)	00:00:01	KEY	KEY
* 19	INDEX UNIQUE SCAN	SYS_C0009092	1	1	0 (0)	00:00:01		
* 20	TABLE ACCESS BY GLOBAL INDEX ROWID	STUDENTS	1	110	0 (0)	00:00:01	ROWID	ROWID
21	TABLE ACCESS FULL	QUIDDITCH_TEAM_MEMBERS	1	47	3 (0)	00:00:01		
* 22	INDEX UNIQUE SCAN	SYS_C0009102	1	1	0 (0)	00:00:01		
* 23	TABLE ACCESS BY INDEX ROWID	HOUSES	1	23	0 (0)	00:00:01		
24	PARTITION RANGE OR		1	35	10901 (1)	00:00:01	[KEY(OR)]	[KEY(OR)]
* 25	TABLE ACCESS FULL	POINTS	1	35	10901 (1)	00:00:01	[KEY(OR)]	[KEY(OR)]
26	VIEW		1	39	10902 (1)	00:00:01		
27	HASH GROUP BY		1	39	10902 (1)	00:00:01		
28	PARTITION RANGE ALL		1	39	10901 (1)	00:00:01	1	11
29	TABLE ACCESS FULL	POINTS	1	39	10901 (1)	00:00:01	1	11
30	TABLE ACCESS BY INDEX ROWID	TEACHERS	1	81	0 (0)	00:00:01		
* 31	INDEX UNIQUE SCAN	SYS_C0009112	1	1	0 (0)	00:00:01		
* 32	INDEX UNIQUE SCAN	SYS_C0009112	1	1	0 (0)	00:00:01		
33	TABLE ACCESS BY INDEX ROWID	TEACHERS	1	81	0 (0)	00:00:01		

11.3.5 Best students without partitions

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time	
0	SELECT STATEMENT		10	1330		22M (1)	00:14:52	
1	SORT ORDER BY		10	1330		22M (1)	00:14:52	
* 2	VIEW		10	1330		22M (1)	00:14:52	
* 3	WINDOW SORT PUSHED RANK		631M	45G	52G	22M (1)	00:14:52	
4	HASH GROUP BY		631M	45G	52G	22M (1)	00:14:52	
* 5	HASH JOIN		631M	45G	79M	38180 (12)	00:00:02	
* 6	TABLE ACCESS FULL	POINTS	2986K	45M		8895 (1)	00:00:01	
* 7	HASH JOIN		1977K	115M		14265 (2)	00:00:01	
* 8	HASH JOIN		9346	355K		22 (0)	00:00:01	
9	TABLE ACCESS FULL	HOUSES	4	56		3 (0)	00:00:01	
10	TABLE ACCESS FULL	STUDENTS	9346	228K		19 (0)	00:00:01	
11	VIEW		1977K	41M		14229 (2)	00:00:01	
12	HASH GROUP BY		1977K	35M	60M	14229 (2)	00:00:01	
* 13	HASH JOIN		1977K	35M		2521 (3)	00:00:01	
14	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01	
* 15	TABLE ACCESS FULL	GRADES	1977K	26M		2505 (2)	00:00:01	

11.3.6 Best students with partitions

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	Pstart	Pstop
0	SELECT STATEMENT		10	1330	13368 (1)	00:00:01		
1	SORT ORDER BY		10	1330	13368 (1)	00:00:01		
* 2	VIEW		10	1330	13367 (1)	00:00:01		
* 3	WINDOW SORT PUSHED RANK		1	135	13367 (1)	00:00:01		
4	HASH GROUP BY		1	135	13367 (1)	00:00:01		
5	VIEW	VM_NWWV_1	1	135	13365 (1)	00:00:01		
6	HASH GROUP BY		1	240	13365 (1)	00:00:01		
7	NESTED LOOPS		1	240	13364 (1)	00:00:01		
8	MERGE JOIN CARTESIAN		1	216	10904 (1)	00:00:01		
9	NESTED LOOPS		1	201	10901 (1)	00:00:01		
10	NESTED LOOPS		1	201	10901 (1)	00:00:01		
11	NESTED LOOPS		1	166	10901 (1)	00:00:01		
12	PARTITION RANGE ITERATOR		1	47	10901 (1)	00:00:01	1	10
* 13	TABLE ACCESS FULL	POINTS	1	47	10901 (1)	00:00:01	1	10
14	TABLE ACCESS BY GLOBAL INDEX ROWID	STUDENTS	1	119	0 (0)	00:00:01	ROWID	ROWID
* 15	INDEX UNIQUE SCAN	SYS_C009092	1		0 (0)	00:00:01		
* 16	INDEX UNIQUE SCAN	SYS_C009102	1		0 (0)	00:00:01		
17	TABLE ACCESS BY INDEX ROWID	HOUSES	1	35	0 (0)	00:00:01		
18	BUFFER SORT		1	15	10904 (1)	00:00:01		
19	TABLE ACCESS FULL	GRADES_ENUM	1	15	3 (0)	00:00:01		
20	PARTITION LIST ITERATOR		1	24	2460 (1)	00:00:01	KEY	KEY
* 21	TABLE ACCESS FULL	GRADES	1	24	2460 (1)	00:00:01	KEY	KEY

11.3.7 Assign subjects without partitions

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time	
0	INSERT STATEMENT		1017	52920		12318 (3)	00:00:01	
1	LOAD TABLE CONVENTIONAL	STUDENTS SUBJECTS						
2	HASH UNIQUE		1017	52920		12318 (3)	00:00:01	
3	UNION-ALL							
* 4	HASH JOIN		1016	26416		6155 (3)	00:00:01	
* 5	HASH JOIN ANTI		93	1860		6152 (3)	00:00:01	
* 6	TABLE ACCESS FULL	STUDENTS	9346	65422		19 (0)	00:00:01	
7	VIEW	VW_NSO_1	25113	318K		6133 (3)	00:00:01	
* 8	FILTER							
9	SORT GROUP BY		25113	343K	45M	6133 (3)	00:00:01	
* 10	HASH JOIN		1977K	26M		2511 (2)	00:00:01	
11	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01	
12	TABLE ACCESS FULL	GRADES	1977K	16M		2495 (2)	00:00:01	
* 13	TABLE ACCESS FULL	SUBJECTS	76	456		3 (0)	00:00:01	
* 14	HASH JOIN ANTI		1	44		6158 (3)	00:00:01	
* 15	HASH JOIN		97	3007		25 (0)	00:00:01	
16	MERGE JOIN CARTESIAN		97	2328		6 (0)	00:00:01	
* 17	TABLE ACCESS FULL	SUBJECTS	1	20		3 (0)	00:00:01	
18	BUFFER SORT		97	388		3 (0)	00:00:01	
19	TABLE ACCESS FULL	QUIDDITCH_TEAM_MEMBERS	97	388		3 (0)	00:00:01	
* 20	TABLE ACCESS FULL	STUDENTS	8011	56077		19 (0)	00:00:01	
* 21	TABLE ACCESS FULL	SUBJECTS	1	20		3 (0)	00:00:01	
22	VIEW	VW_NSO_2	25113	318K		6133 (3)	00:00:01	
* 23	FILTER							
24	SORT GROUP BY		25113	343K	45M	6133 (3)	00:00:01	
* 25	HASH JOIN		1977K	26M		2511 (2)	00:00:01	
26	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01	
27	TABLE ACCESS FULL	GRADES	1977K	16M		2495 (2)	00:00:01	

11.3.8 Assign subjects with partitions

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	Pstart	Pstop
0	INSERT STATEMENT		2	302	4942 (1)	00:00:01		
1	LOAD TABLE CONVENTIONAL	STUDENTS SUBJECTS						
2	HASH UNIQUE		2	302	4942 (1)	00:00:01		
3	UNION-ALL							
* 4	HASH JOIN		1	65	2468 (1)	00:00:01		
5	MERGE JOIN ANTI		1	39	2465 (1)	00:00:01		
* 6	TABLE ACCESS BY GLOBAL INDEX ROWID	STUDENTS	1	26	0 (0)	00:00:01	ROWID	ROWID
7	INDEX FULL SCAN	SYS_CO09092	1	1	0 (0)	00:00:01		
* 8	SORT UNIQUE		1	13	2465 (1)	00:00:01		
9	VIEW	VW_NSO_1	1	13	2464 (1)	00:00:01		
* 10	FILTER							
11	SORT GROUP BY		1	43	2464 (1)	00:00:01		
12	NESTED LOOPS		1	43	2463 (1)	00:00:01		
13	TABLE ACCESS FULL	GRADES_ENUM	1	15	3 (0)	00:00:01		
14	PARTITION LIST ITERATOR		1	28	2460 (1)	00:00:01	KEY	KEY
* 15	TABLE ACCESS FULL	GRADES	1	28	2460 (1)	00:00:01	KEY	KEY
* 16	TABLE ACCESS FULL	SUBJECTS	1	26	3 (0)	00:00:01		
17	MERGE JOIN CARTESIAN		1	86	6 (0)	00:00:01		
18	NESTED LOOPS		1	39	3 (0)	00:00:01		
19	NESTED LOOPS		1	39	3 (0)	00:00:01		
20	TABLE ACCESS FULL	QUIDDITCH_TEAM_MEMBERS	1	13	3 (0)	00:00:01		
* 21	INDEX UNIQUE SCAN	SYS_CO09092	1	1	0 (0)	00:00:01		
* 22	FILTER							
23	HASH GROUP BY		1	43	2464 (1)	00:00:01		
24	NESTED LOOPS		1	43	2463 (1)	00:00:01		
25	TABLE ACCESS FULL	GRADES_ENUM	1	15	3 (0)	00:00:01		
26	PARTITION LIST ITERATOR		1	28	2460 (1)	00:00:01	KEY	KEY
* 27	TABLE ACCESS FULL	GRADES	1	28	2460 (1)	00:00:01	KEY	KEY
* 28	TABLE ACCESS BY GLOBAL INDEX ROWID	STUDENTS	1	26	0 (0)	00:00:01	ROWID	ROWID
* 29	TABLE ACCESS FULL	SUBJECTS	1	47	3 (0)	00:00:01		
30	BUFFER SORT		1	47	6 (0)	00:00:01		
* 31	TABLE ACCESS FULL	SUBJECTS	1	47	3 (0)	00:00:01		

11.3.9 Remove points without partitions

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	DELETE STATEMENT		2078	64418	11460	(2) 00:00:01
1	DELETE	POINTS	1	1	1	1
* 2	HASH JOIN RIGHT SEMI		2078	64418	11460	(2) 00:00:01
3	VIEW	VW_NSO_1	13	169	2569	(5) 00:00:01
* 4	SORT GROUP BY		13	221	2569	(5) 00:00:01
5	VIEW		248	4216	2569	(5) 00:00:01
6	FILTER		1	1	1	1
7	HASH GROUP BY		248	5456	2569	(5) 00:00:01
* 8	HASH JOIN		4943	106K	2568	(5) 00:00:01
9	TABLE ACCESS FULL	GRADES_ENUM	6	30	3	(0) 00:00:01
* 10	TABLE ACCESS FULL	GRADES	4943	84031	2565	(5) 00:00:01
* 11	TABLE ACCESS FULL	POINTS	1494K	25M	8881	(1) 00:00:01

11.3.10 Remove points with partitions

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	Pstart	Pstop
0	DELETE STATEMENT		1	65	13365	(1) 00:00:01	1	1
1	DELETE	POINTS	1	1	1	1	1	1
* 2	HASH JOIN SEMI		1	65	13365	(1) 00:00:01	1	1
3	PARTITION RANGE ALL		1	52	10901	(1) 00:00:01	1	11
* 4	TABLE ACCESS FULL	POINTS	1	52	10901	(1) 00:00:01	1	11
5	VIEW	VW_NSO_1	1	13	2464	(1) 00:00:01	1	1
* 6	SORT GROUP BY		1	26	2464	(1) 00:00:01	1	1
7	VIEW		1	26	2464	(1) 00:00:01	1	1
* 8	FILTER		1	1	1	1	1	1
9	HASH GROUP BY		1	52	2464	(1) 00:00:01	1	1
10	NESTED LOOPS		1	52	2463	(1) 00:00:01	1	1
11	TABLE ACCESS FULL	GRADES_ENUM	1	15	3	(0) 00:00:01	1	1
12	PARTITION LIST ITERATOR		1	37	2460	(1) 00:00:01	KEY	KEY
* 13	TABLE ACCESS FULL	GRADES	1	37	2460	(1) 00:00:01	KEY	KEY

11.4 Experiments

Some experiments were conducted, checking differences in performance between two different types of partitions for each test. For Grades analysis procedure we compared List partition on GRADES.value (default) with Hash partition on GRADES.student_id. For Points analysis we compared Range partition on POINTS.award_date with composite partition - firstly Range partition on POINTS.award_date with Hash subpartition on POINTS.teacher_id.

Procedure name	Min. measured time				Max. measured time				Average measured time				Query plan cost			
	Default	Alt.	Diff.	Profit	Default	Alt.	Diff.	Profit	Default	Alt.	Diff.	Profit	Default	Alt.	Diff.	Profit
Grades analysis	498	1108	-610	-122 %	537	1178	-641	-119 %	514.4	1144.0	-629.6	-122 %	18300	493 M	-493 M	-2694 K %
Points summary	468	806	-338	-72 %	540	857	-317	-59 %	507.8	820.6	-312.8	-62 %	24273	246 K	-222 K	-913 %

Both query plans prove that our initially selected partitions are much more efficient than the alternatives.

11.4.1 Grades analysis with default partition

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	Pstart	Pstop
0	SELECT STATEMENT		1	294	18300	(1) 00:00:01		
1	SORT ORDER BY		1	294	18300	(1) 00:00:01		
2	NESTED LOOPS		1	294	18299	(1) 00:00:01		
3	NESTED LOOPS		1	294	18299	(1) 00:00:01		
4	NESTED LOOPS		1	247	18299	(1) 00:00:01		
* 5	HASH JOIN		1	200	18299	(1) 00:00:01		
6	MERGE JOIN CARTESIAN		1	30	6	(0) 00:00:01		
7	TABLE ACCESS FULL	GRADES_ENUM	1	15	3	(0) 00:00:01		
8	BUFFER SORT		1	15	3	(0) 00:00:01		
9	TABLE ACCESS FULL	GRADES_ENUM	1	15	3	(0) 00:00:01		
10	VIEW		1	170	18293	(1) 00:00:01		
11	SORT GROUP BY		1	280	18293	(1) 00:00:01		
* 12	HASH JOIN		1	280	15829	(1) 00:00:01		
* 13	HASH JOIN OUTER		1	228	13364	(1) 00:00:01		
* 14	HASH JOIN		1	181	13361	(1) 00:00:01		
15	NESTED LOOPS		1	146	2460	(1) 00:00:01		
16	NESTED LOOPS		1	146	2460	(1) 00:00:01		
17	NESTED LOOPS		1	123	2460	(1) 00:00:01		
18	PARTITION LIST ALL		1	13	2460	(1) 00:00:01	1	7
19	TABLE ACCESS FULL	GRADES	1	13	2460	(1) 00:00:01	1	7
* 20	TABLE ACCESS BY GLOBAL INDEX ROWID	STUDENTS	1	110	0	(0) 00:00:01	ROWID	ROWID
* 21	INDEX UNIQUE SCAN	SYS_CO09092	1	1	0	(0) 00:00:01		
22	NESTED LOOPS SEMI		1	39	2463	(1) 00:00:01		
* 23	TABLE ACCESS FULL	GRADES_ENUM	1	15	3	(0) 00:00:01		
* 24	TABLE ACCESS FULL	GRADES_ENUM	1	15	3	(0) 00:00:01		
* 25	TABLE ACCESS FULL	GRADES_ENUM	1	15	3	(0) 00:00:01		
26	PARTITION LIST ITERATOR		1	24	2460	(1) 00:00:01	KEY	KEY
* 27	TABLE ACCESS FULL	GRADES	1	24	2460	(1) 00:00:01	KEY	KEY
* 28	INDEX UNIQUE SCAN	SYS_CO09102	1	1	0	(0) 00:00:01		
* 29	TABLE ACCESS BY INDEX ROWID	HOUSES	1	23	0	(0) 00:00:01		
30	PARTITION RANGE ITERATOR		1	35	10901	(1) 00:00:01	1	10
* 31	TABLE ACCESS FULL	POINTS	1	35	10901	(1) 00:00:01	1	10
32	TABLE ACCESS FULL	QUIDDITCH_TEAM_MEMBERS	1	47	3	(0) 00:00:01		
33	VIEW		1	52	2464	(1) 00:00:01		
34	HASH GROUP BY		1	43	2464	(1) 00:00:01		
35	NESTED LOOPS		1	43	2463	(1) 00:00:01		
36	TABLE ACCESS FULL	GRADES_ENUM	1	15	3	(0) 00:00:01		
37	PARTITION LIST ITERATOR		1	28	2460	(1) 00:00:01	KEY	KEY
* 38	TABLE ACCESS FULL	GRADES	1	28	2460	(1) 00:00:01	KEY	KEY
39	TABLE ACCESS BY INDEX ROWID	SUBJECTS	1	47	0	(0) 00:00:01		
* 40	INDEX UNIQUE SCAN	SYS_CO09130	1	1	0	(0) 00:00:01		
* 41	INDEX UNIQUE SCAN	SYS_CO09130	1	1	0	(0) 00:00:01		
42	TABLE ACCESS BY INDEX ROWID	SUBJECTS	1	47	0	(0) 00:00:01		

11.4.2 Grades analysis with alternative partition

Id	Operation	Name	Rows	Bytes	TempSpcl	Cost (%CPU)	Time	Pstart	Pstop
0	SELECT STATEMENT							5764M	1578G
1	SORT ORDER BY							5764M	1578G
* 2	HASH JOIN							5764M	1578G
3	TABLE ACCESS FULL	SUBJECTS	76	3572				963K	(77)
* 4	HASH JOIN							5764M	1325G
5	MERGE JOIN CARTESIAN							924K	(76)
6	MERGE JOIN CARTESIAN							2736	205K
7	TABLE ACCESS FULL	GRADES_ENUM	6	90				64	(0)
8	BUFFER SORT							36	1080
9	TABLE ACCESS FULL	GRADES_ENUM	6	90				13	(0)
10	BUFFER SORT							76	3572
11	TABLE ACCESS FULL	SUBJECTS	76	3572				1	(0)
12	VIEW							5764M	912G
13	SORT GROUP BY							5764M	1513G
* 14	HASH JOIN RIGHT OUTER							5764M	1513G
15	TABLE ACCESS FULL	QUIDDITCH_TEAM_MEMBERS	97	4559				326K	(30)
* 16	HASH JOIN							5764M	1261G
17	PART JOIN FILTER CREATE	:BF0000						97M	288K
18	VIEW							1590K	78M
19	HASH GROUP BY							1590K	65M
20	HASH JOIN							1590K	65M
21	TABLE ACCESS FULL	GRADES_ENUM	6	90				3319	(2)
22	PARTITION HASH ALL							1590K	42M
23	TABLE ACCESS FULL	GRADES	1590K	42M				3305	(1)
* 24	HASH JOIN							21M	3690M
25	PART JOIN FILTER CREATE	:BF0001						21M	3690M
26	HASH JOIN							77571	12M
* 27	NESTED LOOPS ANTI							268	36180
28	HASH JOIN							1459	189K
29	TABLE ACCESS FULL	HOUSES	1	23				1096	(1)
30	PARTITION HASH JOIN-FILTER							1	3
31	TABLE ACCESS FULL	STUDENTS	5835	626K				1093	(1)
32	VIEW PUSHED PREDICATE	VW_NSO_1	5835	626K				1093	(1)
33	NESTED LOOPS		1	2				7	(0)
34	TABLE ACCESS FULL	GRADES_ENUM	1	39				7	(0)
35	TABLE ACCESS FULL	GRADES_ENUM	1	15				3	(0)
36	TABLE ACCESS FULL	GRADES_ENUM	1	15				3	(0)
37	PARTITION HASH SINGLE							4	(0)
38	TABLE ACCESS FULL	GRADES	2608	62592				4	(0)
39	PARTITION RANGE ITERATOR							2400K	80M
40	PARTITION HASH ALL							29684	(64)
41	TABLE ACCESS FULL	POINTS	2400K	80M				29684	(64)
42	PARTITION HASH JOIN-FILTER							1590K	19M
43	TABLE ACCESS FULL	GRADES	1590K	19M				3303	(1)

11.4.3 Points summary with default partition

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	Pstart	Pstop
0	SELECT STATEMENT		1	334	24273	(1)	00:00:01	
1	NESTED LOOPS		1	334	24273	(1)	00:00:01	
2	NESTED LOOPS		1	334	24273	(1)	00:00:01	
3	NESTED LOOPS		1	253	24273	(1)	00:00:01	
* 4	HASH JOIN		1	172	24273	(1)	00:00:01	
5	TABLE ACCESS FULL	GRADES_ENUM	1	15	3	(0)	00:00:01	
6	VIEW		1	157	24270	(1)	00:00:01	
7	SORT GROUP BY		1	284	24270	(1)	00:00:01	
* 8	HASH JOIN		1	284	24269	(1)	00:00:01	
* 9	HASH JOIN		1	245	13367	(1)	00:00:01	
10	NESTED LOOPS		1	210	2466	(1)	00:00:01	
11	NESTED LOOPS		1	210	2466	(1)	00:00:01	
* 12	HASH JOIN		1	187	2466	(1)	00:00:01	
13	NESTED LOOPS		1	140	2463	(1)	00:00:01	
14	NESTED LOOPS		1	140	2463	(1)	00:00:01	
15	NESTED LOOPS		1	30	2463	(1)	00:00:01	
16	TABLE ACCESS FULL	GRADES_ENUM	1	15	3	(0)	00:00:01	
17	PARTITION LIST ITERATOR		1	15	2460	(1)	00:00:01	KEY KEY
* 18	TABLE ACCESS FULL	GRADES	1	15	2460	(1)	00:00:01	KEY KEY
19	INDEX UNIQUE SCAN	SYS_CO09092	1	1	0	(0)	00:00:01	
* 20	TABLE ACCESS BY GLOBAL INDEX ROWID	STUDENTS	1	110	0	(0)	00:00:01	ROWID ROWID
21	TABLE ACCESS FULL	QUIDDITCH_TEAM_MEMBERS	1	47	3	(0)	00:00:01	
* 22	INDEX UNIQUE SCAN	SYS_CO09102	1	1	0	(0)	00:00:01	
* 23	TABLE ACCESS BY INDEX ROWID	HOUSES	1	23	0	(0)	00:00:01	
24	PARTITION RANGE OR		1	35	10901	(1)	00:00:01	KEY(OR) KEY(OR)
* 25	TABLE ACCESS FULL	POINTS	1	35	10901	(1)	00:00:01	KEY(OR) KEY(OR)
26	VIEW		1	39	10902	(1)	00:00:01	
27	HASH GROUP BY		1	39	10902	(1)	00:00:01	
28	PARTITION RANGE ALL		1	39	10901	(1)	00:00:01	1 11
29	TABLE ACCESS FULL	POINTS	1	39	10901	(1)	00:00:01	1 11
30	TABLE ACCESS BY INDEX ROWID	TEACHERS	1	1	81	0	(0)	00:00:01
* 31	INDEX UNIQUE SCAN	SYS_CO09112	1	1	0	(0)	00:00:01	
* 32	INDEX UNIQUE SCAN	SYS_CO09112	1	1	0	(0)	00:00:01	
33	TABLE ACCESS BY INDEX ROWID	TEACHERS	1	1	81	0	(0)	00:00:01

11.4.4 Points summary with alternative partition

Id	Operation	Name	Rows	Bytes	TempSpcl	Cost (%CPU)	Time	Pstart	Pstop
0	SELECT STATEMENT		1336M	415G		246K	(61)	00:00:10	
* 1	HASH JOIN	TEACHERS	1097	88857		5	(0)	00:00:01	
2	TABLE ACCESS FULL		1336M	315G		237K	(59)	00:00:10	
* 3	HASH JOIN		6582	617K		26	(0)	00:00:01	
4	MERGE JOIN CARTESIAN		6	90		3	(0)	00:00:01	
5	TABLE ACCESS FULL	GRADES_ENUM	1097	88857		23	(0)	00:00:01	
6	BUFFER SORT		1097	88857		4	(0)	00:00:01	
7	TABLE ACCESS FULL	TEACHERS	1336M	195G		228K	(58)	00:00:09	
8	VIEW		1336M	353G		228K	(58)	00:00:09	
9	SORT GROUP BY		1336M	353G		113M	107K	(9)	00:00:05
* 10	HASH JOIN		2536K	84M		11323	(1)	00:00:01	KEY(OR) KEY(OR)
11	PARTITION RANGE OR		2536K	84M		11323	(1)	00:00:01	1 4
12	PARTITION HASH ALL		2536K	84M		11323	(1)	00:00:01	KEY KEY
* 13	TABLE ACCESS FULL	POINTS	5170K	1227M		17321	(3)	00:00:01	
* 14	HASH JOIN		5170K	1153M		5904K	17284	(3)	00:00:01
15	TABLE ACCESS FULL	GRADES_ENUM	6	90		3	(0)	00:00:01	
* 16	HASH JOIN		5170K	1153M		17284	(3)	00:00:01	
17	PART JOIN FILTER CREATE	:BF0000	5170K	1153M		26164	5595K	11618	(3)
* 18	HASH JOIN		73	13140		103	(0)	00:00:01	
19	JOIN FILTER CREATE	:BF0001	73	13140		103	(0)	00:00:01	
* 20	HASH JOIN		73	13140		103	(0)	00:00:01	
* 21	TABLE ACCESS FULL	HOUSES	3	69		3	(0)	00:00:01	
22	NESTED LOOPS		97	15229		100	(0)	00:00:01	
23	NESTED LOOPS		97	15229		100	(0)	00:00:01	
24	TABLE ACCESS FULL	QUIDDITCH_TEAM_MEMBERS	97	4559		3	(0)	00:00:01	
* 25	INDEX UNIQUE SCAN	SYS_CO09169	1	1		0	(0)	00:00:01	
* 26	TABLE ACCESS BY GLOBAL INDEX ROWID	STUDENTS	1	110		1	(0)	00:00:01	ROWID ROWID
27	VIEW		2894K	107M		11496	(3)	00:00:01	
28	HASH GROUP BY		2894K	107M		11496	(3)	00:00:01	
29	JOIN FILTER USE	:BF0001	2894K	107M		11309	(1)	00:00:01	
30	PARTITION RANGE ALL		2894K	107M		11309	(1)	00:00:01	1 11
31	PARTITION HASH ALL		2894K	107M		11309	(1)	00:00:01	1 4
* 32	TABLE ACCESS FULL	POINTS	2894K	107M		11309	(1)	00:00:01	1 44
33	PARTITION HASH JOIN-FILTER		1590K	22M		3303	(1)	00:00:01	:BF0000 :BF0000
34	TABLE ACCESS FULL	GRADES	1590K	22M		3303	(1)	00:00:01	:BF0000 :BF0000

12 Columnar store - implementation

12.1 Implemented columnar storage

We implemented Oracle In-Memory columnar storage for our databases. We configured it for key tables that are frequently accessed in our workload and benefit from columnar compression and parallel processing. These tables are often subject to aggregate functions, so the columnar storage should improve the efficiency of our workload.

First, we selected the in-memory size to be 250 MB using:

```
ALTER SYSTEM SET inmemory_size=250M SCOPE=SPFILE;
```

We selected tables for columnar storage based on the usage of these tables in our workload. For all of the selected tables, we used MEMCOMPRESS FOR QUERY HIGH, because experiments showed that it is usually the best option for us. Analysis of the workload made us select the following tables:

- **GRADES and POINTS** (Critical Priority): Large fact tables with millions of records. Heavy analytical processing with aggregations, filtering, and joins. Core tables for most workload queries. However, column POINTS.description should be excluded from columnar storage as it is a column with relatively big textual data that is not being used in any query.

- **STUDENTS** (High Priority): Central dimensional table frequently joined with fact tables. Medium size but high join frequency across all analytical queries.
- **SUBJECTS and STUDENTS SUBJECTS** (High Priority): Key dimensional and junction tables essential for query performance. Moderate size but critical for join operations.

Small lookup tables (HOUSES, TEACHERS, GRADES_ENUM) remained in traditional row storage due to their size and access patterns being more suitable for OLTP operations.

ALTER TABLE Points

```
INMEMORY MEMCOMPRESS FOR QUERY HIGH PRIORITY CRITICAL  
NO INMEMORY (description);
```

ALTER TABLE Grades

```
INMEMORY MEMCOMPRESS FOR QUERY HIGH PRIORITY CRITICAL;
```

ALTER TABLE Students

```
INMEMORY MEMCOMPRESS FOR QUERY HIGH PRIORITY HIGH;
```

ALTER TABLE Subjects

```
INMEMORY MEMCOMPRESS FOR QUERY HIGH PRIORITY HIGH;
```

ALTER TABLE Students_Subjects

```
INMEMORY MEMCOMPRESS FOR QUERY HIGH PRIORITY HIGH;
```

12.2 Cost comparison

We ran the whole workload on the database with columnar store enabled and disabled, **10 times** in each case. Enabling columnar store influenced positively 5 out of 6 queries in our workload in terms of query plan cost. In terms of actual running times, the improvement is only visible in 3 out of 6 queries. For the whole workload, however, both running time and query plan cost were overall improved, so the usage of columnar store in our database is completely justified.

Proc. name	Min. measured time				Max. measured time				Average measured time				Query plan cost			
	Disa.	Ena.	Gain	Profit	Disa.	Ena.	Gain	Profit	Disa.	Ena.	Gain	Profit	Disa.	Ena.	Gain	Profit
Grades analysis	564	517	-47	-8.3 %	813	685	-128	-15.7 %	677.0	565.8	-111.2	-16.4 %	3628 K	816 K	-2812 K	-77.5 %
Points summary	377	492	115	30.5 %	556	659	103	18.5 %	468.6	558.7	90.1	19.2 %	39 M	260 K	-38740 K	-99.3 %
Best students	221	291	70	31.7 %	297	357	60	20.2 %	250.0	319.9	69.9	28.0 %	22 M	38839	-21961 K	-99.8 %
Raise grades	26694	1767	-24927	-93.4 %	33966	9302	-24664	-72.6 %	30823.8	7923.3	-22900.5	-74.3 %	11460	833 K	821540	7168.8 %
Assign subjects	2735	2232	-503	-18.4 %	3932	3250	-682	-17.3 %	3403.8	2571.8	-832	-24.4 %	12318	2095	-10223	-83.0 %
Remove points	26694	41439	14745	55.2 %	33966	56561	22595	66.5 %	30823.8	44774.3	13950.5	45.3 %	11460	6819	-4641	-40.5 %
Whole workload	57285	46738	-10547	-18.4 %	73530	70814	-2716	-3.7 %	66447	56713.8	-9733.2	-14.6 %	64663 K	1957 K	-62706 K	-97.0 %

12.3 Query plan comparison

Enabling columnar store significantly changed all the query plans. Below we present their comparison.

12.3.1 Grades analysis without columnar store

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT		47M	9981M		3628K (1)	00:02:22
1	SORT ORDER BY		47M	9981M	10G	3628K (1)	00:02:22
* 2	HASH JOIN		47M	9981M		1359K (1)	00:00:54
3	TABLE ACCESS FULL	SUBJECTS	76	1520		3 (0)	00:00:01
* 4	HASH JOIN		47M	9074M		1359K (1)	00:00:54
5	MERGE JOIN CARTESIAN		2736	82080		64 (0)	00:00:01
6	MERGE JOIN CARTESIAN		36	360		13 (0)	00:00:01
7	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01
8	BUFFER SORT		6	30		10 (0)	00:00:01
9	TABLE ACCESS FULL	GRADES_ENUM	6	30		2 (0)	00:00:01
10	BUFFER SORT		76	1520		62 (0)	00:00:01
11	TABLE ACCESS FULL	SUBJECTS	76	1520		1 (0)	00:00:01
12	VIEW		47M	7712M		1358K (1)	00:00:54
13	SORT GROUP BY		47M	5671M	6092M	1358K (1)	00:00:54
* 14	HASH JOIN		47M	5671M	28M	24853 (3)	00:00:01
* 15	HASH JOIN		224K	25M		19118 (2)	00:00:01
* 16	HASH JOIN		711	74655	10204	(3) (0)	00:00:01
17	JOIN FILTER CREATE	:BF0000	13	806		2548 (3)	00:00:01
* 18	HASH JOIN OUTER		13	806		2548 (3)	00:00:01
* 19	HASH JOIN ANTI		13	585		2545 (3)	00:00:01
* 20	HASH JOIN		1323	54243	22	(0) (0)	00:00:01
* 21	TABLE ACCESS FULL	HOUSES	1	14		3 (0)	00:00:01
* 22	TABLE ACCESS FULL	STUDENTS	5291	139K		19 (0)	00:00:01
23	VIEW	VW_NSO_1	190K	745K		2521 (3)	00:00:01
* 24	HASH JOIN		190K	3541K		2521 (3)	00:00:01
* 25	TABLE ACCESS FULL	GRADES_ENUM	1	5		3 (0)	00:00:01
* 26	TABLE ACCESS FULL	GRADES_ENUM	1	5		3 (0)	00:00:01
* 27	TABLE ACCESS FULL	GRADES_ENUM	1	5		3 (0)	00:00:01
* 28	TABLE ACCESS FULL	GRADES	1957K	26M		2505 (2)	00:00:01
29	TABLE ACCESS FULL	QUIDDITCH_TEAM_MEMBERS	97	1649		3 (0)	00:00:01
30	VIEW		502K	20M		7653 (3)	00:00:01
31	HASH GROUP BY		502K	6866K	45M	7653 (3)	00:00:01
* 32	HASH JOIN		1977K	26M		2511 (2)	00:00:01
33	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01
34	JOIN FILTER USE	:BF0000	1977K	16M		2495 (2)	00:00:01
* 35	TABLE ACCESS FULL	GRADES	1977K	16M		2495 (2)	00:00:01
* 36	TABLE ACCESS FULL	POINTS	2956K	45M		8895 (1)	00:00:01
37	TABLE ACCESS FULL	GRADES	1977K	7723K		2492 (2)	00:00:01

12.3.2 Grades analysis with columnar store

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT		110	32340		816K (74)	00:00:32
1	SORT ORDER BY		110	32340		816K (74)	00:00:32
* 2	HASH JOIN		110	32340		816K (74)	00:00:32
3	TABLE ACCESS INMEMORY FULL	SUBJECTS	76	3572		3 (0)	00:00:01
* 4	HASH JOIN		110	27170		816K (74)	00:00:32
5	TABLE ACCESS INMEMORY FULL	SUBJECTS	76	3572		3 (0)	00:00:01
* 6	HASH JOIN		110	22000		816K (74)	00:00:32
7	TABLE ACCESS FULL	GRADES_ENUM	6	90		3 (0)	00:00:01
* 8	HASH JOIN		110	20350		816K (74)	00:00:32
9	TABLE ACCESS FULL	GRADES_ENUM	6	90		3 (0)	00:00:01
10	VIEW		110	18700		816K (74)	00:00:32
11	SORT GROUP BY		110	32230		816K (74)	00:00:32
* 12	HASH JOIN RIGHT OUTER		5468M	1492G		287K (26)	00:00:12
13	TABLE ACCESS FULL	QUIDDITCH_TEAM_MEMBERS	97	4559		3 (0)	00:00:01
* 14	HASH JOIN		5468M	1252G	39M	250K (15)	00:00:10
15	TABLE ACCESS INMEMORY FULL	GRADES	1656K	20M		100 (7)	00:00:01
* 16	HASH JOIN		17M	3939M	18M	6856 (6)	00:00:01
17	JOIN FILTER CREATE	:BF0000	101K	17M		556 (23)	00:00:01
18	HASH JOIN		101K	17M		556 (23)	00:00:01
19	JOIN FILTER CREATE	:BF0001	315	45990		136 (27)	00:00:01
* 20	HASH JOIN ANTI		315	45990		136 (27)	00:00:01
21	JOIN FILTER CREATE	:BF0002	1343	174K		4 (0)	00:00:01
22	HASH JOIN		1343	174K		4 (0)	00:00:01
23	JOIN FILTER CREATE	:BF0003	1	23		3 (0)	00:00:01
24	TABLE ACCESS FULL	HOUSES	1	23		3 (0)	00:00:01
25	JOIN FILTER USE	:BF0003	5370	576K		1 (0)	00:00:01
26	TABLE ACCESS INMEMORY FULL	STUDENTS	5370	576K		1 (0)	00:00:01
27	VIEW	VW_NSO_1	158K	2012K		130 (27)	00:00:01
28	HASH JOIN		158K	6038K		130 (27)	00:00:01
29	JOIN FILTER CREATE	:BF0004	1	15		3 (0)	00:00:01
30	TABLE ACCESS FULL	GRADES_ENUM	1	15		3 (0)	00:00:01
31	TABLE ACCESS FULL	GRADES_ENUM	1	15		3 (0)	00:00:01
* 32	TABLE ACCESS FULL	GRADES_ENUM	1	15		3 (0)	00:00:01
33	JOIN FILTER USE	:BF0002	1626K	37M		116 (20)	00:00:01
34	JOIN FILTER USE	:BF0004	1626K	37M		116 (20)	00:00:01
* 35	TABLE ACCESS INMEMORY FULL	GRADES	1626K	37M		116 (20)	00:00:01
36	JOIN FILTER USE	:BF0001	2958K	98M		401 (19)	00:00:01
37	TABLE ACCESS INMEMORY FULL	POINTS	2958K	98M		401 (19)	00:00:01
38	VIEW		1656K	82M		229 (59)	00:00:01
39	HASH GROUP BY		1656K	170M		229 (59)	00:00:01
* 40	HASH JOIN		1656K	170M		126 (24)	00:00:01
41	VIEW	VW_GBC_8	6	168		4 (25)	00:00:01
42	HASH GROUP BY		6	90		4 (25)	00:00:01
43	TABLE ACCESS FULL	GRADES_ENUM	6	90		3 (0)	00:00:01
44	VIEW	VW_GBF_9	1656K	126M		111 (17)	00:00:01
45	HASH GROUP BY		1656K	44M		111 (17)	00:00:01
46	JOIN FILTER USE	:BF0000	1656K	44M		110 (16)	00:00:01
* 47	TABLE ACCESS INMEMORY FULL	GRADES	1656K	44M		110 (16)	00:00:01

12.3.3 Points summary without columnar store

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT		1105M	205G		39M (1)	00:25:29
* 1	HASH JOIN		1105M	205G		39M (1)	00:25:29
2	TABLE ACCESS FULL	TEACHERS	1097	20843		5 (0)	00:00:01
* 3	HASH JOIN		1105M	186G		39M (1)	00:25:29
4	MERGE JOIN CARTESIAN		6582	154K		26 (0)	00:00:01
5	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01
6	BUFFER SORT		1097	20843		23 (0)	00:00:01
7	TABLE ACCESS FULL	TEACHERS	1097	20843		4 (0)	00:00:01
8	VIEW		1105M	161G		39M (1)	00:25:29
9	SORT GROUP BY		1105M	118G	135G	39M (1)	00:25:29
* 10	HASH JOIN		1120M	120G	119M	65365 (13)	00:00:03
11	VIEW		2986K	85M	22617	(2)	00:00:01
12	HASH GROUP BY		2986K	34M	68M	22617	(2) 00:00:01
13	TABLE ACCESS FULL	POINTS	2986K	34M		8879 (1)	00:00:01
* 14	HASH JOIN		3506K	284M		13234 (2)	00:00:01
15	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01
* 16	HASH JOIN		3506K	267M	1400K	13207 (2)	00:00:01
* 17	HASH JOIN		16576	1197K		8928 (1)	00:00:01
* 18	HASH JOIN		73	4234		25 (0)	00:00:01
* 19	HASH JOIN		97	4268		22 (0)	00:00:01
20	TABLE ACCESS FULL	QUIDDITCH_TEAM_MEMBERS	97	1649		3 (0)	00:00:01
* 21	TABLE ACCESS FULL	STUDENTS	5504	145K		19 (0)	00:00:01
* 22	TABLE ACCESS FULL	HOUSES	3	42		3 (0)	00:00:01
* 23	TABLE ACCESS FULL	POINTS	2129K	32M		8888 (1)	00:00:01
24	TABLE ACCESS FULL	GRADES	1977K	11M		2492 (2)	00:00:01

12.3.4 Points summary with columnar store

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT		1668M	518G		260K (72)	00:00:11
* 1	HASH JOIN		1668M	518G		260K (72)	00:00:11
2	TABLE ACCESS FULL	TEACHERS	1097	88857		5 (0)	00:00:01
* 3	HASH JOIN		1668M	393G		249K (71)	00:00:10
4	MERGE JOIN CARTESIAN		6582	617K		26 (0)	00:00:01
5	TABLE ACCESS FULL	GRADES_ENUM	6	90		3 (0)	00:00:01
6	BUFFER SORT		1097	88857		23 (0)	00:00:01
7	TABLE ACCESS FULL	TEACHERS	1097	88857		4 (0)	00:00:01
8	VIEW		1668M	243G		238K (69)	00:00:10
9	SORT GROUP BY		1668M	441G		238K (69)	00:00:10
* 10	HASH JOIN		1668M	441G	146M	85682 (14)	00:00:04
11	VIEW		3005K	111M		388 (16)	00:00:01
12	HASH GROUP BY		3005K	111M		388 (16)	00:00:01
13	TABLE ACCESS INMEMORY FULL	POINTS	3005K	111M		387 (16)	00:00:01
* 14	HASH JOIN		5260K	1229M		2967 (7)	00:00:01
15	TABLE ACCESS FULL	GRADES_ENUM	6	90		3 (0)	00:00:01
* 16	HASH JOIN		5260K	1153M	4744K	2929 (6)	00:00:01
17	JOIN FILTER CREATE		:BF0000	21400	4493K	432 (23)	00:00:01
* 18	HASH JOIN		21400	4493K		432 (23)	00:00:01
19	JOIN FILTER CREATE		:BF0001	73	13140	8 (13)	00:00:01
* 20	HASH JOIN		73	13140		8 (13)	00:00:01
21	JOIN FILTER CREATE		:BF0002	3	69	3 (0)	00:00:01
* 22	TABLE ACCESS FULL	HOUSES	3	69		3 (0)	00:00:01
* 23	HASH JOIN		97	15229		5 (20)	00:00:01
24	JOIN FILTER CREATE		:BF0003	97	4559	3 (0)	00:00:01
25	TABLE ACCESS FULL	QUIDDITCH_TEAM_MEMBERS	97	4559		3 (0)	00:00:01
26	JOIN FILTER USE		:BF0002	6739	723K	1 (0)	00:00:01
27	JOIN FILTER USE		:BF0003	6739	723K	1 (0)	00:00:01
* 28	TABLE ACCESS INMEMORY FULL	STUDENTS	6739	723K		1 (0)	00:00:01
29	JOIN FILTER USE		:BF0001	2704K	90M	406 (20)	00:00:01
* 30	TABLE ACCESS INMEMORY FULL	POINTS	2704K	90M		406 (20)	00:00:01
31	JOIN FILTER USE		:BF0000	1656K	23M	105 (12)	00:00:01
* 32	TABLE ACCESS INMEMORY FULL	GRADES	1656K	23M		105 (12)	00:00:01

12.3.5 Best students without columnar store

Id	Operation	Name	Rows	Bytes	TempSpcl	Cost (%CPU)	Time
0	SELECT STATEMENT		10	1330		93493 (69)	00:00:04
1	SORT ORDER BY		10	1330		93493 (69)	00:00:04
* 2	VIEW		10	1330		93492 (69)	00:00:04
3	TEMP TABLE TRANSFORMATION						
4	LOAD AS SELECT (CURSOR DURATION MEMORY) SYS_TEMP_0FD9D664F_95BC7EA						
5	HASH GROUP BY		3174	133K		24 (9)	00:00:01
6	KEY VECTOR CREATE BUFFERED :KV0000		3174	133K		22 (0)	00:00:01
* 7	HASH JOIN		9346	355K		22 (0)	00:00:01
8	TABLE ACCESS FULL HOUSES		4	56		3 (0)	00:00:01
9	TABLE ACCESS FULL STUDENTS		9346	228K		19 (0)	00:00:01
* 10	WINDOW SORT PUSHED RANK		9346	857K	1048K	93468 (69)	00:00:04
11	HASH GROUP BY		9346	857K	1048K	93468 (69)	00:00:04
* 12	HASH JOIN		9346	857K		93059 (69)	00:00:04
13	TABLE ACCESS FULL SYS_TEMP_0FD9D664F_95BC7EA		3174	130K		7 (0)	00:00:01
14	VIEW	VW_VT_9E908A56	9346	474K		93052 (69)	00:00:04
15	VECTOR GROUP BY		9346	282K		93052 (69)	00:00:04
16	HASH GROUP BY		9346	282K		93052 (69)	00:00:04
17	KEY VECTOR USE :KV0000		631M	18G		33571 (14)	00:00:02
* 18	HASH JOIN		631M	15G	43M	33466 (14)	00:00:02
19	VIEW		1977K	20M		14229 (2)	00:00:01
20	HASH GROUP BY		1977K	35M	60M	14229 (2)	00:00:01
* 21	HASH JOIN		1977K	35M		2521 (3)	00:00:01
22	TABLE ACCESS FULL GRADES_ENUM		6	30		3 (0)	00:00:01
* 23	TABLE ACCESS FULL GRADES		1977K	26M		2505 (2)	00:00:01
* 24	TABLE ACCESS FULL POINTS		2986K	45M		8895 (1)	00:00:01

12.3.6 Best students with columnar store

Id	Operation	Name	Rows	Bytes	TempSpcl	Cost (%CPU)	Time
0	SELECT STATEMENT		10	1330		38839 (33)	00:00:02
1	SORT ORDER BY		10	1330		38839 (33)	00:00:02
* 2	VIEW		10	1330		38838 (33)	00:00:02
3	TEMP TABLE TRANSFORMATION						
4	LOAD AS SELECT (CURSOR DURATION MEMORY) SYS_TEMP_0FD9D6640_95BC7EA						
5	HASH GROUP BY						
6	KEY VECTOR CREATE BUFFERED :KV0000		5310	694K		6 (34)	00:00:01
7	HASH JOIN		5310	694K		4 (0)	00:00:01
8	TABLE ACCESS FULL HOUSES		409	9407		3 (0)	00:00:01
9	TABLE ACCESS INMEMORY FULL STUDENTS		5310	554K		1 (0)	00:00:01
* 10	WINDOW SORT PUSHED RANK						
11	HASH GROUP BY						
* 12	HASH JOIN						
13	VIEW	VW_VT_9E908A56					
14	VECTOR GROUP BY						
15	HASH GROUP BY						
16	KEY VECTOR USE :KV0000						
17	MERGE JOIN						
18	SORT JOIN						
19	VIEW						
20	HASH GROUP BY						
* 21	HASH JOIN						
22	TABLE ACCESS FULL GRADES_ENUM						
* 23	TABLE ACCESS INMEMORY FULL GRADES						
* 24	SORT JOIN						
* 25	TABLE ACCESS INMEMORY FULL POINTS						
26	TABLE ACCESS FULL SYS_TEMP_0FD9D6640_95BC7EA						

12.3.7 Raise grades without columnar store

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	MERGE STATEMENT		9346	37384		21196 (2)	00:00:01
1	MERGE	GRADES					
2	VIEW						
* 3	HASH JOIN		9346	520K		21196 (2)	00:00:01
* 4	VIEW		9346	255K		18686 (2)	00:00:01
* 5	WINDOW SORT PUSHED RANK		1115K	22M	34M	18686 (2)	00:00:01
* 6	HASH JOIN		1115K	22M		11451 (2)	00:00:01
* 7	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01
* 8	HASH JOIN		1115K	17M		11441 (2)	00:00:01
9	VIEW		7467	29868		8926 (2)	00:00:01
* 10	FILTER						
11	SORT GROUP BY		7467	320K		8926 (2)	00:00:01
* 12	HASH JOIN OUTER		149K	6416K		8917 (1)	00:00:01
* 13	HASH JOIN		467	16812		22 (0)	00:00:01
* 14	TABLE ACCESS FULL	HOUSES	1	14		3 (0)	00:00:01
15	TABLE ACCESS FULL	STUDENTS	9346	200K		19 (0)	00:00:01
16	TABLE ACCESS FULL	POINTS	2986K	22M		8875 (1)	00:00:01
* 17	TABLE ACCESS FULL	GRADES	1396K	15M		2506 (2)	00:00:01
18	TABLE ACCESS FULL	GRADES	1977K	54M		2497 (2)	00:00:01

12.3.8 Raise grades with columnar store

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	MERGE STATEMENT					110 (1)	00:00:33
1	MERGE	GRADES					
2	VIEW						
3	NESTED LOOPS					110 (1)	00:00:33
4	NESTED LOOPS					110 (1)	00:00:33
* 5	VIEW					110 (1)	00:00:33
* 6	WINDOW SORT PUSHED RANK					60M (1)	00:00:33
* 7	HASH JOIN					60M (1)	00:00:33
8	VIEW					751K (1)	00:00:01
* 9	FILTER						
10	SORT GROUP BY					751K (1)	00:00:01
* 11	HASH JOIN OUTER					751K (1)	00:00:01
12	JOIN FILTER CREATE	:BF0000				2370 (1)	00:00:01
* 13	HASH JOIN					2370 (1)	00:00:01
14	JOIN FILTER CREATE	:BF0001				1 (1)	23 (1)
15	TABLE ACCESS FULL	HOUSES				1 (1)	23 (1)
16	JOIN FILTER USE	:BF0001				9480 (1)	870K (1)
* 17	TABLE ACCESS INMEMORY FULL	STUDENTS				9480 (1)	870K (1)
18	JOIN FILTER USE	:BF0000				3005K (1)	74M (1)
* 19	TABLE ACCESS INMEMORY FULL	POINTS				3005K (1)	74M (1)
* 20	HASH JOIN					764K (1)	31M (1)
21	JOIN FILTER CREATE	:BF0002				4 (1)	60 (1)
* 22	TABLE ACCESS FULL	GRADES_ENUM				4 (1)	60 (1)
23	JOIN FILTER USE	:BF0002				1146K (1)	30M (1)
* 24	TABLE ACCESS INMEMORY FULL	GRADES				1146K (1)	30M (1)
* 25	INDEX UNIQUE SCAN	SYS_C009281				1 (1)	0 (0)
26	TABLE ACCESS BY INDEX ROWID	GRADES				1 (1)	75 (1)

12.3.9 Assign subjects without columnar store

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	INSERT STATEMENT		1017	52920		12318 (3)	00:00:01
1	LOAD TABLE CONVENTIONAL	STUDENTS SUBJECTS	1	1		1	1
2	HASH UNIQUE		1017	52920		12318 (3)	00:00:01
3	UNION-ALL		1	1		1	1
* 4	HASH JOIN		1016	26416		6155 (3)	00:00:01
* 5	HASH JOIN ANTI		93	1860		6152 (3)	00:00:01
* 6	TABLE ACCESS FULL	STUDENTS	9346	65422		19 (0)	00:00:01
7	VIEW	VW_NSO_1	25113	318K		6133 (3)	00:00:01
* 8	FILTER		1	1		1	1
9	SORT GROUP BY		25113	343K	45M	6133 (3)	00:00:01
* 10	HASH JOIN		1977K	26M		2511 (2)	00:00:01
11	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01
12	TABLE ACCESS FULL	GRADES	1977K	16M		2495 (2)	00:00:01
* 13	TABLE ACCESS FULL	SUBJECTS	76	456		3 (0)	00:00:01
* 14	HASH JOIN ANTI		1	44		6158 (3)	00:00:01
* 15	HASH JOIN		97	3007		25 (0)	00:00:01
16	MERGE JOIN CARTESIAN		97	2328		6 (0)	00:00:01
* 17	TABLE ACCESS FULL	SUBJECTS	1	20		3 (0)	00:00:01
18	BUFFER SORT		97	388		3 (0)	00:00:01
19	TABLE ACCESS FULL	QUIDDITCH_TEAM_MEMBERS	97	388		3 (0)	00:00:01
* 20	TABLE ACCESS FULL	STUDENTS	8011	56077		19 (0)	00:00:01
* 21	TABLE ACCESS FULL	SUBJECTS	1	20		3 (0)	00:00:01
22	VIEW	VW_NSO_2	25113	318K		6133 (3)	00:00:01
* 23	FILTER		1	1		1	1
24	SORT GROUP BY		25113	343K	45M	6133 (3)	00:00:01
* 25	HASH JOIN		1977K	26M		2511 (2)	00:00:01
26	TABLE ACCESS FULL	GRADES_ENUM	6	30		3 (0)	00:00:01
27	TABLE ACCESS FULL	GRADES	1977K	16M		2495 (2)	00:00:01

12.3.10 Assign subjects with columnar store

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	INSERT STATEMENT		102K	12M	2095 (14)	00:00:01
1	LOAD TABLE CONVENTIONAL	STUDENTS SUBJECTS	1	1	1	1
2	HASH UNIQUE		102K	12M	2095 (14)	00:00:01
3	UNION-ALL		1	1	1	1
* 4	HASH JOIN		102K	6533K	243 (59)	00:00:01
* 5	TABLE ACCESS INMEMORY FULL	SUBJECTS	76	1976	3 (0)	00:00:01
* 6	HASH JOIN ANTI		9480	361K	239 (60)	00:00:01
* 7	TABLE ACCESS INMEMORY FULL	STUDENTS	9480	240K	1 (0)	00:00:01
8	VIEW	VW_NSO_1	1656K	20M	227 (58)	00:00:01
* 9	FILTER		1	1	1	1
10	SORT GROUP BY		1656K	67M	227 (58)	00:00:01
* 11	HASH JOIN		1656K	67M	124 (23)	00:00:01
12	TABLE ACCESS FULL	GRADES_ENUM	6	90	3 (0)	00:00:01
13	TABLE ACCESS INMEMORY FULL	GRADES	1656K	44M	110 (16)	00:00:01
* 14	HASH JOIN ANTI		1	99	245 (58)	00:00:01
* 15	HASH JOIN RIGHT SEMI		97	8342	7 (0)	00:00:01
16	TABLE ACCESS FULL	QUIDDITCH_TEAM_MEMBERS	97	1261	3 (0)	00:00:01
17	MERGE JOIN CARTESIAN		474	34602	4 (0)	00:00:01
* 18	TABLE ACCESS INMEMORY FULL	SUBJECTS	1	47	3 (0)	00:00:01
19	BUFFER SORT		474	12324	1 (0)	00:00:01
* 20	TABLE ACCESS INMEMORY FULL	STUDENTS	474	12324	1 (0)	00:00:01
* 21	TABLE ACCESS INMEMORY FULL	SUBJECTS	1	47	3 (0)	00:00:01
22	VIEW	VW_NSO_2	1656K	20M	227 (58)	00:00:01
* 23	FILTER		1	1	1	1
24	SORT GROUP BY		1656K	67M	227 (58)	00:00:01
* 25	HASH JOIN		1656K	67M	124 (23)	00:00:01
26	TABLE ACCESS FULL	GRADES_ENUM	6	90	3 (0)	00:00:01
27	TABLE ACCESS INMEMORY FULL	GRADES	1656K	44M	110 (16)	00:00:01

12.3.11 Remove points without columnar store

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	DELETE STATEMENT		2078	64418	11495 (2)	00:00:01
1	DELETE	POINTS	1	1	1	1
* 2	HASH JOIN RIGHT SEMI		2078	64418	11495 (2)	00:00:01
3	VIEW	VW_NSO_1	13	169	2604 (6)	00:00:01
* 4	SORT GROUP BY		13	221	2604 (6)	00:00:01
5	VIEW		248	4216	2604 (6)	00:00:01
* 6	FILTER				1	1
7	HASH GROUP BY		248	5456	2604 (6)	00:00:01
* 8	HASH JOIN		4943	106K	2603 (6)	00:00:01
9	TABLE ACCESS FULL	GRADES_ENUM	6	30	3 (0)	00:00:01
* 10	TABLE ACCESS FULL	GRADES	4943	84031	2600 (6)	00:00:01
* 11	TABLE ACCESS FULL	POINTS	1494K	25M	8881 (1)	00:00:01

12.3.12 Remove points with columnar store

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	DELETE STATEMENT		1487K	92M		6819 (5)	00:00:01
1	DELETE	POINTS	1	1		1	1
* 2	HASH JOIN RIGHT SEMI		1487K	92M	31M	6819 (5)	00:00:01
3	VIEW	VW_NSO_1	1323K	16M		300 (68)	00:00:01
* 4	SORT GROUP BY		1323K	32M		300 (68)	00:00:01
5	VIEW		1323K	32M		300 (68)	00:00:01
* 6	FILTER				1	1	1
7	HASH GROUP BY		1323K	65M		300 (68)	00:00:01
* 8	HASH JOIN		1323K	65M		219 (57)	00:00:01
9	TABLE ACCESS FULL	GRADES_ENUM	6	90		3 (0)	00:00:01
* 10	TABLE ACCESS INMEMORY FULL	GRADES	1323K	46M		207 (56)	00:00:01
* 11	TABLE ACCESS INMEMORY FULL	POINTS	1487K	73M		416 (22)	00:00:01

12.4 Experiment 1: In-Memory Buffer Size Impact Analysis

We conducted the first experiment to evaluate the impact of different In-Memory buffer sizes on columnar storage performance. The experiment compared two configurations: **256MB In-Memory buffer**, and **1GB In-Memory buffer**, measuring performance **10 times** for each configuration to ensure statistical reliability.

12.4.1 Experiment 1 - System Configuration Parameters

Memory Component	256MB In-Memory (bytes)	1GB In-Memory (bytes)
Total System Global Area	2,147,482,608	2,147,480,608
Fixed Size	9,858,032	9,876,512
Variable Size	1,358,954,496	721,420,288
Database Buffers	503,316,480	339,738,624
Redo Buffers	6,918,144	2,703,360
In-Memory Area	268,435,456	1,073,741,824
In-Memory Size (MB)	256	1024

12.4.2 Experiment 1 - Performance Results

As it turns out - a bigger in-memory buffer doesn't always provide improved performance. In fact, usually both running time and query plan cost are better for the smaller buffer, at least on our computers. Below we show the results of our experiment (we excluded two queries for which the difference was too small).

Procedure name	Min execution time (ms)				Max execution time (ms)				Avg execution time (ms)				Query plan cost			
	256 MB	1 GB	Gain	Profit	256 MB	1 GB	Gain	Profit	256 MB	1 GB	Gain	Profit	256 MB	1 GB	Gain	Profit
Grades analysis	723	664	-59	-8.2 %	868	1598	730	84.1 %	778.6	1194.8	416.2	53.5 %	2446 K	3347 K	901 K	36.8 %
Best students	438	487	49	11.2 %	548	836	288	52.6 %	497.8	672.9	175.1	35.2 %	44600	42704	-1896	-4.3 %
Raise grades	2206	1469	-737	-33.4 %	2442	5376	2934	120.1 %	2306.6	4162.9	1856.3	80.5 %	498 K	746 K	248 K	49.8 %
Remove points	3587	5414	1827	50.9 %	5910	6985	1075	18.2 %	4046.1	6399.2	2353.1	58.2 %	6479	559	-5920	-91.4 %
Whole workload	6954	8034	1080	15.5 %	9768	14795	5027	51.5 %	7629.1	12429.8	4800.7	62.9 %	2995 K	4136 K	1141 K	38.1 %

12.4.3 Experiment 1 - Performance Analysis

The results demonstrate that Oracle's In-Memory optimization effectiveness varies significantly with buffer size allocation:

- **256 MB Configuration:** Overall better than 1 GB buffer, both in running time and in query plan cost. Better choice for our scenario.
- **1 GB Configuration:** Significantly cheaper query plan only for query that involves deleting many rows. For the whole workload however on average it runs 62.9% slower and has 38.1% bigger cost.

12.5 Query plans comparison

Implementing columnar storage dramatically changed query execution plans. The Oracle optimizer leveraged in-memory columnar operations, vector processing, and specialized algorithms for analytical workloads.

12.5.1 Grades analysis - 256MB In-Memory

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT		106	37948		2446K (66)	00:01:36
1	SORT ORDER BY		106	37948		2446K (66)	00:01:36
* 2	HASH JOIN		106	37948		2446K (66)	00:01:36
3	TABLE ACCESS INMEMORY FULL	SUBJECTS	76	6004		3 (0)	00:00:01
* 4	HASH JOIN		106	29574		2446K (66)	00:01:36
5	TABLE ACCESS INMEMORY FULL	SUBJECTS	76	6004		3 (0)	00:00:01
* 6	HASH JOIN		106	21200		2446K (66)	00:01:36
7	TABLE ACCESS FULL	GRADES_ENUM	6	90		3 (0)	00:00:01
* 8	HASH JOIN		106	19610		2446K (66)	00:01:36
9	TABLE ACCESS FULL	GRADES_ENUM	6	90		3 (0)	00:00:01
10	VIEW		106	18020		2446K (66)	00:01:36
11	SORT GROUP BY		106	33602		2446K (66)	00:01:36
* 12	HASH JOIN RIGHT OUTER		13G	4118G		1039K (19)	00:00:41
13	TABLE ACCESS FULL	QUIDDITCH_TEAM_MEMBERS	1	47		3 (0)	00:00:01
* 14	HASH JOIN		13G	3508G	101M	946K (10)	00:00:37
15	VIEW					1666K	82M
16	HASH GROUP BY					231 (58)	00:00:01
* 17	HASH JOIN					1666K	151M
18	TABLE ACCESS FULL	GRADES_ENUM	6	90		231 (23)	00:00:01
19	VIEW	VW_GBF_7	1666K	127M		3 (0)	00:00:01
20	HASH GROUP BY					113 (16)	00:00:01
21	TABLE ACCESS INMEMORY FULL	GRADES	1666K	44M		113 (16)	00:00:01
* 22	HASH JOIN		76M	15G	39M	13423 (5)	00:00:01
23	TABLE ACCESS INMEMORY FULL	GRADES	1666K	20M		102 (7)	00:00:01
* 24	HASH JOIN		422K	82M		6489 (2)	00:00:01
25	JOIN FILTER CREATE	:BF0000	2292	380K		138 (27)	00:00:01
* 26	HASH JOIN ANTI		2292	380K		138 (27)	00:00:01
27	JOIN FILTER CREATE	:BF0001	2292	351K		5 (20)	00:00:01
* 28	HASH JOIN		2292	351K		5 (20)	00:00:01
29	JOIN FILTER CREATE	:BF0002	1	47		3 (0)	00:00:01
* 30	TABLE ACCESS FULL	HOUSES	1	47		3 (0)	00:00:01
31	JOIN FILTER USE	:BF0002	9167	984K		1 (0)	00:00:01
* 32	TABLE ACCESS INMEMORY FULL	STUDENTS	9167	984K		1 (0)	00:00:01
33	VIEW	VW_NSO_1	159K	2028K		133 (27)	00:00:01
* 34	HASH JOIN		159K	6085K		133 (27)	00:00:01
35	JOIN FILTER CREATE	:BF0003	1	15		3 (0)	00:00:01
* 36	TABLE ACCESS FULL	GRADES_ENUM	1	15		3 (0)	00:00:01
37	TABLE ACCESS BY INDEX ROWID GRADES_ENUM		1	15		2 (0)	00:00:01
* 38	INDEX UNIQUE SCAN	SYS_C009232	1	1		1 (0)	00:00:01
39	TABLE ACCESS BY INDEX ROWID GRADES_ENUM		1	15		2 (0)	00:00:01
* 40	INDEX UNIQUE SCAN	SYS_C009232	1	1		1 (0)	00:00:01
41	JOIN FILTER USE	:BF0001	1638K	37M		119 (21)	00:00:01
42	JOIN FILTER USE	:BF0003	1638K	37M		119 (21)	00:00:01
* 43	TABLE ACCESS INMEMORY FULL GRADES		1638K	37M		119 (21)	00:00:01
44	JOIN FILTER USE	:BF0000	1755K	58M		6339 (1)	00:00:01
* 45	TABLE ACCESS INMEMORY FULL	POINTS	1755K	58M		6339 (1)	00:00:01

12.5.2 Grades analysis - 1GB In-Memory

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT		106	37948		3347K(100)	00:02:11
1	SORT ORDER BY		106	37948		3347K(100)	00:02:11
* 2	HASH JOIN		106	37948		3347K(100)	00:02:11
3	TABLE ACCESS INMEMORY FULL	SUBJECTS	76	6004		3 (0)	00:00:01
* 4	HASH JOIN		106	29574		3347K(100)	00:02:11
5	TABLE ACCESS INMEMORY FULL	SUBJECTS	76	6004		3 (0)	00:00:01
* 6	HASH JOIN		106	21200		3347K(100)	00:02:11
7	TABLE ACCESS FULL	GRADES_ENUM	6	90		3 (0)	00:00:01
* 8	HASH JOIN		106	19610		3347K(100)	00:02:11
9	TABLE ACCESS FULL	GRADES_ENUM	6	90		3 (0)	00:00:01
10	VIEW		106	18020		3347K(100)	00:02:11
11	SORT GROUP BY		106	33602		3347K(100)	00:02:11
* 12	HASH JOIN RIGHT OUTER		21G	6277G		1163K (98)	00:00:46
13	TABLE ACCESS FULL	HOUSES	1	47		3 (0)	00:00:01
* 14	HASH JOIN		85G	20T		598K (95)	00:00:24
15	VIEW		467M	11G		28849 (13)	00:00:02
16	HASH GROUP BY		1666K	352M		1367 (14)	00:00:01
* 17	HASH JOIN		1666K	332M		468 (30)	00:00:01
18	TABLE ACCESS FULL	GRADES_ENUM	6	90		464 (30)	00:00:01
19	VIEW	VW_GBF_7	1666K	127M		113 (16)	00:00:01
20	HASH GROUP BY		1666K	82M		231 (58)	00:00:01
21	TABLE ACCESS INMEMORY FULL	GRADES	1666K	151M		231 (58)	00:00:01
* 22	HASH JOIN		1666K	151M		127 (23)	00:00:01
23	TABLE ACCESS FULL	GRADES_ENUM	6	90		3 (0)	00:00:01
24	VIEW	VW_GBF_7	1666K	127M		113 (16)	00:00:01
25	HASH GROUP BY		1666K	44M		113 (16)	00:00:01
26	TABLE ACCESS INMEMORY FULL	GRADES	1666K	44M		112 (16)	00:00:01
27	SORT JOIN		9167	984K	2344K	233 (1)	00:00:01
28	TABLE ACCESS INMEMORY FULL	STUDENTS	9167	984K		1 (0)	00:00:01
29	SORT JOIN		97	4559		4 (25)	00:00:01
30	TABLE ACCESS FULL	QUIDDITCH_TEAM_MEMBERS	97	4559		3 (0)	00:00:01
31	SORT UNIQUE		159K	2028K	6280K	899 (6)	00:00:01
32	VIEW	VW_NSO_1	159K	2028K		133 (27)	00:00:01
33	HASH JOIN		159K	6085K		133 (27)	00:00:01
34	JOIN FILTER CREATE	:BF0000	1	15		3 (0)	00:00:01
35	TABLE ACCESS FULL	GRADES_ENUM	1	15		3 (0)	00:00:01
36	TABLE ACCESS BY INDEX ROWID GRADES_ENUM		1	15		2 (0)	00:00:01
* 37	INDEX UNIQUE SCAN	SYS_C009232	1	1		1 (0)	00:00:01
38	TABLE ACCESS BY INDEX ROWID GRADES_ENUM		1	15		2 (0)	00:00:01
* 39	INDEX UNIQUE SCAN	SYS_C009232	1	1		1 (0)	00:00:01
40	JOIN FILTER USE	:BF0000	1638K	37M		119 (21)	00:00:01
* 41	TABLE ACCESS INMEMORY FULL	GRADES	1638K	37M		119 (21)	00:00:01
42	SORT JOIN		2572K	85M	236M	24391 (2)	00:00:01
* 43	TABLE ACCESS INMEMORY FULL POINTS		2572K	85M		404 (18)	00:00:01
44	SORT JOIN		1666K	20M	63M	8082 (2)	00:00:01
* 45	TABLE ACCESS INMEMORY FULL	GRADES	1666K	20M		102 (7)	00:00:01

12.5.3 Best students display - 256MB In-Memory

=====Best students display=====										
Plan hash value: 3211729349										
Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time			
0	SELECT STATEMENT		10	1570		44600 (29)	00:00:02			
1	SORT ORDER BY		10	1570		44600 (29)	00:00:02			
* 2	VIEW		10	1570		44599 (29)	00:00:02			
3	TEMP TABLE TRANSFORMATION									
4	LOAD AS SELECT (CURSOR DURATION MEMORY) SYS_TEMP_0FD9D6611_1836B28									
5	HASH GROUP BY		5310	819K		6 (34)	00:00:01			
6	KEY VECTOR CREATE BUFFERED	:KV0000	5310	819K		4 (0)	00:00:01			
* 7	HASH JOIN		5310	798K		4 (0)	00:00:01			
8	TABLE ACCESS FULL	HOUSES	409	19223		3 (0)	00:00:01			
9	TABLE ACCESS INMEMORY FULL	STUDENTS	5310	554K		1 (0)	00:00:01			
* 10	WINDOW SORT PUSHED RANK		319K	57M	64M	44594 (29)	00:00:02			
11	HASH GROUP BY		319K	57M	64M	44594 (29)	00:00:02			
* 12	HASH JOIN		319K	57M	18223	69K	00:00:01			
13	VIEW	VW_VT_9E908A56	6019	358K		18188 (69)	00:00:01			
14	VECTOR GROUP BY		6019	346K		18188 (69)	00:00:01			
15	HASH GROUP BY		6019	346K		18188 (69)	00:00:01			
16	KEY VECTOR USE	:KV0000	6018	346K		18186 (69)	00:00:01			
17	MERGE JOIN		6019	323K		5924 (4)	00:00:01			
18	SORT JOIN		1655K	31M		236 (59)	00:00:01			
19	VIEW		1655K	31M		236 (59)	00:00:01			
20	HASH GROUP BY		1655K	61M		236 (59)	00:00:01			
* 21	HASH JOIN		1655K	61M		133 (27)	00:00:01			
22	TABLE ACCESS FULL	GRADES_ENUM	6	90		3 (0)	00:00:01			
* 23	TABLE ACCESS INMEMORY FULL	GRADES	1655K	37M		119 (21)	00:00:01			
* 24	SORT JOIN		6019	205K		5688 (2)	00:00:01			
* 25	TABLE ACCESS INMEMORY FULL	POINTS	6019	205K		5687 (2)	00:00:01			
26	TABLE ACCESS FULL	SYS_TEMP_0FD9D6611_1836B28	5310	663K		33 (0)	00:00:01			

Predicate Information (identified by operation id):

```

2 - filter("RANK"><=10)
7 - access("$"."HOUSE_ID"="H"."ID")
10 - filter(ROW_NUMBER() OVER ( ORDER BY ROUND(DECODE(NVL(SUM("ITEM_5"),0),0,TO_NUMBER(NULL),SUM("ITEM_4")/NVL(SUM("ITEM_5"),0)),2) DESC ,NVL(SUM("ITEM_3"),0) DESC )<=10)
12 - access("ITEM_11"="INTERNAL_FUNCTION("CO"))
21 - access("GRADES"."VALUE"="GRADES_ENUM"."SYMBOL")
23 - inmemory("GRADES"."AWARD_DATE">>=TO_DATE(' 2023-09-02 00:00:00 ', 'yyyy-mm-dd hh24:mi:ss') AND
      "GRADES"."AWARD_DATE"=<TO_DATE(' 2024-06-28 00:00:00 ', 'yyyy-mm-dd hh24:mi:ss'))
      filter("GRADES"."AWARD_DATE">>=TO_DATE(' 2023-09-02 00:00:00 ', 'yyyy-mm-dd hh24:mi:ss') AND
      "GRADES"."AWARD_DATE"=<TO_DATE(' 2024-06-28 00:00:00 ', 'yyyy-mm-dd hh24:mi:ss'))
24 - access("P"."STUDENT_ID"="MA"."STUDENT_ID")
      filter("P"."STUDENT_ID"="MA"."STUDENT_ID")
25 - inmemory("P"."AWARD_DATE"><TO_DATE(' 2024-06-28 00:00:00 ', 'yyyy-mm-dd hh24:mi:ss') AND
      "P"."AWARD_DATE">>=TO_DATE(' 2023-09-02 00:00:00 ', 'yyyy-mm-dd hh24:mi:ss'))
      filter("P"."AWARD_DATE"><TO_DATE(' 2024-06-28 00:00:00 ', 'yyyy-mm-dd hh24:mi:ss') AND
      "P"."AWARD_DATE">>=TO_DATE(' 2023-09-02 00:00:00 ', 'yyyy-mm-dd hh24:mi:ss'))

```

Note

- dynamic statistics used: dynamic sampling (level=2)
- vector transformation used for this statement

12.5.4 Best students display - 1GB In-Memory

=====Best students display=====										
Plan hash value: 355475729										
Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time			
0	SELECT STATEMENT		10	1570		42704 (31)	00:00:02			
1	SORT ORDER BY		10	1570		42704 (31)	00:00:02			
* 2	VIEW		10	1570		42703 (31)	00:00:02			
3	TEMP TABLE TRANSFORMATION									
4	LOAD AS SELECT (CURSOR DURATION MEMORY) SYS_TEMP_0FD9D6610_187B43B									
5	HASH GROUP BY		5310	819K		6 (34)	00:00:01			
6	KEY VECTOR CREATE BUFFERED	:KV0000	5310	819K		4 (0)	00:00:01			
* 7	HASH JOIN		5310	798K		4 (0)	00:00:01			
8	TABLE ACCESS FULL	HOUSES	409	19223		3 (0)	00:00:01			
9	TABLE ACCESS INMEMORY FULL	STUDENTS	5310	554K		1 (0)	00:00:01			
* 10	WINDOW SORT PUSHED RANK		319K	57M	64M	44594 (29)	00:00:02			
11	HASH GROUP BY		319K	57M	64M	44594 (29)	00:00:02			
* 12	HASH JOIN		319K	57M	18223	69K	00:00:01			
13	VIEW	VW_VT_9E908A56	6019	358K		18188 (69)	00:00:01			
14	VECTOR GROUP BY		6019	346K		18188 (69)	00:00:01			
15	HASH GROUP BY		6019	346K		18188 (69)	00:00:01			
16	KEY VECTOR USE	:KV0000	6018	346K		18186 (69)	00:00:01			
* 17	HASH JOIN		6019	323K		5924 (4)	00:00:01			
18	JOIN FILTER CREATE	:BF0000	1655K	31M		229K	00:00:01			
* 19	TABLE ACCESS INMEMORY FULL	POINTS	1655K	239K		404 (18)	00:00:01			
20	VIEW		1655K	31M		236 (59)	00:00:01			
21	HASH GROUP BY		1655K	61M		236 (59)	00:00:01			
* 22	HASH JOIN		1655K	61M		133 (27)	00:00:01			
23	TABLE ACCESS FULL	GRADES_ENUM	6	90		3 (0)	00:00:01			
* 24	JOIN FILTER USE	:BF0000	1655K	37M		119 (21)	00:00:01			
* 25	TABLE ACCESS INMEMORY FULL	GRADES	1655K	37M		119 (21)	00:00:01			
26	TABLE ACCESS FULL	SYS_TEMP_0FD9D6610_187B43B	5310	663K		33 (0)	00:00:01			

Predicate Information (identified by operation id):

```

2 - filter("RANK"><=10)
7 - access("$"."HOUSE_ID"="H"."ID")
10 - filter(ROW_NUMBER() OVER ( ORDER BY ROUND(DECODE(NVL(SUM("ITEM_5"),0),0,TO_NUMBER(NULL),SUM("ITEM_4")/NVL(SUM("ITEM_5"),0)),2) DESC ,NVL(SUM("ITEM_3"),0) DESC )<=10)
12 - access("ITEM_11"="INTERNAL_FUNCTION("CO"))
21 - access("P"."AWARD_DATE">>=TO_DATE(' 2024-06-28 00:00:00 ', 'yyyy-mm-dd hh24:mi:ss') AND
      "P"."AWARD_DATE"=<TO_DATE(' 2023-09-02 00:00:00 ', 'yyyy-mm-dd hh24:mi:ss'))
      filter("P"."AWARD_DATE">>=TO_DATE(' 2023-09-02 00:00:00 ', 'yyyy-mm-dd hh24:mi:ss') AND
      "P"."AWARD_DATE"=<TO_DATE(' 2024-06-28 00:00:00 ', 'yyyy-mm-dd hh24:mi:ss'))
22 - access("P"."STUDENT_ID"="MA"."STUDENT_ID")
      filter("P"."STUDENT_ID"="MA"."STUDENT_ID")
25 - inmemory("P"."AWARD_DATE">>=TO_DATE(' 2024-06-28 00:00:00 ', 'yyyy-mm-dd hh24:mi:ss') AND
      "P"."AWARD_DATE"=<TO_DATE(' 2023-09-02 00:00:00 ', 'yyyy-mm-dd hh24:mi:ss'))
      filter("P"."AWARD_DATE">>=TO_DATE(' 2024-06-28 00:00:00 ', 'yyyy-mm-dd hh24:mi:ss') AND
      "P"."AWARD_DATE"=<TO_DATE(' 2023-09-02 00:00:00 ', 'yyyy-mm-dd hh24:mi:ss'))
      SYS_OP_BLOOM_FILTER(:BF0000, "P"."AWARD_DATE", "P"."STUDENT_ID")
26 - inmemory("P"."AWARD_DATE">>=TO_DATE(' 2023-09-02 00:00:00 ', 'yyyy-mm-dd hh24:mi:ss') AND
      "P"."AWARD_DATE"=<TO_DATE(' 2024-06-28 00:00:00 ', 'yyyy-mm-dd hh24:mi:ss')) AND
      "P"."AWARD_DATE">>=TO_DATE(' 2023-09-02 00:00:00 ', 'yyyy-mm-dd hh24:mi:ss') AND
      "P"."AWARD_DATE"=<TO_DATE(' 2024-06-28 00:00:00 ', 'yyyy-mm-dd hh24:mi:ss')) AND
      SYS_OP_BLOOM_FILTER(:BF0000, "P"."AWARD_DATE", "P"."STUDENT_ID")

```

Note

- dynamic statistics used: dynamic sampling (level=2)
- vector transformation used for this statement

12.6 Experiment 2: Compression Methods and Storage Strategy Comparison

We conducted a comprehensive second experiment to evaluate different Oracle In-Memory compression methods and storage strategies. This experiment compared five distinct configurations: **clean database** (no In-Memory), **FOR CAPACITY HIGH** (maximum compression), **FOR QUERY HIGH** (query-optimized), **FOR CAPACITY LOW** (balanced approach), and **5-columnar storage** (full In-Memory implementation), measuring performance **10 times** for each configuration.

12.6.1 Experiment 2 - In-Memory Configuration Details

In this experiment, we applied In-Memory columnar storage **selectively** to the same three critical analytical tables across all test variants, while keeping remaining tables in traditional row storage. The target tables for In-Memory optimization were:

- **GRADES** - Primary fact table containing millions of grade records with high analytical query frequency
- **POINTS** - Secondary fact table with student achievement data requiring frequent aggregations
- **STUDENTS** - Central dimensional table serving as the primary join point for analytical operations

We tested different combinations of Oracle In-Memory parameters to evaluate compression effectiveness and query performance trade-offs:

1. FOR CAPACITY HIGH Configuration:

```
CREATE TABLE students (...)  
INMEMORY PRIORITY HIGH MEMCOMPRESS FOR CAPACITY HIGH;
```

```
CREATE TABLE grades (...)  
INMEMORY PRIORITY CRITICAL MEMCOMPRESS FOR CAPACITY HIGH;
```

```
CREATE TABLE points (...)  
INMEMORY PRIORITY CRITICAL MEMCOMPRESS FOR CAPACITY HIGH;
```

2. FOR QUERY HIGH Configuration:

```
CREATE TABLE students (...)  
INMEMORY PRIORITY HIGH MEMCOMPRESS FOR QUERY HIGH;
```

```
CREATE TABLE grades (...)  
INMEMORY PRIORITY CRITICAL MEMCOMPRESS FOR QUERY HIGH;
```

```
CREATE TABLE points (...)  
INMEMORY PRIORITY CRITICAL MEMCOMPRESS FOR QUERY HIGH;
```

3. FOR CAPACITY LOW Configuration:

```
CREATE TABLE students (...)  
INMEMORY PRIORITY LOW MEMCOMPRESS FOR CAPACITY LOW;
```

```
CREATE TABLE grades (...)  
INMEMORY PRIORITY MEDIUM MEMCOMPRESS FOR CAPACITY LOW;
```

```
CREATE TABLE points (...)  
INMEMORY PRIORITY MEDIUM MEMCOMPRESS FOR CAPACITY LOW;
```

The parameter combinations control two critical aspects of In-Memory behavior:

- **PRIORITY** (HIGH/CRITICAL/MEDIUM/LOW): Determines loading sequence and memory allocation precedence during population
- **MEMCOMPRESS** (FOR CAPACITY HIGH/FOR QUERY HIGH/FOR CAPACITY LOW): Controls compression algorithm selection balancing memory usage versus query performance

12.6.2 Experiment 2 - System Configuration Parameters

All configurations in this experiment utilized a consistent memory allocation optimized for compression method comparison:

Memory Component	Configuration (bytes)
Total System Global Area	1,333,788,432
Fixed Size	9,867,024
Variable Size	515,899,392
Database Buffers	268,435,456
In-Memory Area	536,870,912
In-Memory Size (MB)	512
Total SGA Size (GB)	1.24

12.6.3 Experiment 2 - Performance Results

The compression method comparison revealed significant performance variations across different In-Memory optimization strategies. All 25 test cases were executed 10 times each to ensure statistical reliability. Results are grouped by procedure for easy comparison across columnar storage technologies:

Procedure	Columnar Technology	Min (ms)	Max (ms)	Avg (ms)	vs Clean (%)	Query Cost	Rank
Grades analysis	Clean Database (baseline)	752	1011	851.4	baseline	4851K	baseline
	FOR CAPACITY HIGH	1469	1801	1638.2	+92.4%	3801K	-21.6%
	FOR QUERY HIGH	752	882	818.8	-3.8%	1267K	-73.9%
	FOR CAPACITY LOW	759	909	832.2	-2.3%	6875K	+41.7%
Best students	Clean Database (baseline)	492	571	535.8	baseline	53205	baseline
	FOR CAPACITY HIGH	519	612	576.8	+7.6%	45698	-14.1%
	FOR QUERY HIGH	410	540	476.9	-11.0%	25348	-52.4%
	FOR CAPACITY LOW	471	561	524.4	-2.1%	44668	-16.0%
Raise grades	Clean Database (baseline)	2200	2745	2459.1	baseline	1243K	baseline
	FOR CAPACITY HIGH	2443	2734	2619.3	+6.5%	1124K	-9.6%
	FOR QUERY HIGH	2346	2754	2587.4	+5.2%	532K	-57.2%
	FOR CAPACITY LOW	2465	2786	2581.4	+5.0%	1662K	+33.7%
Assign subjects	Clean Database (baseline)	3587	3987	3853.3	baseline	7421	baseline
	FOR CAPACITY HIGH	3447	3897	3636.1	-5.6%	2691	-63.7%
	FOR QUERY HIGH	3687	4100	3847.3	-0.2%	7181	-3.2%
	FOR CAPACITY LOW	3568	3971	3765.2	-2.3%	7443	+0.3%
Remove points	Clean Database (baseline)	3492	5368	4340.5	baseline	15863	baseline
	FOR CAPACITY HIGH	3683	4333	4004.2	-7.7%	9021	-43.1%
	FOR QUERY HIGH	3695	5397	4342.6	+0.05%	11347	-28.5%
	FOR CAPACITY LOW	3811	4458	4077.9	-6.1%	16329	+2.9%

12.6.4 Experiment 2 - Results Summary

Performance comparison results:

- **FOR QUERY HIGH:** Best overall - 73.9% cost reduction for Grades analysis, minimal execution time impact
- **FOR CAPACITY HIGH:** Excellent for DML operations - 63.7% cost reduction for Assign subjects
- **FOR CAPACITY LOW:** Balanced approach with consistent modest improvements